

# 1 Biomodelling.jl 2.0: mechanistic simulation of heritable gene-expression 2 states, cell division and drug response for single-cell benchmarking and 3 persister biology

4 Ayoub Lasri<sup>1</sup> Marc Sturrock<sup>1</sup>

5 <sup>1</sup> Department of Physiology and Medical Physics, Royal College of Surgeons in Ireland, Dublin,  
6 Ireland

## 7 **Abstract**

8 Non-genetic heterogeneity, and in particular expression states that are inherited across cell divisions, decides  
9 which cancer cells survive drug treatment, yet the simulators used to benchmark single-cell methods model  
10 neither cell growth nor division, and therefore cannot generate such states from first principles. We present  
11 Biomodelling.jl 2.0, a rewrite of a Julia framework for stochastic simulation of gene regulatory networks  
12 inside growing, dividing and drug-treated cell populations. The framework couples exact, hybrid and adaptive  
13 stochastic kinetics to exponential growth, volume-scaled transcription, gene replication, division with binomial  
14 partitioning and promoter inheritance, and population control; adds dose schedules, state-dependent death,  
15 growth inhibition and genetic perturbations; records every cell's lineage; converts molecule counts into scRNA-  
16 seq, smFISH and time-lapse observations; and provides likelihood-free inference. We validate the engine  
17 against exact stationary distributions and an independent implementation, show that growth and division  
18 alone reproduce published results on cell-size scaling, partitioning noise, lineage-versus-population statistics  
19 and cell-cycle-dependent bursting, and demonstrate that a single resistance gene with slow promoter switching  
20 reproduces the signatures of drug-tolerant persisters reported in recent experiments and models: dose-dependent  
21 decay with unchanged division times, correlated fates of related cells, preserved or collapsed barcode diversity,  
22 loss of resistance when memory is disrupted during exposure, an intermediate-dose optimum for drug-induced  
23 tolerance, and stable inheritance of a selected phenotype. Used as a ground-truth generator, the framework  
24 evaluates memory-gene detection, network inference under the confounds of growth and sequencing, imputation  
25 and perturbation-prediction baselines; used as an inference engine, it recovers kinetic and drug parameters  
26 from snapshot and lineage data and quantifies the bias of ignoring division. Biomodelling.jl 2.0 is open source  
27 (MIT) at <https://github.com/ayoublasri/Biomodelling.jl>.

## 28 **Introduction**

29 Single-cell technologies have made non-genetic heterogeneity a central object of study in cell biology and drug  
30 discovery. Isogenic cells differ in their transcriptional states, these states are often inherited across several divisions,  
31 and the inherited states decide cell fate under stress: drug-tolerant persister cells that survive targeted or cytotoxic  
32 therapy arise from pre-existing or drug-induced expression states rather than from mutations (Shaffer et al. 2017,  
33 2020; Harmange et al. 2023; Iyer et al. 2025; Lin et al. 2026). Lineage-resolved single-cell experiments now  
34 quantify this memory directly, showing that the fates of sisters and cousins are correlated, that clonal barcode  
35 diversity is preserved or collapsed depending on whether tolerance pre-exists or is induced, and that the timescale

of expression fluctuations relative to the cell cycle determines which genes carry memory (Iyer et al. 2025; Lin et al. 2026; Ghosh et al. 2025). In parallel, mechanistic analyses of single-cell RNA-sequencing data have established that transcriptional kinetics can only be inferred correctly when gene copy number, cell age and partitioning at division are modelled explicitly (Sukys and Grima 2025; Zhang et al. 2025), and evaluations of “virtual cell” models for perturbation prediction have found that deep-learning approaches still fail to beat simple baselines, largely because they lack mechanism and causal structure (Ahlmann-Eltze et al. 2025; Vollenweider and Bühlmann 2026; Dibaeinia et al. 2026).

These developments create a demand that existing simulators do not meet. Simulators used to benchmark single-cell methods, such as SERGIO, dyngen, scMultiSim and GRouNdGAN (Dibaeinia and Sinha 2020; Cannoodt et al. 2021; Li et al. 2025; Zinati et al. 2024), model regulation and technical noise but not cell growth, volume-dependent transcription, gene replication or the partitioning of molecules at division; lineage simulators such as TedSim (Pan et al. 2022) generate division trees but attach phenomenological, largely Markovian expression to them. As a consequence none of them can produce expression memory, cell-size scaling, cell-cycle-dependent bursting or persister dynamics from first principles, and none provides the ground truth needed to test the methods that are now being developed for exactly those phenomena. Stochastic simulators of growing and dividing cells exist (Bertaux et al. 2018; Piho and Thomas 2026) but stop at the single-cell state: they do not provide drug response, lineage statistics, observation models, inference or benchmark generation.

Biomodelling.jl was introduced as a generator of synthetic scRNA-seq data from stochastic gene regulatory networks in growing and dividing cells, and was used to benchmark imputation methods for network inference (Lasri et al. 2022). Its distinguishing feature, the coupling of stochastic reaction kinetics to exponential growth, volume-scaled transcription and binomial partitioning at division, is precisely the mechanism behind the heritable states discussed above. Here we present version 2.0, a rewrite that turns this feature into a complete framework for the mechanistic study of heritable expression states, drug tolerance and single-cell benchmarking. The package adds (i) a typed model layer with mass-action, Hill and custom kinetics and explicit promoter states; (ii) exact, hybrid and adaptive stochastic kernels validated against analytic laws and an independent implementation; (iii) a population layer with sizer, adder and timer division control, binomial or over-dispersed partitioning, promoter inheritance, gene replication and constant-size, free-growth or logistic population control; (iv) a drug and perturbation layer with dose schedules, pharmacokinetics, state-dependent death, growth inhibition, drug-induced rate changes and genetic perturbations; (v) complete lineage records with heritability, autocorrelation, fluctuation-test and clonal-memory statistics; (vi) observation models for scRNA-seq, smFISH and time-lapse reporters with AnnData output; and (vii) likelihood-free inference and the exact telegraph likelihood.

We validate the engine against exact stationary distributions and against JumpProcesses.jl, show that growth, division and partitioning alone reproduce published results on cell-size scaling, partitioning noise, lineage-versus-population statistics and cell-cycle-dependent bursting, and demonstrate that a single resistance gene with slow promoter switching reproduces the quantitative signatures of drug-tolerant persisters reported in recent experimental and theoretical work: dose-dependent population decay with unchanged single-cell division times, correlated fates of related cells, preserved or collapsed barcode diversity, loss of resistance when memory is disrupted during exposure, an intermediate-dose optimum for drug-induced tolerance, and stable inheritance of a selected phenotype. We then use the framework as a ground-truth generator to evaluate memory-gene detection, network inference under the confounds of growth and sequencing, imputation, and perturbation-prediction baselines, and finally show that the same forward model supports parameter inference from snapshot and lineage data and quantifies the bias incurred by ignoring division.

## Results

### A layered framework for stochastic simulation in growing, dividing and drug-treated populations

Biomodelling.jl 2.0 is organised in layers (Figure 1). A model is a list of reactions with mass-action, Hill or custom kinetics and named parameters; promoter states are explicit species grouped so that they are inherited at division and doubled at replication. Four kernels advance a cell’s state at fixed volume: the exact direct method, fixed-step tau-leaping, a hybrid scheme that falls back to the exact method whenever a leap would produce a negative count, and adaptive tau-leaping with critical-reaction handling. The population layer grows each cell’s volume exponentially, scales propensities with volume so that concentrations are invariant, replicates genes at a set point of the cycle, divides cells by a sizer, adder or timer rule with binomial (or over-dispersed) partitioning of molecules, and controls the population size (constant size by Moran-like replacement, free growth, or logistic). A drug and perturbation layer applies dose schedules, state-dependent death, growth inhibition, drug-induced rate changes and genetic perturbations. Every cell’s ancestry, fate and state at birth and division are recorded in a lineage table, from which heritability, autocorrelation, memory timescales, fluctuation-test and clonal-memory statistics are computed. Observation models convert true counts into scRNA-seq, smFISH or time-lapse measurements, and the same forward model drives likelihood-free inference. Table 1 compares these capabilities with those of the simulators most used to benchmark single-cell methods.

Table 1. Capabilities of single-cell simulators (✓ native; ○ partial or phenomenological; – absent). ABM.jl, AgentBasedModeling.jl; v1.x and v2.0, Biomodelling.jl versions.

| Capability                                           | SERGIO | dyngen | scMulti-Sim | GRouNd-GAN | TedSim | Lineage-Sim | ABM.jl | v1.x | v2.0 |
|------------------------------------------------------|--------|--------|-------------|------------|--------|-------------|--------|------|------|
| Mechanistic reaction kinetics (SSA / tau-leaping)    | ○      | ✓      | ○           | –          | –      | –           | ✓      | ✓    | ✓    |
| Gene regulatory networks with known ground truth     | ✓      | ✓      | ✓           | ✓          | –      | –           | ○      | ✓    | ✓    |
| Cell growth, volume-scaled transcription             | –      | –      | –           | –          | –      | –           | ✓      | ✓    | ✓    |
| Gene replication / cell- cycle copy number           | –      | –      | ○           | –          | –      | –           | –      | –    | ✓    |
| Division with molecular partitioning                 | –      | –      | –           | –          | ○      | ○           | ✓      | ✓    | ✓    |
| Heritable promoter states, expression memory         | –      | –      | –           | –          | ○      | ○           | ○      | ○    | ✓    |
| Lineage tree with per-cell fates                     | –      | –      | –           | –          | ✓      | ✓           | ○      | –    | ✓    |
| Drug response: dose schedules, state-dependent death | –      | –      | –           | –          | –      | –           | –      | –    | ✓    |
| Genetic perturbations (knockdown / overexpression)   | ✓      | ✓      | ○           | ✓          | –      | –           | ○      | –    | ✓    |
| scRNA-seq observation model                          | ✓      | ✓      | ✓           | ✓          | ✓      | ✓           | –      | ○    | ✓    |
| smFISH / time-lapse observation                      | –      | –      | –           | –          | –      | –           | –      | –    | ✓    |
| Parameter inference from data                        | –      | –      | –           | –          | –      | –           | –      | –    | ✓    |
| AnnData output                                       | ✓      | ✓      | ✓           | ✓          | ○      | ✓           | –      | –    | ✓    |

### The engine is exact and agrees with an independent implementation

We first verified the kernels against exact stationary laws (Figure 2). For a birth-death process the simulated distribution of 50 000 cells is indistinguishable from the Poisson law (Figure 2a); for the telegraph model in

an intermediate and in a bursty regime it matches the Beta-Poisson distribution of Peccoud and Ycart (1995) (Figure 2b); and for bursty protein production it matches the negative binomial (Figure 2c). Across kernels, Kolmogorov-Smirnov distances to the exact law with 20 000 cells are below the 99% critical value for the direct method and for the hybrid and adaptive schemes at their default tolerances, at wall-clock costs of 0.25 s to 0.91 s per  $10^4$  cells (Figure 2d). Simulating the same telegraph model with JumpProcesses.jl gives distributions whose two-sample Kolmogorov-Smirnov statistic against Biomodelling.jl is 0.0034 (99% critical value 0.0163; Figure 2e), with a 6.2-fold (in favour of JumpProcesses.jl) difference in serial wall-clock time. Population simulations scale linearly with the number of cells and approximately linearly with the number of reactions: 2.6 s for 5000 cells  $\times$  10 telegraph genes  $\times$  200 steps (hybrid kernel) with four threads (Figure 2f).

## **Growth, division and partitioning generate expression memory, size scaling and cell-cycle effects**

The physiology of the population layer is sufficient to produce the properties of single-cell data that phenomenological simulators must impose by hand (Figure 3). In a single lineage of a slowly switching gene, mRNA and protein track volume within a cycle and are halved at division, while the promoter state, inherited by both daughters, persists over many divisions (Figure 3a). For a fast-switching gene without replication, transcript numbers scale linearly with cell volume so that concentration is homeostatic; adding mid-cycle gene replication doubles transcription in the second half of the cycle and lifts counts of large cells above the volume-scaling line (Figure 3b), the copy-number effect that mechanistic fits to scRNA-seq data must account for (Sukys and Grima 2025). Correlations of protein concentration at division between mothers and daughters, sisters and cousins decrease from 0.94 to 0.02 as the promoter switching rate crosses the inverse cell-cycle time (Figure 3c), with sisters more correlated than mother-daughter pairs and cousins less, and mRNA less heritable than protein. The memory timescale extracted from the generation-resolved lineage autocorrelation is set by the promoter switching time and is nearly independent of the cell-cycle time as long as memory outlasts the cycle; when the two are comparable the estimate is bounded below by one cell cycle, the resolution of the autocorrelation (Figure 3d). Following Zhang et al. (2025), the coefficient of variation of a stable protein’s concentration is larger across the population than along single lineages, whether or not cell-cycle times are noisy, and over-dispersed partitioning adds variability in both perspectives (Figure 3e). Finally, telegraph parameters fitted separately to cells before and after replication differ by the expected factor of two in burst frequency, and a fit that ignores the cell cycle returns intermediate values (Figure 3f).

## **A single resistance gene with slow promoter switching reproduces persister biology**

To test whether heritable expression states suffice to reproduce the observed phenomenology of drug-tolerant persisters, we coupled a resistance gene with slow promoter switching (memory timescale 45 time units, cell-cycle time 20) to a death hazard that acts on the drug dose and is suppressed by the concentration of the gene’s protein product (Methods). Under continuous, pulsed and reduced-dose schedules the population first collapses as low-expressing cells die and then regrows from the high-expressing survivors (Figure 4a); kill curves are biphasic with a surviving fraction that decreases with dose (Figure 4b). The net rate of population decline over the first 30 time units of treatment rises from 0.023 to 0.050 per time unit between doses 0.5 and 2 (2.1-fold), while the mean division time of cells born under drug changes by only 0.0% and the mean time to death of killed cells only moves from 8.5 to 6.2 time units (Figure 4c), reproducing the observation of Iyer et al. (2025) that population decay depends strongly on dose even though single-cell division and death times barely do, because fate is decided by cell states present before the drug. Accordingly, related cells born in the 25 time units before the drug share the

fate of their lineages (extinction, or survival of at least one descendant): a cell whose sister founded a surviving lineage does so itself with probability 0.39, against a marginal probability of 0.08 (4.8-fold enrichment over 1830 sister pairs; 2.9-fold for cousins), whereas in the fast-switching control, in which survival is rarer, the conditional and marginal probabilities coincide (0.03 versus 0.02; Figure 4d). When tolerance pre-exists, survivors descend from a minority of founder clones and the effective number of clones drops from 491 to 89, whereas when tolerance is drug-induced (the drug raises the promoter activation rate of every cell) barcode diversity is largely preserved (165 of 492; Figure 4e), the discriminating signature used in barcoding experiments (Shaffer et al. 2017; Iyer et al. 2025). Disrupting expression memory changes this outcome: accelerating promoter switching 20-fold from two cell-cycle times before the drug, as memory-disrupting epigenetic compounds do (Harmange et al. 2023), leaves the number of surviving clones unchanged when the acceleration stops at the moment the drug arrives (98 versus 120 clones), because it reshuffles which cells are in the resistant state without changing their number, but reduces it to 7 clones (10 cells, against 792) when the acceleration continues during exposure (Figure 4f), because resistant cells now leave the protected state within a fraction of the time they need to establish a colony: in the model, memory must be disrupted while the drug is present. A transient drug pulse enriches high-expressing cells and, with slow switching, the enrichment persists for many generations after the drug is removed, whereas with fast switching it relaxes within a few cycles (Figure 4g): the phenotypic selection with stable inheritance of the original MGMT model (Lasri et al. 2020). Finally, scanning release periods and doses of intermittent schedules (Figure 4h; net growth rate of the population over a 200-time-unit course), the most effective schedule is release period 0, dose 2 (growth rate +0.004 per time unit) for pre-existing tolerance, release period 10, dose 2 (growth rate -0.016 per time unit) when the resistant state carries a fitness cost, and release period 10, dose 1 (growth rate +0.009 per time unit) for drug-induced tolerance (the drug raising the promoter activation rate 100-fold at dose 1). Without a fitness cost, therefore, nothing beats continuous dosing at the highest dose and every release period lets the resistant subpopulation regrow. With a fitness cost of the resistant state the population declines under every schedule with dose  $\geq 0.5$ , and release periods of up to 10 time units are as effective as continuous dosing (-0.013 per time unit under continuous dose 2), because cells that leave the resistant state during the holiday are killed on re-exposure. For drug-induced tolerance the ranking of doses reverses: the highest dose induces tolerance in more cells than it kills, so dose 1 outperforms dose 2 at every release period (release period 10, dose 1 (growth rate +0.009 per time unit), against +0.012 for continuous dose 2), the intermediate-dose optimum reported by Corigliano et al. (2025); the benefit of release periods reported in that work does not appear in this model, in which the induced state decays at the same rate with and without drug, so that a holiday only postpones the death of reverting cells while all cells grow.

## Ground truth for memory genes, network inference, imputation and perturbation prediction

Because every cell's clone, lineage and true state are known, the framework supplies ground truth that observational data cannot (Figure 5). For 40 independent genes whose memory times span three decades, MemorySeq-style clonal variance scores computed on true concentrations rise from 1.0 for genes with memory times far below the cell-cycle time to 29.6 for the slowest genes, and the ranking is preserved when scores are computed on sequenced counts at 20% capture efficiency (Figure 5a); the score separates genes with memory times above the cycle time from the rest with an area under the receiver operating curve of 1.00 (true) and 0.99 (sequenced). For a 30-gene network simulated at fixed volume, in a dividing population and after sequencing, correlation-based and GENIE3 inference lose most of their accuracy in the dividing population when raw counts are used (area under the precision-recall curve 0.52 to 0.14 for Pearson correlation and 0.32 to 0.13 for GENIE3, against 0.08 for a random ranking), because volume and gene copy number induce positive correlations between all genes. Normalising

counts to concentrations removes the volume effect but recovers little (0.18 and 0.18), since the dependence of expression on copy number and cell age remains, whereas regressing expression on the true cell-cycle covariates (copy number and age, which the simulation provides and which cell-cycle scoring approximates in real data) before inference recovers most of the loss (0.46 and 0.30). Sequencing at 15% capture efficiency lowers accuracy further (0.11 for raw and 0.24 for library-size-normalised counts with Pearson correlation; 0.07 and 0.08 with GENIE3; Figure 5b). Imputation with kNN-smoothing and MAGIC before inference changes the area under the precision-recall curve by +0.05 to +0.21 (Pearson correlation) (Figure 5c), consistent with the conclusions of the original benchmark (Lasri et al. 2022). Finally, knocking down the five most connected regulators of a 20-gene network produces true fold changes that a no-change baseline and the propagation of observational correlations predict with coefficients of determination of -0.16 and -0.23 respectively (Figure 5d), illustrating how causal, out-of-distribution ground truth can be generated for the evaluation of perturbation-prediction models (Ahlmann-Eltze et al. 2025; Vollenweider and Bühlmann 2026).

## The forward model supports inference and quantifies the bias of ignoring division

On a snapshot of 2 000 non-dividing cells, ABC-SMC with log-uniform priors concentrates on the true telegraph parameters and agrees with the exact maximum-likelihood estimate (Figure 6a; posterior medians (0.47, 1.40, 26.8) for truth (0.3, 0.6, 20)). On a snapshot of a growing and dividing population with volume-scaled transcription and mid-cycle replication, the naive telegraph fit is biased ( $k_{\text{on}} = 0.47$ ,  $k_{\text{off}} = 1.95$ ,  $k_{\text{tx}} = 71.5$ ) because volume and copy number inflate counts and their dispersion, whereas ABC with the population simulator as forward model recovers the truth ((0.58, 1.87, 25.1); Figure 6b). Drug parameters are also identifiable from population data: from a kill curve and the concordance of sister fates, the death rate and the protection threshold are recovered with posterior medians [abc\_c] for truth (0.5, 150) (Figure 6c), and posterior predictive kill curves simulated at the posterior median envelope the observed one (Figure 6d).

## Discussion

Biomodelling.jl 2.0 occupies a niche that the single-cell field currently lacks: a mechanistic simulator in which expression memory, cell-size effects, cell-cycle dependence and drug-tolerant persistence are consequences of explicit biophysics rather than assumptions of a generative model. Three uses follow. First, as a hypothesis engine for non-genetic drug tolerance, it lets one ask which schedule, which combination or which memory-modulating pretreatment minimises the long-term fitness of a treated population, with cell-level rules that can be matched to time-lapse and lineage-tracing data. Second, as a ground-truth generator, it produces data with known memory genes, known networks, known perturbation effects and known lineage trees, which is what methods such as MemorySeq-style scores, Power-Seek, lineage-aware models and virtual-cell benchmarks require to be evaluated fairly, and which existing simulators cannot provide. Third, as a likelihood-free inference engine, it allows mechanistic parameters to be estimated from the kinds of data that are now routinely collected, while quantifying the bias of simpler models.

Several limitations delimit the current release. Cells are well-mixed compartments without spatial structure or cell-cell signalling; the pharmacology is a one-compartment model; growth is exponential with phenomenological size control; and the inference layer, although general, relies on user-chosen summary statistics, so posterior calibration should be checked with predictive simulations as in Figure 6. The persister case study is deliberately minimal, a single resistance gene with slow switching, and reproduces qualitative and semi-quantitative signatures reported in the literature; fitting it to a specific cell line will require the corresponding time-lapse or lineage-tracing data, for which the inference layer provides the machinery. The Catalyst.jl bridge planned for interoperability with

the SciML ecosystem is not part of this release. Finally, the benchmark results should be read as demonstrations of what the framework makes possible rather than as verdicts on the methods compared, which were run with default settings on small networks.

Future work will add spatial coupling through neighbourhood-dependent rates, cell-cell communication, chromatin states with their own inheritance rules, and neural simulation-based inference on top of the existing forward model. Because the framework is deterministic given a seed, fast enough to simulate tens of thousands of cells in minutes, and outputs AnnData files, it can be dropped into existing single-cell analysis pipelines as a source of mechanistically grounded synthetic data, and into virtual-cell training and evaluation as a source of causal, out-of-distribution contexts.

## Methods

### Reaction model and propensities

A model is a set of species and reaction channels. Species are molecules (partitioned at division) or promoter states (inherited). For a reaction with reactant counts  $x_i$ , stoichiometric coefficients  $\nu_i$ , rate constant  $k$  and volume exponent  $e$ , the mass-action propensity is

$$a = k V^e \prod_i \binom{x_i}{\nu_i},$$

where by default  $e = 1 - \sum_i \nu_i$ , so that zero-order production scales with volume, first-order reactions are volume independent and bimolecular reactions scale as  $1/V$ ; this keeps concentrations invariant under growth. Regulated reactions multiply this by  $b + (1 - b)F(\mathbf{c})$ , where  $b$  is a basal fraction and  $F$  combines Hill terms of the regulator concentrations  $c_j = x_j/V$ ,  $h_j = c_j^{n_j}/(K_j^{n_j} + c_j^{n_j})$  for activators and  $1 - h_j$  for inhibitors, either as a product (AND logic) or as  $1 - \prod_j (1 - h_j)$  (OR logic). Promoter switching is represented by explicit promoter species (for instance  $G_{\text{off}} \rightleftharpoons G_{\text{on}}$ ) grouped into promoter groups. Custom propensities can be supplied as functions of counts, volume, parameters and time.

### Stochastic simulation kernels

Within one update step the cell volume and gene copy number are held fixed and the reaction network is advanced by one of four kernels: the direct method (Gillespie 1977) with propensities updated through a species-reaction dependency graph; fixed-step Poisson tau-leaping (Gillespie 2001); a hybrid scheme in which a leap that would drive any species negative is rejected and the interval is instead simulated exactly (this is the exact version of the “switching” scheme of version 1); and adaptive tau-leaping with the step-size selection of Cao et al. (2006) and the critical-reaction partition of Cao et al. (2005), falling back to a fixed number of exact events when the proposed leap is shorter than  $10/a_0$ . Every kernel takes an explicit random number generator.

### Cells, growth, replication and division

Each cell carries its volume  $V$ , age, generation, lineage identifiers, gene copy state, growth rate, molecule counts, a cell-specific parameter vector and its own random number generator (seeded from the parent’s generator, which makes results independent of the number of threads). Volume grows exponentially,  $\dot{V} = \lambda V$ , with an optional log-normal spread of  $\lambda$  between cells. Division is triggered by a sizer ( $V \geq V_{\text{div}}$ ), an adder ( $V - V_{\text{birth}} \geq \Delta$ ) or a timer ( $\text{age} \geq T$ ), each with a per-cell noise factor drawn at birth. At division the first daughter inherits a

volume fraction  $f \sim \mathcal{N}(0.5, \sigma^2)$  (clipped to  $[0.05, 0.95]$ ) and each molecule is assigned to it independently with probability  $f$  (binomial partitioning); a beta-binomial option models over-dispersed partitioning. Promoter groups are inherited by both daughters when unreplicated and split one copy per daughter, in random assignment, when replicated. Gene replication, if enabled, doubles the promoter counts (and the propensity of reactions flagged as copy-number dependent) once the cell has completed a set fraction of its cycle, and is reset at division. Populations can be started from per-founder states and volumes (for instance a snapshot of a burnt-in constant-size population). Population control is either constant size (every second daughter replaces a uniformly chosen cell, the behaviour of version 1), free growth (all daughters kept, with uniform subsampling above a cap and book-keeping of the true size) or logistic (free growth with an additional hazard  $\lambda N/K$ ).

## Drug and perturbation layer

A perturbation combines a dose schedule  $d(t)$  (constant, pulsed, piecewise, one-compartment bolus pharmacokinetics or an arbitrary function) with effects evaluated for every cell at every step: a death hazard

$$h = h_{\max} \frac{d^m}{\text{EC}_{50}^m + d^m} \cdot \frac{K^q}{K^q + c_P^q},$$

where  $c_P$  is the concentration of a protective species (survival probability over a step is  $e^{-h \Delta t}$ ); growth inhibition by  $1/(1 + (d/\text{IC}_{50})^m)$ ; a dose-independent fitness cost of a cell state,  $1 - c_{\max} c_P^q/(K^q + c_P^q)$ , on the growth rate; multiplicative modulation of any parameter by a function of the dose; and genetic perturbations (knockdown, knockout, over-expression) of any parameter in a random fraction of cells, the flag being inherited by daughters.

## Lineage statistics

The lineage table records, for every cell, its parent, founder (clone), generation, birth and end times, fate (divided, died, removed, alive) and its volume and state at birth and at the end of its life. From it we compute Pearson correlations of expression at division between mothers and daughters, sisters and cousins; the lineage autocorrelation  $r_g$  between a cell and its ancestor  $g$  generations earlier, and the memory timescale  $\tau$  from  $r_g = e^{-g/\tau}$ ; single-lineage traces by following the ancestral line of a surviving cell; the coefficient of variation squared of concentration across the population versus along single lineages; Luria-Delbrück fluctuation tests (variance of a phenotype across clones grown from single cells compared with the binomial expectation); and MemorySeq-style clonal scores (Shaffer et al. 2020), the variance of clone-mean expression divided by its permutation-null expectation.

## Observation models

True counts are converted to observations by a scRNA-seq model in which each cell captures every molecule independently with a Beta-distributed efficiency (mean and coefficient of variation set by the protocol), optional multinomial down-sampling to a Poisson-distributed depth, optional additional dropout and batch scaling; an smFISH model with a fixed detection efficiency; and a time-lapse model that samples a species along lineages at a fixed interval with additive or multiplicative measurement noise. Because counts scale with volume, library size correlates with cell size, reproducing the transcriptome-size confound of real data. Outputs are written as CSV or as AnnData-compatible HDF5 files.

## 293 Inference

294 Snapshot counts of a gene in non-dividing cells follow the Beta-Poisson law of the telegraph model (Peccoud  
295 and Ycart 1995), which we evaluate through Kummer’s confluent hypergeometric function with rescaled series  
296 summation, and maximise by a grid-initialised Nelder-Mead search in log-parameter space. For any model,  
297 including populations, we implement sequential Monte Carlo approximate Bayesian computation (Toni et al. 2009)  
298 with an adaptive tolerance schedule (the  $\alpha$  quantile of the previous generation’s distances), Gaussian perturbation  
299 kernels of variance twice the weighted population variance, importance weights against the prior, and parallel  
300 evaluation of the user-supplied distance function.

## 301 Random regulatory networks

302 Random networks are built with a chosen number of activating and inhibiting edges under Erdős-Rényi or  
303 preferential-attachment (scale-free) sampling; regulation acts on the transcription rate of constitutive genes or on  
304 the promoter activation rate of telegraph genes through AND-combined Hill functions; per-gene rate constants  
305 can be drawn log-uniformly from ranges.


## 306 Validation protocols

307 Exactness was tested by Kolmogorov-Smirnov distances between simulated stationary distributions and the  
308 analytic laws (Poisson for birth-death, Beta-Poisson for the telegraph model, negative binomial for bursty protein  
309 production (Friedman et al. 2006; Shahrezaei and Swain 2008)), by agreement between kernels, and by a  
310 two-sample comparison against JumpProcesses.jl (Rackauckas and Nie 2017) on the same model. Population-level  
311 properties were compared with the analytic expectations for cell-size scaling and concentration homeostasis  
312 (Bertaux et al. 2018), partitioning noise (Huh and Paulsson 2011), lineage versus population statistics (Zhang  
313 et al. 2025) and cell-cycle-dependent burst parameters (Sukys and Grima 2025). Persister simulations were  
314 compared with the quantitative signatures reported by Iyer et al. (2025), Corigliano et al. (2025) and Harmange  
315 et al. (2023), and with the phenotypic-selection model of Lasri et al. (2020). All scripts, seeds and outputs are in  
316 the `paper/` directory of the repository.


## 317 Software

318 Biomodelling.jl 2.0 is written in Julia ( $\geq 1.10$ ) and depends on Distributions.jl, PoissonRandom.jl, QuadGK.jl,  
319 Optim.jl, SpecialFunctions.jl and StatsBase.jl; HDF5.jl is an optional extension. Benchmarks used absolute  
320 Pearson and Spearman correlations and scikit-learn random forests re-implementing GENIE3 (Huynh-Thu et al.  
321 2010), MAGIC (Dijk et al. 2018) and kNN-smoothing (Wagner et al. 2018). The test suite (Aqua quality checks  
322 and statistical tests) runs under GitHub Actions on Julia 1.10 and the latest release.


## Figure legends



**Figure 1. The Biomodelling.jl 2.0 framework.** (a) Each cell runs the stochastic kinetics of a reaction network (mass-action, Hill or custom propensities, explicit promoter states) at its current volume; volume grows exponentially, genes replicate at a set point of the cycle, and division follows a sizer, adder or timer rule with binomial partitioning of molecules and inheritance of promoter states. (b) Population control (constant size, free growth or logistic), the drug and perturbation layer (dose schedules, state-dependent death, growth inhibition, rate modulation, genetic perturbations) and the lineage table that records every cell's ancestry, fate and state at birth and division. (c) Observation models (scRNA-seq with capture, depth, dropout and batch effects; smFISH; time-lapse reporters) and outputs (CSV, Newick, AnnData). (d) Downstream uses: lineage statistics, fluctuation tests, clonal memory scores, likelihood-free inference and benchmark generation. (e) A complete example in ten lines of code.




**Figure 2. The engine is exact and agrees with an independent implementation.** (a) Stationary distribution of a birth-death process (bars: 50 000 simulated cells; line: Poisson). (b) Telegraph model in an intermediate and a bursty regime against the Beta-Poisson law. (c) Bursty protein production against the negative binomial. (d) Kolmogorov-Smirnov distance to the exact law for every kernel (n = 20 000; dotted line: 99% critical value) with wall-clock time per 10<sup>4</sup> cells. (e) Telegraph model simulated with Biomodelling.jl and with JumpProcesses.jl on the same model; the two-sample KS statistic is below its 99% critical value. (f) Runtime of population simulations against the number of cells (10 genes) and genes (1 000 cells) for the exact and hybrid kernels.




344

345 **Figure 3. Growth, division and partitioning generate memory, size scaling and cell-cycle effects.** (a)  
 346 A single lineage: volume, mRNA and protein of a slowly switching gene, with the promoter-on state shaded. (b)  
 347 mRNA count against cell volume for a fast-switching gene without replication (points, line: mean concentration  
 348  $\times V$ ) and with gene replication at mid-cycle (binned means). (c) Correlation of protein concentration at division  
 349 between mothers and daughters, sisters and cousins as a function of the promoter switching rate; dotted: one over  
 350 the cell-cycle time. (d) Memory timescale in time units against the measured cell-cycle time for two switching  
 351 rates; dotted lines:  $1/(k_{\text{on}} + k_{\text{off}})$ . (e) Coefficient of variation squared of a stable protein's concentration  
 352 across the population and along single lineages, for increasing partitioning noise and cell-cycle-time noise. (f)  
 353 Telegraph parameters fitted to cells before replication (one gene copy), after replication (two copies) and pooled;  
 354 dotted: per-allele values.



**Figure 4. A single resistance gene with slow promoter switching reproduces persister biology.** (a) Population size under continuous, pulsed and reduced-dose schedules (shading: drug present). (b) Kill curves at increasing dose. (c) Net population decay rate over the first 30 time units of treatment against dose, with the mean division time of cells born under drug (inset), which is unchanged. (d) Probability that a cell's lineage survives the treatment (at least one descendant alive at the end of the run) given that its sister's or cousin's lineage survived, against the marginal survival probability, for cells born in the 25 time units before the drug; memory gene and fast-switching control, with the fold enrichment above each bar. (e) Effective number of clones (barcodes) before and after drug for pre-existing and drug-induced tolerance (mean  $\pm$  s.d. of three seeds). (f) Surviving clones after drug without memory disruption, with promoter switching accelerated 20-fold from two cell-cycle times before the drug until the drug arrives, and with the acceleration maintained during exposure (mean  $\pm$  s.d. of five seeds). (g) Fraction of high-expressing cells before, during and after a drug pulse for slow and fast switching (MGMT-like phenotypic selection). (h) Net population growth rate over a 200-time-unit course as a function of release period (time units without drug between 20-time-unit exposures) and dose, for pre-existing tolerance without and with a fitness cost of the resistant state and for drug-induced tolerance; the box marks the

most effective schedule.



**Figure 5. Ground-truth benchmarks.** (a) Clonal variance score against the known memory time of 40 independent genes, from true concentrations and from sequenced counts; dotted: cell-cycle time. (b) Area under the precision-recall curve of correlation (Pearson, Spearman) and GENIE3 network inference on the same 30-gene network observed at fixed volume, in a dividing population (raw counts, concentrations, and concentrations regressed on the cell-cycle covariates copy number and age) and after sequencing (raw and library-size normalised counts); dotted: random ranking. (c) Effect of kNN-smoothing and MAGIC imputation before inference. (d) Coefficient of determination of predicted  $\log_2$  fold changes after knocking down the five most connected regulators, for a no-change baseline and for propagation of observational correlations.

**Figure 6. Inference with the forward model.** (a) Posterior distributions of telegraph parameters from ABC-SMC on a snapshot of non-dividing cells, with the truth and the exact maximum-likelihood estimate. (b) The same for a snapshot of a growing and dividing population: the naive telegraph fit is biased while division-aware ABC recovers the truth. (c) Posterior particles for the drug death rate and protection threshold inferred from a kill curve and sister fate concordance. (d) Posterior predictive kill curves simulated at the posterior median against the observed curve.

## Data and code availability

Biomodelling.jl 2.0 is available under the MIT license at <https://github.com/ayoublasri/Biomodelling.jl>. All data in this article are synthetic and are regenerated, with fixed seeds, by the scripts in the `paper/` directory

of the repository (`julia paper/run_all.jl` followed by `python paper/plot_all.py`); the generated tables are archived with the release.

## Author contributions

A.L. and M.S. conceived the project. [To be completed.]

## Competing interests

The authors declare no competing interests.

## Funding

[To be completed.]

## References

- Ahlmann-Eltze, Constantin, Wolfgang Huber, and Simon Anders. 2025. “Deep-Learning-Based Gene Perturbation Effect Prediction Does Not yet Outperform Simple Linear Baselines.” *Nature Methods* 22. <https://doi.org/10.1038/s41592-025-02772-6>.
- Bertaux, François, Samuel Marguerat, and Vahid Shahrezaei. 2018. “Division Rate, Cell Size and Proteome Allocation: Impact on Gene Expression Noise and Implications for the Dynamics of Genetic Circuits.” *Royal Society Open Science* 5 (3): 172234. <https://doi.org/10.1098/rsos.172234>.
- Cannoodt, Robrecht, Wouter Saelens, Louise Deconinck, and Yvan Saeys. 2021. “Spearheading Future Omics Analyses Using Dyngen, a Multi-Modal Simulator of Single Cells.” *Nature Communications* 12: 3942. <https://doi.org/10.1038/s41467-021-24152-2>.
- Cao, Yang, Daniel T. Gillespie, and Linda R. Petzold. 2005. “Avoiding Negative Populations in Explicit Poisson Tau-Leaping.” *The Journal of Chemical Physics* 123 (5): 054104. <https://doi.org/10.1063/1.1992473>.
- Cao, Yang, Daniel T. Gillespie, and Linda R. Petzold. 2006. “Efficient Step Size Selection for the Tau-Leaping Simulation Method.” *The Journal of Chemical Physics* 124 (4): 044109. <https://doi.org/10.1063/1.2159468>.
- Corigliano, Mattia, Alessandro Di Bernardo, Marco Cosentino Lagomarsino, and Simone Pompei. 2025. “Optimal Treatment for Drug-Induced Cancer Persists Involves Release Periods and Intermediate Drug Doses.” *PRX Life*, ahead of print. <https://doi.org/10.1103/mh5v-lcyw>.
- Dibaeinia, Payam, S. Babu, M. Knudson, et al. 2026. “Virtual Cells Need Context, Not Just Scale.” *bioRxiv*, ahead of print. <https://doi.org/10.64898/2026.02.04.703804>.
- Dibaeinia, Payam, and Saurabh Sinha. 2020. “SERGIO: A Single-Cell Expression Simulator Guided by Gene Regulatory Networks.” *Cell Systems* 11 (3): 252–71. <https://doi.org/10.1016/j.cels.2020.08.003>.

- 418 Dijk, David van, Roshan Sharma, Juozas Nainys, et al. 2018. “Recovering Gene Interactions from Single-Cell  
419 Data Using Data Diffusion.” *Cell* 174 (3): 716–29. <https://doi.org/10.1016/j.cell.2018.05.061>.
- 420 Friedman, Nir, Long Cai, and X. Sunney Xie. 2006. “Linking Stochastic Dynamics to Population Distribution:  
421 An Analytical Framework of Gene Expression.” *Physical Review Letters* 97 (16): 168302. <https://doi.org/10.1103/PhysRevLett.97.168302>.  
422
- 423 Ghosh, S., Shaon Chakrabarti, and Archishman Raju. 2025. “Identifying Memory Gene Expression from Single  
424 Sample scRNA-seq Data Using Power Law Signatures.” *bioRxiv*, ahead of print. <https://doi.org/10.1101/2025.01.15.633183>.  
425
- 426 Gillespie, Daniel T. 1977. “Exact Stochastic Simulation of Coupled Chemical Reactions.” *The Journal of Physical  
427 Chemistry* 81 (25): 2340–61. <https://doi.org/10.1021/j100540a008>.
- 428 Gillespie, Daniel T. 2001. “Approximate Accelerated Stochastic Simulation of Chemically Reacting Systems.” *The  
429 Journal of Chemical Physics* 115 (4): 1716–33. <https://doi.org/10.1063/1.1378322>.
- 430 Harmange, Guillaume, Raúl A. Reyes Hueros, Dylan L. Schaff, et al. 2023. “Disrupting Cellular Memory to  
431 Overcome Drug Resistance.” *Nature Communications* 14: 7130. <https://doi.org/10.1038/s41467-023-41811-8>.
- 432 Huh, Dann, and Johan Paulsson. 2011. “Non-Genetic Heterogeneity from Stochastic Partitioning at Cell Division.”  
433 *Nature Genetics* 43 (2): 95–100. <https://doi.org/10.1038/ng.729>.
- 434 Huynh-Thu, Vân Anh, Alexandre Irrthum, Louis Wehenkel, and Pierre Geurts. 2010. “Inferring Regulatory  
435 Networks from Expression Data Using Tree-Based Methods.” *PLoS ONE* 5 (9): e12776. <https://doi.org/10.1371/journal.pone.0012776>.  
436
- 437 Iyer, Anton, Adrian Alva, Adrián E. Granada, and Shaon Chakrabarti. 2025. “Inheritable Cell-States Shape  
438 Drug-Persister Correlations and Population Dynamics in Cancer Cells.” *PLoS Computational Biology* 21 (9):  
439 e1013446. <https://doi.org/10.1371/journal.pcbi.1013446>.
- 440 Lasri, Ayoub, Viktorija Juric, Maité Verreault, et al. 2020. “Phenotypic Selection Through Cell Death: Stochastic  
441 Modelling of O-6-Methylguanine-DNA Methyltransferase Dynamics.” *Royal Society Open Science* 7 (7): 191243.  
442 <https://doi.org/10.1098/rsos.191243>.
- 443 Lasri, Ayoub, Vahid Shahrezaei, and Marc Sturrock. 2022. “Benchmarking Imputation Methods for Network  
444 Inference Using a Novel Method of Synthetic scRNA-seq Data Generation.” *BMC Bioinformatics* 23: 236.  
445 <https://doi.org/10.1186/s12859-022-04778-9>.
- 446 Li, Hechen, Ziqi Zhang, Michael Squires, Xi Chen, and Xiuwei Zhang. 2025. “scMultiSim: Simulation of Single-Cell  
447 Multi-Omics and Spatial Data Guided by Gene Regulatory Networks and Cell–Cell Interactions.” *Nature  
448 Methods*, ahead of print. <https://doi.org/10.1038/s41592-025-02651-0>.
- 449 Lin, Yongjie, Xiangru Chen, Long Wu, Yutong Zhou, and Yihan Lin. 2026. “Widespread Transcriptional Memory

450 Shapes Heritable States and Functional Heterogeneity in Cancer and Stem Cells.” *Cell Systems* 17 (7): 101621.  
451 <https://doi.org/10.1016/j.cels.2026.101621>.

452 Pan, Xinhai, Hechen Li, and Xiuwei Zhang. 2022. “TedSim: Temporal Dynamics Simulation of Single-Cell RNA  
453 Sequencing Data and Cell Division History.” *Nucleic Acids Research* 50 (8): 4272–88. [https://doi.org/10.1093/](https://doi.org/10.1093/nar/gkac235)  
454 [nar/gkac235](https://doi.org/10.1093/nar/gkac235).

455 Peccoud, Jean, and Bernard Ycart. 1995. “Markovian Modeling of Gene-Product Synthesis.” *Theoretical Population*  
456 *Biology* 48 (2): 222–34. <https://doi.org/10.1006/tpbi.1995.1027>.

457 Piho, Paul, and Philipp Thomas. 2026. “AgentBasedModeling.jl: A Tool for Stochastic Simulation of Structured  
458 Population Dynamics.” *PLoS Computational Biology* 22 (7): e1014366. [https://doi.org/10.1371/journal.pcbi.1](https://doi.org/10.1371/journal.pcbi.1014366)  
459 [014366](https://doi.org/10.1371/journal.pcbi.1014366).

460 Rackauckas, Christopher, and Qing Nie. 2017. “DifferentialEquations.jl – a Performant and Feature-Rich  
461 Ecosystem for Solving Differential Equations in Julia.” In *Journal of Open Research Software*, No. 1, vol. 5.  
462 <https://doi.org/10.5334/jors.151>.

463 Shaffer, Sydney M., Margaret C. Dunagin, Stefan R. Torborg, et al. 2017. “Rare Cell Variability and Drug-Induced  
464 Reprogramming as a Mode of Cancer Drug Resistance.” *Nature* 546: 431–35. [https://doi.org/10.1038/nature](https://doi.org/10.1038/nature22794)  
465 [22794](https://doi.org/10.1038/nature22794).

466 Shaffer, Sydney M., Benjamin L. Emert, Raul A. Reyes Hueros, et al. 2020. “Memory Sequencing Reveals  
467 Heritable Single-Cell Gene Expression Programs Associated with Distinct Cellular Behaviors.” *Cell* 182 (4):  
468 947–59. <https://doi.org/10.1016/j.cell.2020.07.003>.

469 Shahrezaei, Vahid, and Peter S. Swain. 2008. “Analytical Distributions for Stochastic Gene Expression.”  
470 *Proceedings of the National Academy of Sciences* 105 (45): 17256–61. <https://doi.org/10.1073/pnas.0803850105>.

471 Sukys, Augustinas, and Ramon Grima. 2025. “Cell-Cycle Dependence of Bursty Gene Expression: Insights  
472 from Fitting Mechanistic Models to Single-Cell RNA-seq Data.” *Nucleic Acids Research* 53 (7): gkaf295.  
473 <https://doi.org/10.1093/nar/gkaf295>.

474 Toni, Tina, David Welch, Natalja Strelkowa, Andreas Ipsen, and Michael P. H. Stumpf. 2009. “Approximate  
475 Bayesian Computation Scheme for Parameter Inference and Model Selection in Dynamical Systems.” *Journal*  
476 *of the Royal Society Interface* 6 (31): 187–202. <https://doi.org/10.1098/rsif.2008.0172>.

477 Vollenweider, Michael Samir, and Peter Bühlmann. 2026. “Signal, Bounds, and Baselines: Principles for  
478 Rigorous Evaluation of High-Dimensional Biological Perturbation Prediction.” *bioRxiv*, ahead of print. <https://doi.org/10.64898/2026.04.20.719650>.  
479 <https://doi.org/10.64898/2026.04.20.719650>.

480 Wagner, Florian, Yun Yan, and Itai Yanai. 2018. “K-Nearest Neighbor Smoothing for High-Throughput Single-Cell  
481 RNA-Seq Data.” *bioRxiv*, ahead of print. <https://doi.org/10.1101/217737>.

- 482 Zhang, Zhanhao, Iryna Zabaikina, César Nieto, Zahra Vahdat, Pavol Bokes, and Abhyudai Singh. 2025. “Stochastic  
483 Gene Expression in Proliferating Cells: Differing Noise Intensity in Single-Cell and Population Perspectives.”  
484 *PLoS Computational Biology* 21: e1013014. <https://doi.org/10.1371/journal.pcbi.1013014>.
- 485 Zinati, Yazdan, Abdulrahman Takiddeen, and Amin Emad. 2024. “GRouNdGAN: GRN-Guided Simulation of  
486 Single-Cell RNA-seq Data Using Causal Generative Adversarial Networks.” *Nature Communications* 15: 4055.  
487 <https://doi.org/10.1038/s41467-024-48516-6>.