

Mechanistic simulation of heritable expression states, cell division and drug response in single-cell populations

Ayoub Lasri¹

¹ Department of Physiology and Medical Physics, Royal College of Surgeons in Ireland, Dublin, Ireland. Correspondence: alasri@berache.com

Abstract

Heritable differences in gene expression decide which cancer cells survive a drug, yet the simulators used to benchmark single-cell methods model neither growth nor division, and cannot generate such states. Here we present Biomodelling.jl 2.0, in which stochastic reaction kinetics run inside growing, dividing and drug-treated cells. Exact, hybrid and adaptive kernels are coupled to exponential growth, volume-scaled transcription, gene replication and division with binomial partitioning, to dose schedules and state-dependent killing, and to a lineage record, observation models and likelihood-free inference. Expression memory, cell-size scaling and cell-cycle-dependent bursting emerge from this physiology rather than being imposed. A single resistance gene with slow promoter switching reproduces the lineage and memory signatures of heritable drug tolerance, and the framework supplies ground truth that benchmarks of memory genes, network inference and perturbation prediction lack. Calibrated to time-lapse measurements of cisplatin-treated cells, it reproduces lineage correlations and the dose-invariance of single-cell timing, and its residual errors expose the limits of one switching gene. With a built-in optimiser it reproduces the opposite rankings of intermittent dosing reported in a melanoma trial and in xenografts, and shows the ranking to turn on two measurable cell properties, their growth under drug and their fitness without it.

Isogenic cells differ in their transcriptional states, those differences are often inherited across several divisions, and the inherited state biases what a cell does under stress. Cancer cells that survive targeted or cytotoxic therapy frequently do so without a resistance mutation: they occupy a rare, reversible, drug-tolerant expression state that pre-exists the treatment or is induced by it^{1–5}. Lineage-resolved experiments now measure this memory directly. The fates of sisters and cousins are correlated, clonal barcode diversity is preserved or collapsed depending on whether tolerance pre-exists or is induced, and the timescale of expression fluctuations relative to the cell cycle decides which genes carry memory at all^{6–9}. In parallel, mechanistic analyses of single-cell RNA-sequencing data have shown that transcriptional kinetics can be inferred correctly only when gene copy number, cell age and partitioning at division are modelled explicitly^{10–12}, and evaluations of “virtual cell” models for perturbation prediction keep finding that deep-learning approaches fail to beat simple baselines, largely for want of mechanism and causal structure^{13–16}.

These developments create a demand that existing simulators do not meet. The simulators used to benchmark single-cell methods, among them SERGIO, dyngen, scMultiSim and GRouNdGAN^{17–20}, model regulation and technical noise but not cell growth, volume-dependent transcription, gene replication or the partitioning of molecules at division; lineage simulators such as TedSim²¹ generate division trees but attach phenomenological, largely Markovian expression to them. None of them can therefore produce expression memory, cell-size scaling, cell-cycle-dependent bursting or persister dynamics from first principles, and none provides the ground truth that

the methods now being developed for exactly those phenomena require²². Stochastic simulators of growing and dividing cells exist^{23,24} but stop at the single-cell state: they offer no drug response, no lineage statistics, no observation models, no inference and no benchmark generation.

Here we present `Biomodelling.jl` 2.0, a rewrite of a Julia framework whose distinguishing feature, the coupling of stochastic reaction kinetics to exponential growth, volume-scaled transcription and binomial partitioning at division²⁵, is precisely the mechanism behind heritable expression states. Version 2.0 turns that feature into a complete framework for the mechanistic study of drug tolerance and for single-cell benchmarking. We validate the engine against exact stationary laws and an independent implementation, show that growth, division and partitioning alone reproduce published results on cell-size scaling, partitioning noise, lineage-versus-population statistics and cell-cycle-dependent bursting, and demonstrate that a single resistance gene with slow promoter switching reproduces the quantitative lineage and memory signatures of heritable drug tolerance. We then use the framework as a ground-truth generator for memory-gene detection, network inference, imputation and perturbation prediction; as an inference engine that recovers kinetic parameters from snapshot and lineage data, quantifies the bias of ignoring division, and maps the ridge to which drug parameters are determined by population data; as a model calibrated to published time-lapse measurements and confronted with conditions it was not fitted to; and, with a built-in schedule optimiser, as a testbed in which clinical dosing regimens are compared and better ones are sought.

Results

Layers couple reaction kinetics to growth and division

`Biomodelling.jl` 2.0 is organised in layers (Fig. 1). A model is a list of reactions with mass-action, Hill or custom kinetics and named parameters; promoter states are explicit species, grouped so that they are inherited at division and doubled at replication. Four kernels advance a cell's state at fixed volume: the exact direct method, fixed-step tau-leaping, a hybrid scheme that falls back to the exact method whenever a leap would produce a negative count, and adaptive tau-leaping with critical-reaction handling. The population layer grows each cell's volume exponentially, scales propensities with volume so that concentrations are invariant, replicates genes at a set point of the cycle, divides cells by a sizer, adder or timer rule with binomial or over-dispersed partitioning, and controls the population size by Moran-like replacement, free growth or a logistic hazard. A perturbation layer applies dose schedules, state-dependent death, growth inhibition, drug-induced rate changes and genetic perturbations. Every cell's ancestry, fate and state at birth and division enter a lineage table, from which heritability, autocorrelation, memory timescales, fluctuation-test and clonal-memory statistics follow. Observation models convert true counts into scRNA-seq, smFISH or time-lapse measurements, and the same forward model drives likelihood-free inference. Table 1 sets these capabilities against those of the simulators most used to benchmark single-cell methods.

Table 1. Capabilities of single-cell simulators (✓ native; ○ partial or phenomenological; – absent). SERGIO¹⁷, dyngen¹⁸, scMultiSim¹⁹, GRouNdGAN²⁰, TedSim²¹, Cassiopeia²⁶, ABM.jl (AgentBasedModeling.jl)²⁴, Biomodelling.jl v1.x²⁵ and v2.0 (this work). Cassiopeia simulates division topologies with heritable fitness and CRISPR barcodes rather than molecular content, and its inference recovers phylogenies rather than kinetic parameters; entries were read from each tool’s publication and documentation.

Capability	SERGIO	dyngen	scMulti-Sim	GRouNd-GAN	TedSim	Cassio-peia	ABM.jl	v1.x	v2.0
Mechanistic reaction kinetics (SSA / tau-leaping)	○	✓	○	–	–	–	✓	✓	✓
Gene regulatory networks with known ground truth	✓	✓	✓	✓	–	–	○	✓	✓
Cell growth, volume-scaled transcription	–	–	–	–	–	–	✓	✓	✓
Gene replication / cell- cycle copy number	–	–	–	–	–	–	–	–	✓
Division with molecular partitioning	–	–	–	–	○	○	✓	✓	✓
Heritable promoter states, expression memory	–	–	–	–	○	–	○	○	✓
Lineage tree with per-cell fates	–	–	–	–	✓	✓	○	–	✓
Drug response: dose schedules, state-dependent death	–	–	–	–	–	–	–	–	✓
Genetic perturbations (knockdown / overexpression)	✓	✓	○	✓	–	–	○	–	✓
scRNA-seq observation model	✓	✓	✓	✓	✓	–	–	○	✓
smFISH / time-lapse observation	–	–	–	–	–	–	–	–	✓
Parameter inference from data	–	–	–	–	–	○	–	–	✓
AnnData output	✓	✓	✓	✓	○	○	–	–	✓

The engine matches exact laws and an independent package

We first verified the kernels against exact stationary laws (Fig. 2). For a birth-death process the simulated distribution of 50,000 cells is indistinguishable from the Poisson law (Fig. 2a); for the telegraph model in an intermediate and in a bursty regime it matches the Beta-Poisson distribution²⁷ (Fig. 2b); and for bursty protein production it matches the negative binomial (Fig. 2c). Across kernels, Kolmogorov-Smirnov distances to the exact law with 20,000 cells lie below the 99% critical value for the direct method, for the adaptive scheme at both tolerances and for the hybrid scheme at the shorter step, and above it for the hybrid scheme at the longer step (0.0160 and 0.0122 against a critical value of 0.0115; Fig. 2d). That is the expected behaviour of an approximate kernel, and it sets the step below which the approximation is not detectable at this sample size; wall-clock costs are 0.25 s to 0.91 s per 10^4 cells. Simulating the same telegraph model with JumpProcesses.jl gives distributions whose two-sample Kolmogorov-Smirnov statistic against Biomodelling.jl is 0.0034 (99% critical value 0.0163; Fig. 2e), with a 6.2-fold (in favour of JumpProcesses.jl) difference in serial wall-clock time. Population simulations scale linearly with the number of cells and approximately linearly with the number of reactions: 2.6 s for 5000 cells $\times$ 10 telegraph genes $\times$ 200 steps (hybrid kernel) with four threads (Fig. 2f and Supplementary Table 3).

Physiology generates memory, scaling and cycle effects

The physiology of the population layer suffices to produce properties of single-cell data that phenomenological simulators must impose by hand (Fig. 3). In a single lineage of a slowly switching gene, mRNA and protein track

volume within a cycle and are halved at division, while the promoter state, inherited by both daughters, persists over many divisions (Fig. 3a). For a fast-switching gene without replication, transcript numbers scale linearly with cell volume, so that concentration is homeostatic; adding mid-cycle gene replication doubles transcription in the second half of the cycle and lifts counts of large cells above the volume-scaling line (Fig. 3b), the copy-number effect that mechanistic fits to scRNA-seq data must account for¹⁰. Correlations of protein concentration at division between mothers and daughters, sisters and cousins fall from 0.94 to 0.02 as the promoter switching rate crosses the inverse cell-cycle time (Fig. 3c), mother-daughter and sister pairs correlating similarly at slow switching and mother-daughter pairs more strongly as switching quickens, cousins least, and mRNA less heritable than protein. The memory timescale extracted from the generation-resolved lineage autocorrelation is set by the promoter switching time and is nearly independent of the cell-cycle time as long as memory outlasts the cycle; when the two are comparable the estimate is bounded below by one cell cycle, the resolution of the autocorrelation (Fig. 3d). As reported for real lineages¹¹, the coefficient of variation of a stable protein's concentration is larger across the population than along single lineages, whether or not cell-cycle times are noisy, and over-dispersed partitioning adds variability in both perspectives (Fig. 3e). Telegraph parameters fitted separately to cells before and after replication differ by the expected factor of two in burst frequency, and a fit that ignores the cell cycle returns intermediate values (Fig. 3f).

One resistance gene reproduces tolerance signatures

To test whether heritable expression states suffice to reproduce the lineage and memory signatures of drug-tolerant persistence, we coupled a resistance gene with slow promoter switching (memory timescale 45 time units, cell-cycle time 20) to a death hazard that rises with dose and is suppressed by the concentration of the gene's protein product (Methods). Under continuous, pulsed and reduced-dose schedules the population first collapses as low-expressing cells die, then regrows from high-expressing survivors (Fig. 4a); kill curves are biphasic with a surviving fraction that falls with dose (Fig. 4b). Between doses 0.5 and 2 the net rate of population decline over the first 30 time units of treatment rises from 0.023 to 0.050 per time unit (2.1-fold), while the mean division time of cells born under drug changes by 0.0% and the mean time to death of killed cells moves only from 8.5 to 6.2 time units (Fig. 4c; the single-cell timing distributions are in Supplementary Fig. 1). This reproduces the observation that population decay depends strongly on dose even though single-cell division and death times barely do⁶, because fate is biased by states that exist before the drug arrives. Unlike the growth-arrested persisters of the canonical definition¹, the survivors here divide at an essentially unchanged rate (Fig. 4c, inset), so the model reproduces the heritable-state selection signatures of tolerance rather than its proliferative-arrest hallmark.

Related cells accordingly share the fate of their lineages. Among cells born in the 25 time units before the drug, one whose sister founded a surviving lineage does so itself with probability 0.39, against a marginal probability of 0.08 (4.8-fold enrichment over 1830 sister pairs; 2.9-fold for cousins), whereas in a fast-switching control, in which survival is rarer, the conditional and marginal probabilities coincide (0.03 versus 0.02; Fig. 4d). When tolerance pre-exists, survivors descend from a minority of founder clones and the effective number of clones drops from 491 to 89; when tolerance is drug-induced, so that the drug raises the promoter activation rate of every cell, barcode diversity is largely preserved (165 of 492; Fig. 4e). This is the signature that distinguishes the two origins in barcoding experiments^{2,6}.

Disrupting expression memory changes the outcome, and when it is disrupted matters. Accelerating promoter switching 20-fold from two cell-cycle times before the drug, as memory-disrupting epigenetic compounds do⁹, does not reduce the number of surviving clones if the acceleration stops when the drug arrives: it raises them, from 98 to 120 (mean of 5 seeds each, $P = 0.009$, Welch's t test). Reshuffling spreads the resistant state over

134 more lineages and moves the resistant fraction towards its stationary value (0.09 to 0.11), so more clones carry a
 135 protected cell when the drug arrives, even though protection itself is untouched. Maintaining it during exposure
 136 reduces the count to 7 clones (10 cells, against 792; Fig. 4f), because resistant cells now leave the protected state
 137 within a fraction of the time they need to found a colony. In the model, memory must be disrupted while the
 138 drug is present. Compartment theory for phenotypic switching already predicts that an agent which perturbs the
 139 retention of the resistant state can turn treatment failure into success²⁸; what the simulations add is the timing
 140 requirement and a lineage-resolved readout, the number of surviving clones, by which it can be tested. A transient
 141 pulse enriches high-expressing cells and, with slow switching, the enrichment persists for many generations after
 142 the drug is removed, whereas with fast switching it relaxes within a few cycles (Fig. 4g): phenotypic selection with
 143 stable inheritance, as in the original MGMT model²⁹. Finally, scanning release periods and doses of intermittent
 144 schedules (Fig. 4h), the most effective schedule is release period 0, dose 2 (growth rate +0.004 per time unit)
 145 for pre-existing tolerance, release period 10, dose 2 (growth rate -0.016 per time unit) when the resistant state
 146 carries a fitness cost, and release period 10, dose 1 (growth rate +0.009 per time unit) for drug-induced tolerance.
 147 Without a fitness cost nothing beats continuous dosing at the highest dose, since every release period lets the
 148 resistant subpopulation regrow (the drug does not act on the switching rates in this model; when it also stabilises
 149 the resistant state, as in the melanoma case study below, holidays regain a small advantage even without a cost).
 150 With a fitness cost the population declines under every schedule at dose ≥ 0.5 , and release periods of up to 10
 151 time units match continuous dosing (-0.013 per time unit under continuous dose 2), because cells that leave
 152 the resistant state during a holiday are killed on re-exposure. For drug-induced tolerance the ranking of doses
 153 reverses: the highest dose induces tolerance in more cells than it kills, so dose 1 outperforms dose 2 at every
 154 release period (release period 10, dose 1 (growth rate +0.009 per time unit), against +0.012 for continuous dose
 155 2), the intermediate-dose optimum reported previously³⁰; the benefit of release periods reported in that work does
 156 not appear here, because in this model the induced state decays at the same rate with and without drug, so a
 157 holiday only postpones the death of reverting cells while all cells grow.

158 **Simulated ground truth benchmarks inference methods**

159 Because every cell's clone, lineage and true state are known, the framework supplies ground truth that observational
 160 data cannot (Fig. 5). For 40 independent genes whose memory times span three decades, MemorySeq-style clonal
 161 variance scores computed on true concentrations rise from 1.0 for genes with memory far below the cell-cycle
 162 time to 29.6 for the slowest genes, and the ranking survives computation on sequenced counts at 20% capture
 163 efficiency (Fig. 5a); the score separates genes with memory longer than the cycle from the rest with an area under
 164 the receiver operating characteristic (ROC) curve of 1.00 on true concentrations and 0.99 on sequenced counts.
 165 For a 30-gene network simulated at fixed volume, in a dividing population and after sequencing, correlation-based
 166 and GENIE3 inference lose most of their accuracy in the dividing population when raw counts are used (area
 167 under the precision-recall curve 0.52 to 0.14 for Pearson correlation and 0.32 to 0.13 for GENIE3, against 0.10
 168 for a random ranking of undirected pairs and 0.05 for directed edges), because volume and gene copy number
 169 induce positive correlations between all genes. Normalising counts to concentrations removes the volume effect but
 170 recovers little (0.18 and 0.18), since the dependence on copy number and cell age remains; regressing expression
 171 on the true cell-cycle covariates, which the simulation provides and which cell-cycle scoring approximates in real
 172 data, recovers most of the loss (0.46 and 0.30). Sequencing at 15% capture efficiency lowers accuracy further (0.11
 173 raw and 0.24 library-size-normalised for Pearson correlation; 0.07 and 0.08 for GENIE3; Fig. 5b). Imputing the
 174 sequenced counts before inference recovers part of that loss: MAGIC raises the area under the precision-recall
 175 curve from 0.11 to 0.32 for Pearson correlation and from 0.07 to 0.18 for GENIE3, and kNN-smoothing gives 0.16
 176 and 0.12 (Fig. 5c). Neither reaches the accuracy obtained at fixed volume, nor that of regressing out the cell-cycle

covariates, so imputation mitigates the sequencing loss without touching the growth confound that produces most of it²⁵. Knocking down the five most connected regulators of a 20-gene network gives true fold changes that a no-change baseline predicts with a mean coefficient of determination of -0.16. Propagating observational correlations beats that baseline for 4 of 5 knockdowns, with coefficients of determination from 0.09 to 0.43, and fails on the remaining one (-2.39), so its mean over the five is -0.23 against a median of 0.30 (Fig. 5d). A predictor that is useful for most perturbations and badly wrong for one is precisely what an out-of-distribution evaluation has to expose, and exposing it needs the causal ground truth that the simulation supplies^{13,14}.

Inference recovers parameters and exposes division bias

On a snapshot of 2,000 non-dividing cells, ABC-SMC with log-uniform priors over three decades and the Wasserstein distance between simulated and observed counts³¹ concentrates on the true telegraph parameters and agrees with the exact maximum-likelihood estimate (posterior medians (0.31, 0.60, 20.9) for truth 0.3, 0.6, 20; Fig. 6a), the broadest posterior falling on the promoter inactivation rate, which a single snapshot constrains least. On a snapshot of a growing and dividing population with volume-scaled transcription and mid-cycle replication, the naive telegraph fit is biased ($k_{\text{on}} = 0.47$, $k_{\text{off}} = 1.95$, $k_{\text{tx}} = 71.5$) because volume and copy number inflate counts and their dispersion, whereas ABC with the population simulator as forward model recovers the truth ((0.30, 0.68, 20.8); Fig. 6b). Drug parameters follow from population data in the same way: from a kill curve and the concordance of sister fates, the posterior over the maximal death rate and the protection threshold is a ridge (weighted correlation of the two log parameters -0.85) that contains the truth (Fig. 6c), and posterior predictive kill curves simulated at the posterior median reproduce the observed one (Fig. 6d). The two summaries constrain the product of death rate and unprotected fraction, so separating the parameters requires measuring the protective protein: an experimental-design question the forward model answers before an experiment is run (the tolerance and acceptance-rate schedules of all three ABC runs are in Supplementary Fig. 2).

Calibration to time-lapse data is tested on held-out data

A simulator is useful for treatment questions only if it can be fitted to the data a laboratory produces and then predict what it was not fitted to. We therefore calibrated the resistance-gene model to time-lapse measurements in which U2OS cells were tracked for three days under 7, 10 and 13 μM cisplatin and every cell present at drug addition was scored as dying, dividing or surviving without dividing⁶ (Supplementary Table 4). Six parameters, the two promoter switching rates, the maximal death rate, the EC_{50} and Hill coefficient of death and the IC_{50} of growth arrest, were fitted to the fate fractions at 7 and 13 μM by simulated moments with common random numbers, under a weak prior keeping the memory of the resistant state near the two to three generations that the same study inferred from lineage correlations (Methods). The calibrated model reproduces the training fractions with a root-mean-square error of 0.071 (memory 2.8 generations, resistant fraction 15.2%, EC_{50} 16.7 μM ; Fig. 7a,f). That error should be read against the precision of the measurement: the observed fractions carry a mean binomial standard error of 0.028 and repeated simulations differ by 0.015, so the fit is about 2.2 times the combined noise and the model does not reproduce the training data within error. At the held-out concentration the root-mean-square error is 0.093, and the discrepancy is specific rather than diffuse: the model matches the fraction surviving without dividing but over-predicts death and under-predicts division (observed died, divided and survived 0.28, 0.31 and 0.41; predicted 0.39, 0.20 and 0.41; Fig. 7a), so it captures the ordering of the three concentrations while misplacing the split between death and division at the intermediate one.

Six parameters fitted to six numbers are not all determined by them. Moving one parameter at a time from the fitted values, with the others held there, the fate fractions constrain the death parameters to within about

a factor of 2 but leave the two switching rates free over roughly a factor of 6 (Supplementary Fig. 3a), which is why the memory of the resistant state was given a prior rather than left to the fate data. Holding the death parameters at their fitted values, memories from 1.0 to 9.0 generations paired with resistant fractions from 1% to 25% all describe the fate fractions within the noise of the measurement (Supplementary Fig. 3b), because a shorter memory with more resistant cells and a longer memory with fewer are hard to tell apart from fates alone. The parameters quoted above should therefore be read as one member of a family of fits that describe these fate fractions equally well, and the biological reading of the memory rests on the lineage correlations that set the prior, not on the fate fractions. Halving the integration step changes the simulated fractions by 0.005, less than the seed-to-seed spread, so the discretisation is not a factor (Supplementary Fig. 3c).

A second mechanism can produce the same fates. Cisplatin kills preferentially during replication, so where a cell sits in its division cycle when the drug arrives is a competing explanation both for the dose-invariance of death timing and for the cells that survive three days without dividing: a hazard confined to part of the cycle spares whatever never enters it, with no heritable state at all. We therefore refitted the model with a cell-cycle-dependent hazard, adding one parameter, the fraction of the hazard that is independent of cycle position, with killing otherwise peaking at mid-cycle, under the same objective, bounds, budget and seeds (Methods). The two fits describe the training fractions equally well (0.064 against 0.061, against a mean sampling error of 0.028 in the measurement), and the cycle-dependent fit does not improve the held-out prediction either (0.106 against 0.095). The fitted cycle-independent fraction is 0.34, but values from 0.20 to 0.80 all sit within one sampling error of the best fit, so the fate fractions do not determine it (Supplementary Note 12 and Supplementary Fig. 4). Three fate counts per concentration cannot separate inherited protection from cell-cycle position; the measurement that can is the cycle position of each cell at the moment the drug arrives, which a cycle reporter records directly.

Properties that were not fitted can then be compared with the data. The predicted kill curves (Fig. 7b) are 1.5-fold steeper at 13 μ M in their first two days than in their second half, but they do not plateau, as HCT116 cells do after about 100 h. Death fates are correlated between sisters (0.09), first cousins (0.06) and second cousins (0.05), and the correlation decays to 0.02 for third cousins and 0.00 for unrelated cells (Fig. 7c), the pattern reported for HCT116 lineages. The mean time to death changes by 8% between 7 and 13 μ M while the death fraction rises from 0.13 to 0.59 (Fig. 7d), the dose-invariance of single-cell timing that motivated the original study, and it arises here because dose changes the fraction of unprotected cells rather than their kinetics. The plateau the kill curves miss is the clearest failure of the minimal model: reproducing it requires cells to stay in the resistant state for longer than the memory of 2.8 generations allows, and that memory was set by the prior taken from the lineage correlations rather than by the kill curves. The two observations therefore pull in opposite directions within this minimal model: correlations that vanish by the third cousin call for a short memory, and a kill curve that flattens calls for a long one. A single switching gene cannot satisfy both, which is a concrete and testable limitation, and points to a drug-induced stabilisation of the resistant state, for which there is independent evidence², as the mechanism a lineage-tracing experiment should look for. The calibration landscape (Fig. 7e) shows how the fitted death rate and EC_{50} are constrained by the fate fractions alone; the protocol applies to any observable a laboratory records.

Schedule optimisation is consistent with clinical outcomes

With a calibrated model the framework becomes a schedule testbed. We first asked which schedule best delays regrowth of a melanoma-like population with a rare, heritable pre-resistant state² treated at a fractional-killing dose, comparing continuous dosing, the intermittent schedule of the SWOG S1320 trial, in which patients received a continuous 8-week lead-in and were then randomised to continuous dosing or to cycles of 3 weeks off and 5

260 weeks on³², and adaptive therapy, in which the drug is withdrawn when the population has halved and resumed
 261 when it regrows to its initial size³³. Because adaptive therapy deliberately lets the tumour regrow, schedules are
 262 compared by the time after randomisation at which the population exceeds 120% of its pre-treatment size (loss of
 263 control), with the RECIST-like time to progression from the nadir reported alongside (Methods).

264 The answer depends on two properties of the resistant cells (Fig. 8a–d). When the resistant state is fully
 265 protective and carries no fitness cost, as lineage data indicate for pre-existing tolerance⁶, resistant cells grow at
 266 the same rate with or without drug, holidays only let sensitive cells regrow and be killed again, and the three
 267 schedules lose control after 23, 26 and 27 weeks, the intermittent schedule delivering 65% of the cumulative dose.
 268 When resistant cells pay a 50% fitness cost in the absence of drug, a coarse stand-in for the drug dependence
 269 that made intermittent dosing superior in vemurafenib-treated patient-derived xenografts, where withdrawal
 270 drove established resistant tumours to regress rather than merely grow more slowly³⁴, holidays let sensitive cells
 271 outcompete them: none of the three schedules loses control within the 52-week follow-up (52+, 52+ and 52+
 272 weeks, censored), and adaptive therapy achieves this with 36% of the cumulative dose of continuous dosing. When
 273 resistant cells are only partly protected and grow at half speed under drug, holidays release them: continuous
 274 dosing keeps control for 52+ weeks against 49 for the intermittent schedule and 43 for adaptive therapy, the
 275 ranking observed in the trial, where continuous dosing gave a median progression-free survival of 9.0 months
 276 against 5.5 months for intermittent dosing. That the benefit of holidays should depend on a cost of resistance
 277 is established by compartment models of adaptive therapy, which also show that tumour turnover modulates
 278 how large that cost must be³⁵. Here the dependence emerges from single cells whose resistance is a heritable
 279 expression state rather than a fixed compartment, the cost and the degree of protection are properties of cells
 280 that lineage experiments measure, and the third case shows that partial protection reverses the ranking even
 281 when a cost is absent. The optimiser makes the dependence explicit (Fig. 8e): over intermittent schedules of
 282 any period and duty cycle after the lead-in, the best schedule is a period of 2.2 weeks with 39% of the time on
 283 drug (27 weeks to loss of control) without a fitness cost, a period of 1.0 weeks with 47% of the time on drug (52
 284 weeks to loss of control) with a fitness cost and a period of 1.0 weeks with 100% of the time on drug (52 weeks
 285 to loss of control) with partial protection. The model therefore reproduces both the clinical and the preclinical
 286 outcome and shows their difference to depend on measurable properties of the resistant cells, their growth under
 287 drug and their fitness without it, both properties of cells rather than of the tumour, which growth and lineage
 288 experiments can bound, within the identifiability limits of Fig. 6c, before a schedule is chosen. The mechanism
 289 is set in each case, not inferred from the outcome, so this is a demonstration that the ranking follows from cell
 290 properties rather than a claim that these particular cells carry them; and because the trial and the xenograft
 291 also differ in drug and species (dabrafenib plus trametinib in patients, vemurafenib monotherapy in mice), their
 292 contrast is not attributable to resistant-cell fitness alone.

293 The second case study concerns glioblastoma, where the RTOG 0525 trial compared the standard temozolomide
 294 regimen (150–200 mg/m² on days 1–5 of 28-day cycles) with a dose-dense regimen (75–100 mg/m² on days 1–21)
 295 that delivers about twice the cumulative dose and was designed to deplete MGMT, and found no difference
 296 in survival for either MGMT-methylated or unmethylated tumours³⁶. We simulated both regimens with one-
 297 compartment pharmacokinetics (elimination half-life 2.1 h; ref.³⁷) on populations in which 1% (methylated) or
 298 30% (unmethylated) of cells express MGMT, with MGMT protein protecting cells from the drug²⁹ and either
 299 stable or consumed by it stoichiometrically, one molecule per lesion repaired, as a suicide enzyme is, so that
 300 the pool follows the cumulative exposure rather than the peak concentration (Fig. 8f,g). With stable MGMT,
 301 the depth of response over six cycles is 0.20 log₁₀ units for the standard regimen and 1.00 for the dose-dense
 302 regimen in the methylated population (0.08 and 0.10 in the unmethylated one), and a 21-day regimen at the
 303 standard cumulative dose achieves no measurable kill (0.00 log₁₀ units): the death hazard rises steeply with

concentration, so spreading the same dose over more days loses efficacy, and the advantage of the dose-dense arm in the model is entirely its larger cumulative dose. With stable MGMT the two arms therefore differ little in the unmethylated population (0.08 against 0.10 $\log_{10}$ units), as in the trial. Consumption changes that. Because one molecule is spent per lesion repaired, the pool follows the cumulative exposure, which the dose-dense regimen more than doubles: in the methylated population its response deepens to 1.52 $\log_{10}$ units against 0.22 for the standard regimen, and in the unmethylated population, where 30% of cells start protected, the same depletion turns a regimen that barely worked into one that controls the population (1.08 against 0.18 $\log_{10}$ units; final size 0.1 against 1.1 times the initial size, and 4.9 and 3.6 with stable MGMT). Modelled as the suicide enzyme it is, MGMT therefore predicts a clear dose-dense advantage in both populations, which RTOG 0525 did not find. Read as a bound rather than as a fit, that disagreement is informative: either temozolomide consumes MGMT in tumours far more slowly than this stoichiometry allows, or the benefit of depleting it is offset by something the model leaves out. At the standard cumulative dose the number of dosing days per cycle that minimises the final population is 9 (methylated) and 9 (unmethylated) with stable MGMT, and 9 and 5 when MGMT is consumed (Fig. 8h). These simulations are not fits to the trial, whose end points depend on much that the model omits, dose-limiting toxicity and the dose reductions it forces among it. What they add is a sharper version of the rationale for dose-dense dosing: the benefit it is supposed to deliver exists in the model only at consumption rates that the trial’s null result appears to exclude, so the rate at which temozolomide consumes MGMT in the tumour is both the quantity the argument rests on and the one that has not been measured.

Discussion

Biomodelling.jl 2.0 fills a gap in the single-cell toolkit: a simulator in which expression memory, cell-size effects, cell-cycle dependence and drug-tolerant persistence follow from explicit biophysics instead of being assumptions of a generative model. It is worth separating what is new here from what it confirms. The individual ingredients are not new: stochastic kinetics in growing and dividing cells have been simulated before^{23–25}, lineage-tracing data have their own simulators^{21,26}, compartment models of adaptive therapy established that the benefit of treatment holidays depends on a cost of resistance and on turnover^{33,35}, and phenotypic-switching theory established that perturbing the retention of a resistant state can decide a treatment outcome²⁸. Three things are new. First, the combination itself: molecular kinetics, volume-scaled transcription, gene replication, partitioning and promoter inheritance, drug schedules with state-dependent killing, a complete lineage record, observation models and inference in one forward model, which is what allows a resistance phenotype to be a heritable expression state with measurable kinetics rather than a compartment with an assumed rate. Second, the ground truth this makes available: the same simulation supplies known memory genes, known networks, known perturbation effects and the true lineage tree, which is what the quantification of how growth and sequencing destroy network inference, and of how far cell-cycle regression recovers it, depends on. Third, the calibrate-then-predict protocol applied to schedule questions, which turns a mechanism into a falsifiable prediction: here it identifies the growth of resistant cells without drug and their residual growth under drug as the two properties that decide whether holidays help, and points to the rate at which temozolomide consumes MGMT as the measurement that would separate the two glioblastoma regimens. That difference has practical consequences. As a hypothesis engine for non-genetic drug tolerance, the framework answers which schedule, which combination or which memory-modulating pretreatment minimises the long-term fitness of a treated population, with cell-level rules that can be matched to time-lapse and lineage-tracing data. As a ground-truth generator, it produces data with known memory genes, known networks, known perturbation effects and known lineage trees, which is what MemorySeq-style scores, lineage-aware models and virtual-cell benchmarks need in order to be evaluated fairly, and what existing simulators cannot supply. As

346 a likelihood-free inference engine, it estimates mechanistic parameters from data that are now routinely collected,
347 while quantifying the bias of simpler models.

348 The calibration and schedule studies set out a workflow for in silico single-cell pharmacology: declare the training
349 and validation observables before fitting, fit by simulated moments or approximate Bayesian computation with
350 common random numbers, validate on conditions and statistics that were not fitted, and only then compare
351 regimens, reporting outcomes that map onto clinical end points. The case studies show why mechanism matters.
352 The same question, continuous or intermittent dosing, has opposite answers depending on two properties of the
353 resistant cells, their growth under drug and their fitness without it, and a trial result and a xenograft result that
354 appeared to contradict each other are both reproduced once those properties are set — set in each case rather
355 than inferred from the outcomes, so the case studies show that the ranking follows from cell properties, not that
356 any particular tumour has them. Because the framework simulates the lineage-resolved data from which such
357 properties are measured, it can also say which experiment would settle the question before a trial is designed.

358 Several limitations delimit this release. Cells are well-mixed compartments without spatial structure or cell-cell
359 signalling; the pharmacology is a one-compartment model; growth is exponential with phenomenological size
360 control; and the inference layer, although general, relies on user-chosen summary statistics, so posterior calibration
361 should be checked with predictive simulations. The persister case study is deliberately minimal, a single resistance
362 gene with slow switching, and its calibration used fate fractions transcribed from a published study rather than
363 raw tracks. Its death hazard was first written without cell-cycle dependence, which matters for a drug such as
364 cisplatin whose lethality is enhanced during replication. Refitting with a cycle-dependent hazard shows that
365 the two mechanisms are not separable by these data: a model in which killing is confined to the replication
366 window accounts for the fate fractions as well as heritable protective expression does, and the fraction of the
367 hazard that is cycle-independent is not determined by them at all. Cell-cycle position at the moment the drug
368 arrives is therefore a live alternative to inherited protection, and the experiment that distinguishes them is a
369 cycle reporter read at drug addition rather than more fate counts. Six parameters fitted to six fractions leave the
370 switching rates weakly determined, so the memory of the resistant state rests on the lineage correlations used as a
371 prior rather than on the fates themselves; and that memory is carried by one slow promoter, a coarse proxy for
372 chromatin-based inheritance whose single-exponential, symmetric form is a simplification of a process that can be
373 multi-timescale and bistable. The clinical case studies use literature values for cell-cycle times, resistant fractions
374 and pharmacokinetics rather than patient-specific measurements, so their conclusions concern mechanisms and
375 rankings, not absolute survival times. In the glioblastoma study MGMT is consumed stoichiometrically, one
376 molecule per lesion repaired, so its depletion follows the cumulative exposure as a suicide enzyme's should; what
377 remains unmeasured is the rate constant of that consumption, which is why we name it as the quantity that
378 separates the regimens. MGMT is in any case only one determinant of a temozolomide response that is also
379 mismatch-repair-dependent, and the methylated and unmethylated MGMT fractions are illustrative values that
380 bracket low and high expression rather than measured quantities. The memory of the resistant state is expressed
381 in these tumour-scale studies as a number of net population doublings, not of cell divisions, and so is not the
382 division-counted memory of memory-sequencing experiments. A bridge to the Catalyst.jl ecosystem is not part of
383 this release. The benchmark results should be read as demonstrations of what the framework makes possible
384 rather than as verdicts on the methods compared, which were run with default settings on small networks.

385 Spatial coupling through neighbourhood-dependent rates, cell-cell communication, chromatin states with their
386 own inheritance rules and neural simulation-based inference are natural extensions of the existing forward model.
387 Because the framework is deterministic given a seed, fast enough to simulate tens of thousands of cells in minutes,
388 and writes AnnData files, it can be dropped into single-cell analysis pipelines as a source of mechanistically

grounded synthetic data, and into virtual-cell training and evaluation as a source of causal, out-of-distribution contexts.

Methods

Reaction model and propensities

A model is a set of species and reaction channels. Species are molecules (partitioned at division) or promoter states (inherited). For a reaction with reactant counts x_i , stoichiometric coefficients ν_i , rate constant k and volume exponent e , the mass-action propensity is

$$a = k V^e \prod_i \binom{x_i}{\nu_i},$$

where by default $e = 1 - \sum_i \nu_i$, so that zero-order production scales with volume, first-order reactions are volume independent and bimolecular reactions scale as $1/V$; this keeps concentrations invariant under growth. Regulated reactions multiply this by $b + (1 - b)F(\mathbf{c})$, where b is a basal fraction and F combines Hill terms of the regulator concentrations $c_j = x_j/V$, $h_j = c_j^{n_j}/(K_j^{n_j} + c_j^{n_j})$ for activators and $1 - h_j$ for inhibitors, either as a product (AND logic) or as $1 - \prod_j (1 - h_j)$ (OR logic). Promoter switching is represented by explicit promoter species (for instance $G_{\text{off}} \rightleftharpoons G_{\text{on}}$) grouped into promoter groups. Custom propensities can be supplied as functions of counts, volume, parameters and time.

Stochastic simulation kernels

Within one update step the cell volume and gene copy number are held fixed and the reaction network is advanced by one of four kernels: the direct method³⁸ with propensities updated through a species-reaction dependency graph; fixed-step Poisson tau-leaping³⁹; a hybrid scheme in which a leap that would drive any species negative is discarded and that interval is simulated exactly instead (the “switching” scheme of version 1, with the discarded interval resolved exactly rather than clipped); and adaptive tau-leaping with the step-size selection of⁴⁰ and the critical-reaction partition of⁴¹, falling back to a fixed number of exact events when the proposed leap is shorter than $10/a_0$. The hybrid scheme is approximate rather than exact: a leap is kept only when it leaves every species non-negative, so an accepted step follows the Poisson leap truncated to the non-negative region, and only the discarded intervals follow the exact process. The discrepancy falls with τ and is bounded by the Kolmogorov-Smirnov test of Fig. 2d, which the scheme passes at $\tau = 0.05$ and fails at $\tau = 0.2$ for the models tested; the direct method is used wherever exactness is required (Supplementary Table 1 lists the kernel correctness tests). Every kernel takes an explicit random number generator.

Cells, growth, replication and division

Each cell carries its volume V , age, generation, lineage identifiers, gene copy state, growth rate, molecule counts, a cell-specific parameter vector and its own random number generator (seeded from the parent’s generator, which makes results independent of the number of threads). Volume grows exponentially, $\dot{V} = \lambda V$, with an optional log-normal spread of λ between cells. Division is triggered by a sizer ($V \geq V_{\text{div}}$), an adder ($V - V_{\text{birth}} \geq \Delta$) or a timer ($\text{age} \geq T$), each with a per-cell noise factor drawn at birth. At division the first daughter inherits a volume fraction $f \sim \mathcal{N}(0.5, \sigma^2)$ (clipped to $[0.05, 0.95]$) and each molecule is assigned to it independently with probability f (binomial partitioning); a beta-binomial option models over-dispersed partitioning. Promoter groups are inherited by both daughters when unreplicated and split one copy per daughter, in random assignment, when

replicated. Gene replication, if enabled, doubles the promoter counts (and the propensity of reactions flagged as copy-number dependent) once the cell has completed a set fraction of its cycle, and is reset at division. Populations can be started from per-founder states and volumes (for instance a snapshot of a burnt-in constant-size population). Population control is either constant size (every second daughter replaces a uniformly chosen cell, the behaviour of version 1), free growth (all daughters kept, with uniform subsampling above a cap and book-keeping of the true size) or logistic (free growth with an additional hazard $\lambda N/K$).

Drug and perturbation layer

A perturbation combines a dose schedule $d(t)$ (constant, pulsed, piecewise, one-compartment bolus pharmacokinetics or an arbitrary function) with effects evaluated for every cell at every step: a death hazard

$$h = h_{\max} \frac{d^m}{\text{EC}_{50}^m + d^m} \cdot \frac{K^q}{K^q + c_P^q},$$

where c_P is the concentration of a protective species (survival probability over a step is $e^{-h \Delta t}$); growth inhibition by $1/(1 + (d/\text{IC}_{50})^m)$; a dose-independent fitness cost of a cell state, $1 - c_{\max} c_P^q/(K^q + c_P^q)$, on the growth rate; multiplicative modulation of any parameter by a function of the dose; and genetic perturbations (knockdown, knockout, over-expression) of any parameter in a random fraction of cells, the flag being inherited by daughters.

Two further effects couple the drug to the cell cycle and to the protein that resists it. A cell-cycle-dependent hazard multiplies h by

$$\beta + (1 - \beta) \exp \left[-\frac{1}{2} \left(\frac{\varphi - \varphi_0}{w} \right)^2 \right],$$

where φ is the cell's progress through its division cycle (0 at birth, 1 at division), so that killing peaks at φ_0 and falls to a fraction β of the peak elsewhere; $\beta = 1$ is the cycle-independent hazard and $\beta = 0$ confines killing to a window of width w around φ_0 . This represents an agent whose lesions are converted into death during replication. Stoichiometric consumption of a protective protein removes molecules of it at rate

$$k d V \frac{c_P}{c_P + K_m},$$

by a Poisson draw over the step. Because lesions form in proportion to the dose, the rate is set by lesion formation alone while the protein is abundant, so the pool follows the cumulative exposure rather than the peak concentration, and falls with what remains once the pool runs low; resynthesis is whatever the reaction model provides. This is the behaviour of a suicide enzyme such as MGMT, which is inactivated by the very lesion it repairs.

Lineage statistics

The lineage table records, for every cell, its parent, founder (clone), generation, birth and end times, fate (divided, died, removed, alive) and its volume and state at birth and at the end of its life. From it we compute Pearson correlations of expression at division between mothers and daughters, sisters and cousins; the lineage autocorrelation r_g between a cell and its ancestor g generations earlier, and the memory timescale τ from $r_g = e^{-g/\tau}$; single-lineage traces by following the ancestral line of a surviving cell; the coefficient of variation squared of concentration across the population versus along single lineages; Luria-Delbrück fluctuation tests (variance of a phenotype across clones grown from single cells compared with the binomial expectation); and MemorySeq-style clonal scores⁴², the variance of clone-mean expression divided by its permutation-null expectation.

458 Observation models

459 True counts are converted to observations by a scRNA-seq model in which each cell captures every molecule
460 independently with a Beta-distributed efficiency (mean and coefficient of variation set by the protocol), optional
461 multinomial down-sampling to a Poisson-distributed depth, optional additional dropout and batch scaling; an
462 smFISH model with a fixed detection efficiency; and a time-lapse model that samples a species along lineages at a
463 fixed interval with additive or multiplicative measurement noise. Because counts scale with volume, library size
464 correlates with cell size, reproducing the transcriptome-size confound of real data. Outputs are written as CSV,
465 as AnnData-compatible HDF5 files, or as Newick lineage trees.

466 Inference

467 Snapshot counts of a gene in non-dividing cells follow the Beta-Poisson law of the telegraph model²⁷, which we
468 evaluate through Kummer’s confluent hypergeometric function with rescaled series summation, and maximise by
469 a grid-initialised Nelder-Mead search in log-parameter space. For any model, including populations, we implement
470 sequential Monte Carlo approximate Bayesian computation⁴³ with an adaptive tolerance schedule (the α quantile
471 of the previous generation’s distances), Gaussian perturbation kernels of variance twice the weighted population
472 variance, importance weights against the prior, and parallel evaluation of the user-supplied distance function; for
473 count distributions we use the Wasserstein-1 distance between simulated and observed samples³¹.

474 Calibration to laboratory data and validation strategy

475 Reference values were transcribed from the text and tables of the cited articles into `paper/data/` with their
476 provenance (Supplementary Table 4). For the cisplatin case study the model is a resistance gene with slow promoter
477 switching whose protein product protects against a dose-dependent death hazard and against a dose-dependent
478 growth arrest (Methods above), with a 24-h cell cycle (at the fast end of reported U2OS doubling times and
479 consistent with the intermitotic times tracked by⁶). Its six free parameters (activation and inactivation rates,
480 maximal death rate, EC_{50} and Hill coefficient of death, IC_{50} of growth arrest) were fitted by the method of
481 simulated moments: the root-mean-square difference between simulated and observed fractions of cells that died,
482 divided or survived without dividing during 72 h of drug, for the training concentrations (7 and 13 μ M), plus a
483 weak penalty keeping the memory of the resistant state near the two to three generations inferred from lineage
484 correlations, was minimised with the bounded optimiser described below using common random numbers (the
485 same seeds for every candidate). Validation used only quantities not fitted: the fate fractions at the held-out
486 concentration (10 μ M), the biphasic kill curve and its plateau, the extent of lineage correlations in fate (sisters,
487 first, second and third cousins) and the invariance of single-cell division and death times across concentrations.
488 The protocol is deliberately generic: any observable computable from a `PopulationResult` (kill curves, fate
489 fractions, lineage correlations, snapshot distributions) can enter the distance, and the split between training and
490 validation conditions is declared before fitting.

491 Clinical regimens, treatment outcomes and schedule optimisation

492 Regimens are built from the schedule primitives of the perturbation layer: `PulsedDose` with a finite number of
493 cycles, `PiecewiseDose` for arbitrary calendars (the S1320 schedule), `daily_boluses` and `cycle_days` for daily
494 oral doses with one-compartment pharmacokinetics, and `AdaptiveDose`, a feedback schedule that withdraws the
495 drug when the population falls below a fraction of its initial size and resumes it when the population regrows above
496 another fraction, evaluated with the current population size at every step³³. Outcomes computed from a simulation
497 are the net growth rate, the depth of response ($\log_{10}$ of the ratio of the initial to the smallest population), the

time to progression (first time the population exceeds a threshold, 1.2, times its running minimum, the analogue of RECIST progression from the nadir; censored at the end of the simulation), the extinction probability over replicates and the cumulative dose actually applied. `optimize_schedule` minimises any such outcome over a box of schedule parameters: a log-spaced grid (or a Latin hypercube when the grid is large) is evaluated with the same random seeds for every candidate, and the best point is refined by Nelder-Mead in a bounded transform of the box; constraints such as a fixed cumulative dose enter as penalties or as the parameterisation (for instance a per-day dose equal to the budget divided by the number of dosing days). The case studies are run at the scale of a tumour rather than a culture: the cell cycle of the model is the net doubling time of the tumour (4 weeks for melanoma, 40 days for glioblastoma). The glioblastoma value lies within the volume doubling times measured for untreated tumours, a median of 21 days in one magnetic-resonance series and 50 days in another^{44,45}; the melanoma value is an assumption, since net doubling times of treated cutaneous melanoma are not established, and it sets the time axis of the melanoma case study rather than the ranking of schedules, which follows from the resistance mechanism. Growth is exponential throughout, which is a simplification: the same measurements favour Gompertzian slowing in large glioblastomas⁴⁵, so absolute regrowth times in Fig. 8 are shorter than a saturating model would give. Gene-expression rates are scaled to the hour (mRNA half-life 7 h and protein half-life 35 h, generic mammalian values; at the tumour time scale both species turn over many times per generation, so the results are insensitive to them and all heritable memory resides in the promoter), and founders are drawn from the stationary distribution of promoter states. The melanoma case uses a pre-resistant fraction of 1:200 (a per-marker frequency² used as a proxy for the functional pre-resistant fraction, which is not the same quantity), with a memory of five net doublings (the model's generation is the net doubling of the tumour, a modelling choice rather than the cell-division-counted memory of memory sequencing⁴²) and drug effects at the fractional-killing dose: sensitive cells arrest and die with a half-life of two weeks, cells in the resistant state are protected, and the drug stabilises the resistant state ten-fold (the drug-induced reprogramming of²). Three variants of the resistant state are compared: full protection without cost; a 50% fitness cost in the absence of drug³⁴; and partial protection, in which resistant cells grow at half speed under drug. Progression is scored in two ways: from the nadir, as in RECIST (a 20% increase in diameter, taken as a 1.73-fold increase in cell number, after the 8-week lead-in; a well-mixed cell count cannot represent RECIST's absolute 5 mm threshold, its new-lesion and non-target criteria or its sum-of-diameters basis, and it assumes cell number proportional to volume, so this is an analogue rather than the clinical end point), and as loss of control, the time after randomisation at which the population exceeds 120% of its pre-treatment size, the end point of adaptive-therapy studies, which is the one used to compare schedules that deliberately let the tumour regrow. The temozolomide case uses a one-compartment elimination half-life of 2.1 h³⁷ (a plasma value from a three-compartment study; the brain-tumour compartment sees roughly a fifth of the plasma exposure with slower kinetics, so this is a surrogate), the regimens of RTOG 0525³⁶ (200 mg/m² on days 1-5 or 100 mg/m² on days 1-21 of 28-day cycles, six cycles), MGMT-expressing fractions of 1% and 30% for methylated and unmethylated promoters (illustrative values bracketing low and high expression rather than measured fractions; methylation shifts graded expression rather than a bimodal positive fraction), and MGMT either stable or consumed stoichiometrically by the drug, one molecule per lesion repaired ($k = 300$, $K_m = 150$ in the expression above). The rate constant was chosen so that the consumption rate at the standard bolus peak equals that of the dose-scaled first-order degradation used in an earlier version of this model, which isolates the consequence of letting depletion follow the cumulative exposure instead of the peak concentration; this is the suicide-enzyme behaviour that motivated dose-dense schedules.

Random regulatory networks

Random networks are built with a chosen number of activating and inhibiting edges under Erdős-Rényi or preferential-attachment (scale-free) sampling; regulation acts on the transcription rate of constitutive genes or on the promoter activation rate of telegraph genes through AND-combined Hill functions; per-gene rate constants can be drawn log-uniformly from ranges.

Validation protocols

Exactness was tested by Kolmogorov-Smirnov distances between simulated stationary distributions and the analytic laws (Poisson for birth-death, Beta-Poisson for the telegraph model, negative binomial for bursty protein production^{46,47}), by agreement between kernels, and by a two-sample comparison against JumpProcesses.jl⁴⁸ on the same model. Population-level properties were compared with the analytic expectations for cell-size scaling and concentration homeostasis²³, partitioning noise⁴⁹, lineage versus population statistics¹¹ and cell-cycle-dependent burst parameters¹⁰. Persister simulations were compared with the quantitative signatures reported by^{6,30} and⁹, and with the phenotypic-selection model of²⁹. All scripts, seeds and outputs are in the `paper/` directory of the repository.

Identifiability and numerical robustness of the calibration

Because six parameters were fitted to six measured fractions, the calibration was examined in three ways. A conditional profile moves one parameter at a time over its full range on a log grid of eleven points, holding the others at their fitted values, and records the root-mean-square error of the fate fractions without the memory prior; parameters whose error stays within the binomial standard error of the measurement over a wide range are not determined by these data. A two-dimensional slice does the same over the memory of the resistant state and the fraction of cells in it, the two quantities that trade off most directly, with the death parameters held at their fitted values; a slice was used in preference to a profile with re-optimisation because, at this cost per evaluation, the inner search did not reliably recover the fitted solution and reported spuriously poor fits at short memory. Numerical robustness was checked by halving the integration step and comparing the simulated fractions with the seed-to-seed spread at the fitted parameters.

Because the lethality of cisplatin is enhanced during replication, the position of a cell in its division cycle when the drug arrives is a competing explanation for the dose-invariance of death timing and for the fraction of cells that survive without dividing. To test it, the same six parameters were refitted together with a seventh, the cycle-independent fraction β of the death hazard, with the peak of killing at the middle of the cycle ($\varphi_0 = 0.5$, $w = 0.15$), under the same objective, bounds, optimiser budget and seeds as the main fit. The two fits are compared by the root-mean-square error of the fate fractions without the memory prior, by their predictions at the held-out concentration, and by a conditional profile over β with the remaining parameters held at the cycle-dependent fit.

Statistics and reproducibility

No statistical method was used to predetermine sample size, and no experimental data were generated in this study; sample sizes are the numbers of simulated cells, founders or seeds stated in each figure legend and in Supplementary Table 2. No data were excluded. Randomisation and blinding do not apply: every result is a deterministic function of an explicit random seed, and all seeds are fixed in the scripts that generate the figures. Replicates are independent simulations differing only in seed; where a mean and a standard deviation

are reported, the number of seeds is given in the legend. Kolmogorov-Smirnov distances are compared with the 99% critical value for the stated sample size. Distributional agreement between kernels and with an independent implementation was assessed by two-sample Kolmogorov-Smirnov statistics. Correlations are Pearson correlations unless stated otherwise; correlations of binary fates between related cells are phi coefficients. Posterior summaries are medians, because the posteriors are asymmetric on the log scale. Multithreaded runs give results identical to single-threaded runs, because each cell carries its own random number generator seeded from its parent.

Data availability

The laboratory and clinical reference values used for calibration and validation were transcribed from the published articles cited in Supplementary Table 4 and are provided, with their provenance, in the `paper/data/` directory of the software repository (<https://github.com/ayoublasri/Biomodelling.jl>). All other data in this Article are simulated and are regenerated from fixed seeds by the scripts in the `paper/` directory; the generated tables are archived with the software release.


Code availability

Biomodelling.jl 2.0 is open source under the MIT licence at <https://github.com/ayoublasri/Biomodelling.jl>, with documentation, tutorials and a test suite that runs under continuous integration. Every figure is regenerated from fixed seeds by the scripts in `paper/`. After instantiating the environment (`julia --project=paper -e 'using Pkg; Pkg.develop(path="."); Pkg.instantiate()'`), the simulation data are produced by `julia --project=paper -t 4 paper/run_all.jl`, the Python network-inference and imputation benchmarks of Fig. 5b–d by `python paper/scripts/benchmarks_python.py`, and the identifiability analysis of Supplementary Fig. 3 by `julia --project=paper paper/scripts/fig7_calibration.jl` profile; the panels are then assembled by `python paper/fig1_schematic.py`, `python paper/plot_all.py` and `python paper/plot_supplement.py`. Full instructions, including the four-thread timings of Supplementary Table 3, are in `paper/README.md`.

Software


Biomodelling.jl 2.0 is written in Julia (≥ 1.10) and depends on `Distributions.jl`, `PoissonRandom.jl`, `QuadGK.jl`, `Optim.jl`, `SpecialFunctions.jl` and `StatsBase.jl`; `HDF5.jl` is an optional extension. Benchmarks used absolute Pearson and Spearman correlations and scikit-learn random forests re-implementing GENIE3⁵⁰, MAGIC⁵¹ and kNN-smoothing⁵². The test suite (Aqua quality checks and statistical tests) runs under GitHub Actions on Julia 1.10 and the latest release.

Figure legends




607

Fig. 1 | The Biomodelling.jl 2.0 framework. **a**, One cell. Reactions run at the cell's current volume, with propensities scaled so that concentrations are invariant under growth; genes replicate at a set point of the cycle; division splits every molecule binomially while both daughters inherit the promoter state. Inset: exponential volume growth, halved at each division. **b**, Populations under free growth, constant-size or logistic control, with dose schedules and a state-dependent death hazard evaluated for every cell at every step, and the lineage record from which heritability, autocorrelation, memory timescales, clone diversity and fluctuation tests are computed; crosses mark cells killed by the drug. **c**, Observation models turn true counts into scRNA-seq counts with capture, depth, dropout and batch effects, into smFISH spots, or into time-lapse reporter traces; outputs are CSV, AnnData and Newick. **d**, Uses of the forward model: detection of memory genes, benchmarking of network and perturbation inference against known ground truth, likelihood-free parameter inference, and schedule optimisation over the period and the fraction of time on drug (star, optimum). **e**, A complete experiment in twelve lines of code.




620

621 **Fig. 2 | The engine matches exact laws and an independent implementation.** **a**, Stationary distribution
622 of a birth-death process (bars, 50,000 simulated cells; line, Poisson law). **b**, Telegraph model in an intermediate
623 and a bursty regime against the Beta-Poisson law. **c**, Bursty protein production against the negative binomial. **d**,
624 Kolmogorov-Smirnov distance to the exact law for every kernel (n = 20,000; dotted line, 99% critical value) with
625 wall-clock time per 10^4 cells. **e**, Telegraph model simulated with Biomodelling.jl and with JumpProcesses.jl on
626 the same model; the two-sample Kolmogorov-Smirnov statistic lies below its 99% critical value. **f**, Runtime of
627 population simulations against the number of cells (10 genes) and genes (1,000 cells) for the exact and hybrid
628 kernels.




629

630 **Fig. 3 | Growth, division and partitioning generate memory, size scaling and cell-cycle effects.** **a**, A
631 single lineage: volume, mRNA and protein of a slowly switching gene, with the promoter-on state shaded. **b**,
632 mRNA count against cell volume for a fast-switching gene without replication (points; line, mean concentration
633 $\times V$) and with gene replication at mid-cycle (binned means). **c**, Correlation at division between mothers
634 and daughters, sisters and cousins as a function of the promoter switching rate (protein concentration, solid;
635 mother-daughter mRNA, dashed); dotted line, inverse cell-cycle time. **d**, Memory timescale against the measured
636 cell-cycle time for two switching rates; dotted lines, $1/(k_{on} + k_{off})$. **e**, Squared coefficient of variation of a stable
637 protein's concentration across the population and along single lineages, for increasing partitioning noise and
638 cell-cycle-time noise. **f**, Telegraph parameters fitted to cells before replication (one gene copy), after replication
639 (two copies) and pooled; dotted lines, per-allele values.



640


641 **Fig. 4 | A single resistance gene with slow promoter switching reproduces persister biology.** **a**,
642 Population size under continuous, pulsed and reduced-dose schedules (shading, treatment window). **b**, Kill curves
643 at increasing dose. **c**, Net population decay rate over the first 30 time units of treatment against dose, with
644 the mean division time of cells born under drug (inset), which is unchanged. **d**, Probability that a cell's lineage
645 survives the treatment (at least one descendant alive at the end of the run) given that its sister's or cousin's
646 lineage survived, against the marginal survival probability, for cells born in the 25 time units before the drug;
647 memory gene and fast-switching control, with the fold enrichment above each bar. **e**, Effective number of clones
648 (barcodes) before and after drug for pre-existing and drug-induced tolerance (mean $\pm$ s.d. of three seeds). **f**,
649 Surviving clones after drug without memory disruption, with promoter switching accelerated 20-fold from two
650 cell-cycle times before the drug until the drug arrives, and with the acceleration maintained during exposure
651 (mean $\pm$ s.d. of five seeds). **g**, Fraction of high-expressing cells before, during and after a drug pulse for slow
652 and fast switching. **h**, Net population growth rate over a 200-time-unit course as a function of release period
653 (time units without drug between 20-time-unit exposures) and dose, for pre-existing tolerance without and with a
654 fitness cost of the resistant state and for drug-induced tolerance; the box marks the most effective schedule.




655

656 **Fig. 5 | Ground-truth benchmarks.** **a**, Clonal variance score against the known memory time of 40
 657 independent genes, from true concentrations and from sequenced counts; dotted line, cell-cycle time. **b**, Area
 658 under the precision-recall curve of correlation (Pearson, Spearman) and GENIE3 network inference on the same
 659 30-gene network observed at fixed volume, in a dividing population (raw counts, concentrations, and concentrations
 660 regressed on the cell-cycle covariates copy number and age) and after sequencing (raw and library-size normalised
 661 counts); dotted line, random ranking. **c**, Effect of kNN-smoothing and MAGIC imputation before inference. **d**,
 662 Coefficient of determination of predicted $\log_2$ fold changes after knocking down the five most connected regulators,
 663 for a no-change baseline and for propagation of observational correlations.


a non-dividing cells: ABC vs exact MLE




b dividing population: division-aware ABC vs naive fit



c drug parameters from kill curve and sister fates




d posterior predictive check



664

Fig. 6 | Inference with the forward model. **a**, Posterior distributions of telegraph parameters from ABC-SMC on a snapshot of non-dividing cells, with the truth and the exact maximum-likelihood estimate. **b**, The same for a snapshot of a growing and dividing population: the naive telegraph fit is biased while division-aware ABC recovers the truth. **c**, Posterior particles for the maximal death rate and the protection threshold inferred from a kill curve and sister fate concordance: the two parameters trade off along a ridge that contains the truth. **d**, Posterior predictive kill curves simulated at the posterior median against the observed curve.



671

672 **Fig. 7 | Calibration to laboratory time-lapse data and validation on held-out conditions.** **a**, Fates
 673 over 72 h of cisplatin of U2OS cells present at drug addition, observed by time-lapse microscopy (grey; ref.⁶) and
 674 simulated by the calibrated model (colours; mean $\pm$ s.d. of four seeds); the model was calibrated on 7 and 13 μM
 675 and the 10 μM condition is predicted. **b**, Predicted population size at the three concentrations and at 11.5 μM ,
 676 from two days before the drug to seven days of exposure. **c**, Correlation of death fates between sisters, first, second
 677 and third cousins present at drug addition, and between unrelated cells (mean $\pm$ s.d. of three seeds), against the
 678 observation that correlations extend to second but not third cousins. **d**, Mean time to death of cells present at
 679 drug addition against concentration (mean $\pm$ s.d. of four seeds); division times are not shown because too few
 680 cells born under drug divided again to estimate them (at most four per condition). **e**, Calibration landscape: every
 681 parameter set evaluated by the optimiser, coloured by its distance to the training data. **f**, Calibrated parameters.

693 days) temozolomide regimens of the RTOG 0525 trial and a 21-day regimen at the standard cumulative dose,
694 with one-compartment pharmacokinetics (elimination half-life 2.1 h) and stoichiometric consumption of MGMT
695 by the drug, one molecule per lesion repaired, for populations in which 1% (methylated promoter, solid) or 30%
696 (unmethylated, dashed) of cells express MGMT. **g**, Depth of response over six cycles for the three regimens,
697 the two populations and MGMT either stable or consumed by the drug (mean $\pm$ s.d. of two seeds); a bar at
698 zero marks a regimen that achieves no net kill (the 21-day schedule at the standard cumulative dose). **h**, Final
699 population size against the number of dosing days per 28-day cycle at the standard cumulative dose.

700 Acknowledgements

701 Version 1 of the framework was developed with Marc Sturrock and Vahid Shahrezaei, whose contributions to that
702 work are described in ref.²⁵. This Article and version 2.0 of the software were authored with Claude (Anthropic),
703 an AI assistant, under the direction and supervision of the author, who verified every result and takes full
704 responsibility for the content.

705 Author contributions

706 A.L. conceived the study, defined its scope, the validation strategy and the case studies, supervised the work,
707 verified all results and wrote the final version of the manuscript. The implementation of the software, the
708 simulations, the analyses, the figures and the first draft of the manuscript were produced with Claude (Anthropic)
709 under the author's direction.

710 Competing interests

711 The author declares no competing interests.

712 Additional information

713 **Supplementary information** accompanies this Article. Correspondence and requests for materials should be
714 addressed to A.L.

715 References

- 716 1. Sharma, S. V. *et al.* A chromatin-mediated reversible drug-tolerant state in cancer cell subpopulations. *Cell* **141**, 69–80 (2010).
- 717 2. Shaffer, S. M. *et al.* Rare cell variability and drug-induced reprogramming as a mode of cancer drug resistance. *Nature* **546**, 431–435 (2017).
- 718 3. Hata, A. N. *et al.* Tumor cells can follow distinct evolutionary paths to become resistant to epidermal growth factor receptor inhibition. *Nature Medicine* **22**, 262–269 (2016).
- 719 4. Oren, Y. *et al.* Cycling cancer persister cells arise from lineages with distinct programs. *Nature* **596**, 576–582 (2021).
- 720 5. França, G. S. *et al.* Cellular adaptation to cancer therapy along a resistance continuum. *Nature* **631**, 876–883 (2024).

6. Iyer, A., Alva, A., Granada, A. E. & Chakrabarti, S. Inheritable cell-states shape drug-persister correlations and population dynamics in cancer cells. *PLoS Computational Biology* **21**, e1013446 (2025).
7. Lin, Y., Chen, X., Wu, L., Zhou, Y. & Lin, Y. Widespread transcriptional memory shapes heritable states and functional heterogeneity in cancer and stem cells. *Cell Systems* **17**, 101621 (2026).
8. Ghosh, S., Chakrabarti, S. & Raju, A. Identifying memory gene expression from single sample scRNA-seq data using power law signatures. *bioRxiv* <https://doi.org/10.1101/2025.01.15.633183> (2025) doi:10.1101/2025.01.15.633183.
9. Harmange, G. *et al.* Disrupting cellular memory to overcome drug resistance. *Nature Communications* **14**, 7130 (2023).
10. Sukys, A. & Grima, R. Cell-cycle dependence of bursty gene expression: Insights from fitting mechanistic models to single-cell RNA-seq data. *Nucleic Acids Research* **53**, gkaf295 (2025).
11. Zhang, Z. *et al.* Stochastic gene expression in proliferating cells: Differing noise intensity in single-cell and population perspectives. *PLoS Computational Biology* **21**, e1013014 (2025).
12. Okochi, Y., Sawazaki, Y., Kondo, Y. & Naoki, H. Inferring division-associated stochasticity from time-series single-cell transcriptomes. *bioRxiv* <https://doi.org/10.64898/2026.04.14.718485> (2026) doi:10.64898/2026.04.14.718485.
13. Ahlmann-Eltze, C., Huber, W. & Anders, S. Deep-learning-based gene perturbation effect prediction does not yet outperform simple linear baselines. *Nature Methods* **22**, (2025).
14. Vollenweider, M. S. & Bühlmann, P. Signal, bounds, and baselines: Principles for rigorous evaluation of high-dimensional biological perturbation prediction. *bioRxiv* <https://doi.org/10.64898/2026.04.20.719650> (2026) doi:10.64898/2026.04.20.719650.
15. Dibaeinia, P. *et al.* Virtual cells need context, not just scale. *bioRxiv* <https://doi.org/10.64898/2026.02.04.703804> (2026) doi:10.64898/2026.02.04.703804.
16. Bunne, C. *et al.* How to build the virtual cell with artificial intelligence: Priorities and opportunities. *Cell* **187**, 7045–7063 (2024).
17. Dibaeinia, P. & Sinha, S. SERGIO: A single-cell expression simulator guided by gene regulatory networks. *Cell Systems* **11**, 252–271 (2020).
18. Cannoodt, R., Saelens, W., Deconinck, L. & Saeys, Y. Spearheading future omics analyses using dyngen, a multi-modal simulator of single cells. *Nature Communications* **12**, 3942 (2021).
19. Li, H., Zhang, Z., Squires, M., Chen, X. & Zhang, X. scMultiSim: Simulation of single-cell multi-omics and spatial data guided by gene regulatory networks and cell–cell interactions. *Nature Methods* <https://doi.org/10.1038/s41592-025-02651-0> (2025) doi:10.1038/s41592-025-02651-0.
20. Zinati, Y., Takiddeen, A. & Emad, A. GRouNdGAN: GRN-guided simulation of single-cell RNA-seq data using causal generative adversarial networks. *Nature Communications* **15**, 4055 (2024).
21. Pan, X., Li, H. & Zhang, X. TedSim: Temporal dynamics simulation of single-cell RNA sequencing data and cell division history. *Nucleic Acids Research* **50**, 4272–4288 (2022).
22. Pratapa, A., Jalihal, A. P., Law, J. N., Bharadwaj, A. & Murali, T. M. Benchmarking algorithms for gene regulatory network inference from single-cell transcriptomic data. *Nature Methods* **17**, 147–154 (2020).
23. Bertaux, F., Marguerat, S. & Shahrezaei, V. Division rate, cell size and proteome allocation: Impact on gene expression noise and implications for the dynamics of genetic circuits. *Royal Society Open Science* **5**, 172234 (2018).
24. Piho, P. & Thomas, P. AgentBasedModeling.jl: A tool for stochastic simulation of structured population dynamics. *PLoS Computational Biology* **22**, e1014366 (2026).

- 740 25. Lasri, A., Shahrezaei, V. & Sturrock, M. Benchmarking imputation methods for network inference using a novel method of synthetic scRNA-seq data generation. *BMC Bioinformatics* **23**, 236 (2022).
- 741 26. Jones, M. G. *et al.* Inference of single-cell phylogenies from lineage tracing data using Cassiopeia. *Genome Biology* **21**, 92 (2020).
- 742 27. Peccoud, J. & Ycart, B. Markovian modeling of gene-product synthesis. *Theoretical Population Biology* **48**, 222–234 (1995).
- 743 28. Gunnarsson, E. B., De, S., Leder, K. & Foo, J. Understanding the role of phenotypic switching in cancer drug resistance. *Journal of Theoretical Biology* **490**, 110162 (2020).
- 744 29. Lasri, A. *et al.* Phenotypic selection through cell death: Stochastic modelling of O-6-methylguanine-DNA methyltransferase dynamics. *Royal Society Open Science* **7**, 191243 (2020).
- 745 30. Corigliano, M., Di Bernardo, A., Cosentino Lagomarsino, M. & Pompei, S. Optimal treatment for drug-induced cancer persists involves release periods and intermediate drug doses. *PRX Life* <https://doi.org/10.1103/mh5v-lcyw> (2025) doi:10.1103/mh5v-lcyw.
- 746 31. Bernton, E., Jacob, P. E., Gerber, M. & Robert, C. P. Approximate Bayesian computation with the Wasserstein distance. *Journal of the Royal Statistical Society: Series B* **81**, 235–269 (2019).
- 747 32. Algazi, A. P. *et al.* Continuous versus intermittent BRAF and MEK inhibition in patients with BRAF-mutated melanoma: A randomized phase 2 trial. *Nature Medicine* **26**, 1564–1568 (2020).
- 748 33. Zhang, J., Cunningham, J. J., Brown, J. S. & Gatenby, R. A. Integrating evolutionary dynamics into treatment of metastatic castrate-resistant prostate cancer. *Nature Communications* **8**, 1816 (2017).
- 749 34. Das Thakur, M. *et al.* Modelling vemurafenib resistance in melanoma reveals a strategy to forestall drug resistance. *Nature* **494**, 251–255 (2013).
- 750 35. Strobl, M. A. R. *et al.* Turnover modulates the need for a cost of resistance in adaptive therapy. *Cancer Research* **81**, 1135–1147 (2021).
- 751 36. Gilbert, M. R. *et al.* Dose-dense temozolomide for newly diagnosed glioblastoma: A randomized phase III clinical trial. *Journal of Clinical Oncology* **31**, 4085–4091 (2013).
- 752 37. Ostermann, S. *et al.* Plasma and cerebrospinal fluid population pharmacokinetics of temozolomide in malignant glioma patients. *Clinical Cancer Research* **10**, 3728–3736 (2004).
- 753 38. Gillespie, D. T. Exact stochastic simulation of coupled chemical reactions. *The Journal of Physical Chemistry* **81**, 2340–2361 (1977).
- 754 39. Gillespie, D. T. Approximate accelerated stochastic simulation of chemically reacting systems. *The Journal of Chemical Physics* **115**, 1716–1733 (2001).
- 755 40. Cao, Y., Gillespie, D. T. & Petzold, L. R. Efficient step size selection for the tau-leaping simulation method. *The Journal of Chemical Physics* **124**, 044109 (2006).
- 756 41. Cao, Y., Gillespie, D. T. & Petzold, L. R. Avoiding negative populations in explicit Poisson tau-leaping. *The Journal of Chemical Physics* **123**, 054104 (2005).
- 757 42. Shaffer, S. M. *et al.* Memory sequencing reveals heritable single-cell gene expression programs associated with distinct cellular behaviors. *Cell* **182**, 947–959 (2020).
- 758 43. Toni, T., Welch, D., Strelkowa, N., Ipsen, A. & Stumpf, M. P. H. Approximate Bayesian computation scheme for parameter inference and model selection in dynamical systems. *Journal of the Royal Society Interface* **6**, 187–202 (2009).
- 759 44. Ellingson, B. M. *et al.* Contrast-enhancing tumor growth dynamics of preoperative, treatment-naive human glioblastoma. *Cancer* **122**, 1718–1727 (2016).

- 760 45. Stensj en, A. L. *et al.* Growth dynamics of untreated glioblastomas in vivo. *Neuro-Oncology* **17**, 1402–1411 (2015).
- 761 46. Friedman, N., Cai, L. & Xie, X. S. Linking stochastic dynamics to population distribution: An analytical framework of gene expression. *Physical Review Letters* **97**, 168302 (2006).
- 762 47. Shahrezaei, V. & Swain, P. S. Analytical distributions for stochastic gene expression. *Proceedings of the National Academy of Sciences* **105**, 17256–17261 (2008).
- 763 48. Rackauckas, C. & Nie, Q. DifferentialEquations.jl – a performant and feature-rich ecosystem for solving differential equations in Julia. *Journal of Open Research Software* vol. 5 15 (2017).
- 764 49. Huh, D. & Paulsson, J. Non-genetic heterogeneity from stochastic partitioning at cell division. *Nature Genetics* **43**, 95–100 (2011).
- 765 50. Huynh-Thu, V. A., Irrthum, A., Wehenkel, L. & Geurts, P. Inferring regulatory networks from expression data using tree-based methods. *PLoS ONE* **5**, e12776 (2010).
- 766 51. Dijk, D. van *et al.* Recovering gene interactions from single-cell data using data diffusion. *Cell* **174**, 716–729 (2018).
- 767 52. Wagner, F., Yan, Y. & Yanai, I. K-nearest neighbor smoothing for high-throughput single-cell RNA-Seq data. *bioRxiv* <https://doi.org/10.1101/217737> (2018) doi:10.1101/217737.