

Mechanistic simulation of heritable expression states, cell division and drug response in single-cell populations

Ayoub Lasri

Independent researcher, Dublin, Ireland. Correspondence: alasri@berache.com

Abstract

Heritable differences in gene expression decide which cancer cells survive a drug, yet the simulators used to benchmark single-cell methods model neither growth nor division, and stochastic simulators of dividing cells carry no drug layer. Here we present `Biomodelling.jl` 2.0, in which stochastic reaction kinetics run inside growing, dividing, drug-treated cells. Exact, hybrid and adaptive kernels couple to exponential growth, volume-scaled transcription, gene replication, division with partitioning, dose schedules with state-dependent killing, a lineage record, observation models and likelihood-free inference. Expression memory, cell-size scaling and cell-cycle-dependent bursting emerge from this physiology rather than imposed. One resistance gene with slow promoter switching reproduces the lineage and memory signatures of drug tolerance, and supplies ground truth that benchmarks of memory genes, network inference and perturbation prediction lack. Calibrated to time-lapse measurements of cisplatin-treated cells, it reproduces the dose-invariance of single-cell timing and the kin fate correlations, though six parameters fitted to six fate fractions leave that memory undetermined, and a cycle-gated hazard fits as well. Its schedule optimiser reproduces the opposite rankings of intermittent dosing reported in a melanoma trial and in xenografts, and shows the ranking to turn on two measurable cell properties: growth under drug and fitness without it.

Author summary

Two cancer cells with the same genome, given the same drug, can meet opposite fates: one dies and the other survives. Much of that difference comes from how strongly each cell happens to be expressing particular genes when treatment begins, and cells pass those expression states to their daughters when they divide. The simulators used to test single-cell analysis methods do not model growth or division, so they cannot produce such inherited states; the simulators that do model dividing cells carry no drug.

We built `Biomodelling.jl` 2.0, in which chemical reactions play out inside cells that grow, copy their genes, divide, and meet drugs on realistic schedules. Expression memory, cell-size effects and cell-cycle dependence emerge from that physiology rather than being assumed. We checked the simulator against cases whose answers are known exactly in mathematics, fitted it to published time-lapse recordings of cells under chemotherapy, and asked it to predict a concentration it had not been fitted to.

We then put a clinical question to it: should a drug be given continuously, or with breaks? The answer turns on two properties of resistant cells that experiments can measure.

Isogenic cells differ in their transcriptional states, those differences are often inherited across several divisions, and the inherited state biases what a cell does under stress. Cancer cells that survive targeted or cytotoxic therapy frequently do so without a resistance mutation: they occupy a rare, reversible, drug-tolerant expression state that

36 pre-exists the treatment or is induced by it [1–5]. Lineage-resolved experiments now measure this memory directly.
37 The fates of sisters and cousins are correlated, clonal barcode diversity is preserved or collapsed depending on
38 whether tolerance pre-exists or is induced, and the timescale of expression fluctuations relative to the cell cycle
39 decides which genes carry memory at all [6–9]. In parallel, mechanistic analyses of single-cell RNA-sequencing
40 data have shown that transcriptional kinetics can be inferred correctly only when gene copy number, cell age and
41 partitioning at division are modelled explicitly [10–12], and evaluations of “virtual cell” models for perturbation
42 prediction keep finding that deep-learning approaches fail to beat simple baselines, largely for want of mechanism
43 and causal structure [13–16].

44 These developments create a demand that existing simulators do not meet. The simulators used to benchmark
45 single-cell methods divide according to whether they supply a regulatory network as ground truth, and neither
46 kind models cell growth, volume-dependent transcription, gene replication or the partitioning of molecules at
47 division. Splatter and scDesign3 draw counts from statistical models fitted to reference data, and SymSim
48 draws them from a two-state promoter whose kinetic parameters are products of cell-specific and gene-specific
49 factors; none of the three supplies a network [17–19]. SERGIO, dyngen and scMultiSim simulate the kinetics
50 of a user-specified network [20–22], and GRouNdGAN imposes one on an adversarially trained generator, so
51 that each gene’s counts follow from those of its transcription factors [23]. Most of these simulators place cells
52 on a differentiation trajectory, and lineage simulators such as TedSim and Cassiopeia [24,25] generate division
53 trees, but the expression attached to those trees is phenomenological and largely Markovian. Expression memory,
54 cell-size scaling, cell-cycle-dependent bursting and persister dynamics therefore have to be imposed on them rather
55 than emerging, and none provides the ground truth that the methods now being developed for exactly those
56 phenomena require [26]. Stochastic simulators of growing and dividing cells do exist, and the most complete of
57 them couples reaction kinetics to growth, division and inference from lineage data [27–30]; what none of them
58 adds is a drug and schedule layer, observation models for the measurements these questions are asked of, and
59 benchmark generation. Agent-based tumour simulators supply the drug response and the spatial structure but
60 not stochastic transcription in dividing cells [31,32].

61 Here we present Biomodelling.jl 2.0, a rewrite of a Julia framework whose distinguishing feature, the coupling
62 of stochastic reaction kinetics to exponential growth, volume-scaled transcription and binomial partitioning at
63 division [33], is precisely the mechanism behind heritable expression states. Version 2.0 turns that feature into a
64 complete framework for the mechanistic study of drug tolerance and for single-cell benchmarking. We validate
65 the engine against exact stationary laws and an independent implementation, show that growth, division and
66 partitioning alone reproduce published results on cell-size scaling, partitioning noise, lineage-versus-population
67 statistics and cell-cycle-dependent bursting, and demonstrate that a single resistance gene with slow promoter
68 switching reproduces the quantitative lineage and memory signatures of heritable drug tolerance. We then
69 use the framework as a ground-truth generator for memory-gene detection, network inference, imputation and
70 perturbation prediction; as an inference engine that recovers kinetic parameters from snapshot and lineage data,
71 quantifies the bias of ignoring division, and maps the ridge to which drug parameters are determined by population
72 data; as a model calibrated to published time-lapse measurements and confronted with conditions it was not
73 fitted to; and, with a built-in schedule optimiser, as a testbed in which clinical dosing regimens are compared and
74 better ones are sought.

75 Results

76 Layers couple reaction kinetics to growth and division

77 Biomodelling.jl 2.0 is organised in layers (Fig. 1). A model is a list of reactions with mass-action, Hill or custom
78 kinetics and named parameters; promoter states are explicit species, grouped so that they are inherited at division
79 and doubled at replication. Four kernels advance a cell’s state at fixed volume: the exact direct method, fixed-step
80 tau-leaping, a hybrid scheme that falls back to the exact method whenever a leap would produce a negative
81 count, and adaptive tau-leaping with critical-reaction handling. The population layer grows each cell’s volume
82 exponentially, scales propensities with volume so that concentrations are invariant, replicates genes at a set point
83 of the cycle, divides cells by a sizer, adder or timer rule with binomial or over-dispersed partitioning, and controls
84 the population size by Moran-like replacement, free growth or a logistic hazard. A perturbation layer applies dose
85 schedules, state-dependent death, growth inhibition, drug-induced rate changes and genetic perturbations. Every
86 cell’s ancestry, fate and state at birth and division enter a lineage table, from which heritability, autocorrelation,
87 memory timescales, fluctuation-test and clonal-memory statistics follow. Observation models convert true counts
88 into scRNA-seq, smFISH or time-lapse measurements, and the same forward model drives likelihood-free inference.
89 Table 1 sets these capabilities against those of the simulators most used to benchmark single-cell methods.

Table 1. Capabilities of single-cell simulators (✓ native; ◦ partial or phenomenological; – absent). SERGIO [20], dyngen [21], scMultiSim [22], GRouNdGAN [23], TedSim [24], Cassiopeia [25], ABM.jl (AgentBasedModeling.jl) [28], Biomodelling.jl v1.x [33] and v2.0 (this work). Entries were read from each tool’s publication and documentation, and describe what the tool provides rather than what its formalism could express. Three readings need qualifying. SERGIO is a Python package that returns expression matrices as arrays and scMultiSim is an R package returning SingleCellExperiment-style objects, so neither writes AnnData natively; scMultiSim does carry a differentiation tree, so cells are related by a trajectory even though per-cell division and fate are not simulated. ABM.jl is the closest comparator: it simulates stochastic reaction networks in growing, dividing agents with an adder division rule, binomial partitioning, death and environmental coupling, and its companion work infers kinetic and selection parameters from mother-machine and lineage-tree data [29]; its ◦ entries are for gene replication, which its agent formalism can express but does not provide, and for lineage output, which traces a cell back to its ancestor rather than recording a tree with per-cell fates. Cassiopeia simulates division topologies with heritable fitness and CRISPR barcodes rather than molecular content, and its inference recovers phylogenies rather than kinetic parameters.

Capability	SERGIO	dyngen	scMulti-Sim	GRouNd-GAN	TedSim	Cassio-peia	ABM.jl	v1.x	v2.0
Mechanistic reaction kinetics (SSA / tau-leaping)	◦	✓	◦	–	–	–	✓	✓	✓
Gene regulatory networks with known ground truth	✓	✓	✓	✓	–	–	✓	✓	✓
Cell growth, volume-scaled transcription	–	–	–	–	–	–	✓	✓	✓
Gene replication / cell- cycle copy number	–	–	–	–	–	–	◦	–	✓
Division with molecular partitioning	–	–	–	–	◦	◦	✓	✓	✓
Heritable promoter states, expression memory	–	–	–	–	◦	–	✓	◦	✓
Lineage tree with per-cell fates	–	–	◦	–	✓	✓	◦	–	✓
Drug response: dose schedules, state-dependent death	–	–	–	–	–	–	–	–	✓
Genetic perturbations (knockdown / overexpression)	✓	✓	◦	✓	–	–	◦	–	✓
scRNA-seq observation model	✓	✓	✓	✓	✓	–	–	✓	✓
smFISH / time-lapse observation	–	–	–	–	–	–	–	–	✓
Parameter inference from data	–	–	–	–	–	◦	✓	–	✓
AnnData output	–	✓	–	✓	◦	◦	–	–	✓

The engine matches exact laws and an independent package

We first verified the kernels against exact stationary laws (Fig. 2). For a birth-death process the simulated distribution of 50,000 cells is indistinguishable from the Poisson law (Fig. 2a); for the telegraph model in an intermediate and in a bursty regime it matches the Beta-Poisson distribution [34] (Fig. 2b); and for bursty protein production it matches the negative binomial (Fig. 2c). Across kernels, Kolmogorov-Smirnov distances to the exact law with 20,000 cells lie below the 99% critical value for the direct method, for the adaptive scheme at both tolerances and for the hybrid scheme at the shorter step, and above it for the hybrid scheme at the longer step (0.0160 and 0.0122 against a critical value of 0.0115; Fig. 2d). That is the expected behaviour of an approximate kernel, and it sets the step below which the approximation is not detectable at this sample size; wall-clock costs

are 0.25 s to 0.91 s per 10^4 cells. Simulating the same telegraph model with JumpProcesses.jl gives distributions whose two-sample Kolmogorov-Smirnov statistic against Biomodelling.jl is 0.0034 (99% critical value 0.0163; Fig. 2e), with a 6.2-fold (in favour of JumpProcesses.jl) difference in serial wall-clock time. Those tests exercise one cell at fixed volume. The layer above them, growth, replication, partitioning and the difference between a lineage and a population, has exact solutions of its own for this class of model, which take different closed forms in the two settings [35–37]. We reproduced four analytically solvable members of that class, with constitutive, bursty and two-state promoter kinetics under a memoryless interdivision time and with volume-scaled synthesis from a replicating gene under a deterministic cycle, in both modes. Along single lineages, where the mother-machine control makes the recorded cells independent and the Kolmogorov-Smirnov test applies as it stands, all 4 samples match the exact law of their mode (distances 0.0023 to 0.0054 against 99% critical values of 0.0081 to 0.0115). A population snapshot is not an independent sample, because its cells share ancestors, so the independent critical value understates the true one and the population mode is judged on the agreement of its replicate means with the exact mean. Those means sit within 2.5 standard errors of the exact values (18 of the 20 individual samples also fall below the independent critical value, the largest distance being 0.0170). All four population means are nonetheless slightly low, by 0.03% to 0.70%; no single deviation is resolved, the largest being 2.5 standard errors ($p = 0.07$, Student’s t), but the common sign is not accounted for by the discretisation, the subsampling cap or the burn-in, each of which is excluded in Supplementary Note 14. Scored against the exact law of the other mode, the same samples give distances at least 6 times larger, so the comparison has ample power to tell the two settings apart. The two differ because a snapshot of a growing population over-weights cells that have just divided and so just lost half their molecules: the mean molecule number is constitutive production 13.33 against 10.00; bursts of four molecules 13.33 against 10.00; two-state promoter 8.89 against 6.67; volume-scaled synthesis with gene replication 25.64 against 23.96 along a lineage and in a population respectively (Supplementary Note 14 and Supplementary Fig. 6). Population simulations scale linearly with the number of cells: 2.6 s for 5000 cells $\times$ 10 telegraph genes $\times$ 200 steps (hybrid kernel) with four threads (Fig. 2f and Supplementary Table 3). Scaling in the number of reactions depends on the kernel and is close to linear only for the hybrid one (0.58 s to 6.4 s for the hybrid kernel, a factor of 11 for the same factor of 10 in genes; 1.04 s to 31.4 s for the direct kernel, a factor of 30 for the same factor of 10 in genes), because the direct method draws one event at a time while the hybrid kernel leaps over the non-critical channels in a fixed number of operations.

Physiology generates memory, scaling and cycle effects

The physiology of the population layer suffices to produce properties of single-cell data that phenomenological simulators must impose by hand (Fig. 3). In a single lineage of a slowly switching gene, mRNA and protein track volume within a cycle and are halved at division, while the promoter state, inherited by both daughters, persists over many divisions (Fig. 3a). For a fast-switching gene without replication, transcript numbers scale linearly with cell volume, so that concentration is homeostatic; adding mid-cycle gene replication doubles transcription in the second half of the cycle and lifts counts of large cells above the volume-scaling line (Fig. 3b), the copy-number effect that mechanistic fits to scRNA-seq data must account for [10]. Correlations of protein concentration at division between mothers and daughters, sisters and cousins fall from 0.94 to 0.02 as the promoter switching rate crosses the inverse cell-cycle time (Fig. 3c), mother-daughter and sister pairs correlating similarly at slow switching and mother-daughter pairs more strongly as switching quickens, cousins least, and mRNA less heritable than protein. The memory timescale extracted from the generation-resolved lineage autocorrelation is set by the promoter switching time and is nearly independent of the cell-cycle time as long as memory outlasts the cycle; when the two are comparable the estimate is bounded below by one cell cycle, the resolution of the autocorrelation (Fig. 3d). As reported for real lineages [11], the coefficient of variation of a stable protein’s concentration is larger

across the population than along single lineages, whether or not cell-cycle times are noisy, and over-dispersed partitioning adds variability in both perspectives (Fig. 3e). Telegraph parameters fitted separately to cells before and after replication move in the direction gene replication implies, but not by the factor it alone would give (Fig. 3f). Between the one-copy and the two-copy subsets the fitted transcription rate rises 2.0-fold, of which 1.4-fold is the larger volume of cells that have already replicated, leaving 1.4-fold per unit volume rather than the twofold that doubling the copy number implies; the fitted activation rate rises 1.4-fold over the same comparison, and the pooled fit that ignores the cell cycle returns a transcription rate above both subsets. Part of that shortfall is misspecification rather than copy number, and the figure shows it: two independently switching alleles are not one telegraph gene, and neither subset fit recovers the truth (0.5 and 1.5 for the switching rates and 2 for the transcription rate as Fig. 3f scales it, that is the true 40 transcripts per unit time per unit volume per allele divided by twenty, against 0.59, 1.89 and 2.4 fitted to the one-copy cells and 0.65, 3.27 and 4.8 pooled, whose transcription rates are scaled the same way). What the panel demonstrates is that a cell-cycle-blind fit is biased, not how large the copy-number effect is.

One resistance gene reproduces tolerance signatures

To test whether heritable expression states suffice to reproduce the lineage and memory signatures of drug-tolerant persistence, we coupled a resistance gene with slow promoter switching (memory timescale 45 time units, cell-cycle time 20) to a death hazard that rises with dose and is suppressed by the concentration of the gene's protein product (Methods). Under continuous, pulsed and reduced-dose schedules the population first collapses as low-expressing cells die, then regrows from high-expressing survivors (Fig. 4a); kill curves are biphasic with a surviving fraction that falls with dose (Fig. 4b). Between doses 0.5 and 2 the net rate of population decline over the first 30 time units of treatment rises from 0.023 to 0.050 per time unit (2.1-fold), while the mean division time of cells born under drug changes by 0.0% and the mean time to death of killed cells moves only from 8.5 to 6.2 time units (Fig. 4c; the single-cell timing distributions are in Supplementary Fig. 1). This reproduces the observation that population decay depends strongly on dose even though single-cell division and death times barely do [6], because fate is biased by states that exist before the drug arrives. Unlike the growth-arrested persisters of the canonical definition [1], the survivors here divide at an essentially unchanged rate (Fig. 4c, inset), so the model reproduces the heritable-state selection signatures of tolerance rather than its proliferative-arrest hallmark.

Related cells accordingly share the fate of their lineages. Among cells born in the 25 time units before the drug, one whose sister founded a surviving lineage does so itself with probability 0.39, against a marginal probability of 0.08 (4.8-fold enrichment over 1830 sister pairs; 2.9-fold for cousins), whereas in a fast-switching control, in which survival is rarer, the 1.8- and 2.5-fold values the panel prints rest on one sister pair and two cousin pairs in which both lineages survived, and the concordance sits at the independent expectation (+0.000 above it, against +0.051 for the memory gene; Fig. 4d). When tolerance pre-exists, survivors descend from a minority of founder clones and the effective number of clones drops from 491 to 89; when tolerance is drug-induced, so that the drug raises the promoter activation rate of every cell, barcode diversity is largely preserved (165 of 492; Fig. 4e). This is the signature that distinguishes the two origins in barcoding experiments [2,6]; the follow-up method that traces resistant fates back to their drug-naïve precursors puts those precursors at an initial frequency of about 1 in 1,000 to 1 in 10,000 cells [38]. The melanoma case study below does not sit in that regime: it uses 1:200, the per-marker frequency at which rare cells express resistance genes before drug [2], which is 5- to 50-fold higher. That range is in any case not resolvable at the size of this study, since a draw of its 2,000 founders would contain no pre-resistant cell at all with probability 0.8 at 1 in 10,000. The consequence is a larger and less clonal pre-existing resistant pool, so that among the arms which do lose control it is lost earlier than the traced frequency would give; the ranking of the schedules follows from the growth of resistant cells with and without drug rather than from their

199 initial frequency, but it was not re-tested at a lower fraction.

200 Disrupting expression memory changes the outcome, and when it is disrupted matters. Accelerating promoter
201 switching 20-fold from two cell-cycle times before the drug does not reduce the number of surviving clones if the
202 acceleration stops when the drug arrives: it raises them, from 98 to 120 (mean of 5 seeds each, $P = 0.009$, Welch's
203 t test). Reshuffling spreads the resistant state over more lineages and moves the resistant fraction towards its
204 stationary value (0.09 to 0.11), so more clones carry a protected cell when the drug arrives, even though protection
205 itself is untouched. Maintaining it during exposure reduces the count to 7 clones (10 cells, against 792; Fig. 4f),
206 because resistant cells now leave the protected state within a fraction of the time they need to found a colony. In
207 the model, memory must be disrupted while the drug is present.

208 Two published experiments bear on this prediction and they point in opposite directions, so it is worth being
209 explicit about which perturbation each one performs. A CRISPR screen for modulators of single-cell variability
210 found that inhibiting DOT1L *before* adding a BRAF inhibitor produced more therapy resistance than giving
211 the two together [39], which is the ordering the model predicts. A lineage-tracing study that identified PI3K
212 and TGF- β signalling as switching modulators found the opposite for its own pretreatment: a PI3K inhibitor
213 applied before targeted therapy left less resistance than targeted therapy alone [9]. The two results are not in
214 conflict with each other, and only the first is a test of what is simulated here. The perturbation in the model is
215 symmetric: it multiplies the activation and the inactivation rate by the same factor, so it accelerates relaxation
216 towards the stationary resistant fraction without moving that fraction, and its only effect before the drug is to
217 spread the resistant state over more lineages. A modulator that shifts the balance between the states is a different
218 perturbation, and a pretreatment that drives cells out of the primed state lowers the number of protected cells at
219 the moment the drug arrives, which reduces resistance whether or not it is maintained. The prediction here is
220 therefore specific to perturbations that speed switching without biasing it, of which DOT1L inhibition appears to
221 be one and PI3K inhibition is not; distinguishing the two experimentally requires measuring the resistant fraction
222 as well as the switching rate, which memory-sequencing and lineage-tracing assays do. Compartment theory for
223 phenotypic switching already predicts that an agent which perturbs the retention of the resistant state can turn
224 treatment failure into success [40]; what the simulations add is the timing requirement and a lineage-resolved
225 readout, the number of surviving clones, by which it can be tested. A transient pulse enriches high-expressing
226 cells and, with slow switching, the enrichment persists for many generations after the drug is removed, whereas
227 with fast switching it relaxes within a few cycles (Fig. 4g): phenotypic selection with stable inheritance, as in the
228 original MGMT model [41]. Finally, scanning release periods and doses of intermittent schedules (Fig. 4h), the
229 most effective schedule is release period 0, dose 2 (growth rate +0.0038 per time unit, ± 0.0007 over 5 seeds) for
230 pre-existing tolerance, release period 0, dose 2 (growth rate -0.0151 per time unit, ± 0.0012 over 5 seeds) when
231 the resistant state carries a fitness cost, and release period 0, dose 1 (growth rate +0.0089 per time unit, ± 0.0001
232 over 5 seeds) for drug-induced tolerance. Each point of the scan is the mean of 5 independent seeds, and every
233 schedule is run on the same 5 seeds, so schedules are compared by their paired differences, in which the shared
234 founder population cancels; the spread quoted beside each winner above is the seed-to-seed scatter of its absolute
235 growth rate, which is larger than the paired differences and is not the error on them. The winning schedule beats
236 the best schedule with a non-zero release period by 0.0003 to 0.0014 per time unit against paired standard errors
237 of 0.0001 to 0.0003, a margin that a two-sided paired t-test resolves at the 5% level in 1 of the 3 mechanisms, the
238 one in which the resistant state carries a fitness cost. Elsewhere the scan identifies a region of good schedules
239 rather than a single best one (pre-existing tolerance +0.0004 against 0.0002, $P = 0.18$; drug-induced tolerance
240 +0.0003 against 0.0001, $P = 0.11$). Without a fitness cost nothing beats continuous dosing at the highest dose,
241 since every release period lets the resistant subpopulation regrow (the drug does not act on the switching rates in
242 this model; when it also stabilises the resistant state, as in the melanoma case study below, holidays regain a

small advantage even without a cost). With a fitness cost the population declines under every schedule at dose $\geq$ 0.5, and release periods of up to 10 time units at doses of 1 and 2 cost little, though 3 of those 4 are resolvably worse than continuous dosing at the highest dose (+0.0014 to +0.0029 per time unit, -0.0151 per time unit under continuous dose 2), because cells that leave the resistant state during a holiday are killed on re-exposure. For drug-induced tolerance the ranking of doses reverses: the highest dose induces tolerance in more cells than it kills, so dose 1 outperforms dose 2 at every release period (release period 0, dose 1 (growth rate +0.0089 per time unit, $\pm$ 0.0001 over 5 seeds), against +0.0123 for continuous dose 2), the intermediate-dose optimum reported previously [42]; the benefit of release periods reported in that work does not appear here, because in this model the induced state decays at the same rate with and without drug, so a holiday only postpones the death of reverting cells while all cells grow.

The hazard above does not depend on where a cell is in its division cycle, and the calibration below will show that fate data cannot exclude one that does, so we repeated these simulations with three quarters of the hazard confined to a window around mid-cycle. Population decay still rises with dose (-0.001 to +0.023 per time unit over doses 0.5 to 2, against +0.023 to +0.050 without the gate) and the mean time to death stays nearly dose-invariant (8.4 to 9.9 against 6.2 to 8.5 time units), so neither signature depends on a cycle-blind hazard. Sisters share a birth time and so a cycle phase, which lets a gated hazard correlate their fates with no inherited state; it does not do so to any useful degree, lifting the sister concordance of the memoryless fast-switching control only to +0.008 above independence (from +0.000), against +0.054 for the memory gene, so the kin signature still reads expression memory rather than the cycle. In the schedule scan, continuous dosing wins in all 3 models with the gate as without it, at the same dose in 3 of 3. The gate raises the long-term growth rate (+0.0026 to +0.0155 per time unit at all 60 points of the scan), which is what confining three quarters of the hazard to a window of the cycle implies, and it perturbs the order of the schedules below the winner (Spearman correlation 0.977 to 0.998 between the gated and cycle-blind orderings), but it changes neither which schedule wins nor the intermediate-dose optimum under drug-induced tolerance (+0.0156 against +0.0190 under the gate and +0.0089 against +0.0123 without it). Across the two scans a two-sided paired t-test resolves the margin over the best holiday at the 5% level in 4 of the 6 comparisons, the exceptions being pre-existing tolerance without the gate (+0.0004 against a paired standard error of 0.0002) and drug-induced tolerance without the gate (+0.0003 against a paired standard error of 0.0001) (Supplementary Note 13 and Supplementary Fig. 5).

Simulated ground truth benchmarks inference methods

Because every cell's clone, lineage and true state are known, the framework supplies ground truth that observational data cannot (Fig. 5). For 40 independent genes whose memory times span three decades, MemorySeq-style clonal variance scores computed on true concentrations rise from 1.0 for genes with memory far below the cell-cycle time to a mean of 29.6 for genes with memory beyond a hundred time units, five cell-cycle times, the highest single gene reaching 44.5 (38.4 on sequenced counts), and the ranking survives computation on sequenced counts at 20% capture efficiency (Fig. 5a); the score separates genes with memory longer than the cycle from the rest with an area under the receiver operating characteristic (ROC) curve of 1.00 on true concentrations and 0.99 on sequenced counts. For a 30-gene network, small enough that the methods were run with their default settings and the ranking of datasets rather than the absolute accuracy is what should be read, simulated at fixed volume, in a dividing population and after sequencing, correlation-based and GENIE3 inference lose most of their accuracy in the dividing population when raw counts are used (area under the precision-recall curve 0.52 to 0.14 for Pearson correlation and 0.32 to 0.13 for GENIE3, against 0.10 for a random ranking of undirected pairs and 0.05 for directed edges), because volume and gene copy number induce positive correlations between all genes. Normalising counts to concentrations removes the volume effect but recovers little (0.18 and 0.18), since the

dependence on copy number and cell age remains; regressing expression on the true cell-cycle covariates, which the simulation provides and which cell-cycle scoring approximates in real data, recovers most of the loss (0.46 and 0.30). Sequencing at 15% capture efficiency lowers accuracy further (0.11 raw and 0.24 library-size-normalised for Pearson correlation; 0.07 and 0.08 for GENIE3; Fig. 5b). One comparison in that panel runs the wrong way: for Pearson correlation the library-size-normalised sequenced counts score above the true concentrations (0.24 against 0.18). Dividing counts by the library size removes cell-to-cell scale, and in a dividing population that scale is dominated by volume and gene copy number, so the normalisation acts as a coarse correction for the growth confound that dividing by the true volume does not make. It is not evidence that sequencing helps; the fixed-volume ceiling (0.52) is above both, and the cell-cycle regression (0.46) is above both as well. Imputing the sequenced counts before inference recovers part of that loss: MAGIC raises the area under the precision-recall curve from 0.11 to 0.32 for Pearson correlation and from 0.07 to 0.18 for GENIE3, and kNN-smoothing gives 0.16 and 0.12 (Fig. 5c). Neither reaches the accuracy obtained at fixed volume, nor that of regressing out the cell-cycle covariates, so imputation mitigates the sequencing loss without touching the growth confound that produces most of it [33]. Knocking down the five most connected regulators of a 20-gene network, again small and run at default settings, gives true fold changes that a no-change baseline predicts with a mean coefficient of determination of -0.16. Propagating observational correlations beats that baseline for 4 of 5 knockdowns, with coefficients of determination from 0.09 to 0.43, and fails on the remaining one (-2.39), so its mean over the five is -0.23 against a median of 0.30 (Fig. 5d). A predictor that is useful for most perturbations and badly wrong for one is precisely what an out-of-distribution evaluation has to expose, and exposing it needs the causal ground truth that the simulation supplies [13,14].

Inference recovers parameters and exposes division bias

On a snapshot of 2,000 non-dividing cells, ABC-SMC with log-uniform priors over three decades and the Wasserstein distance between simulated and observed counts [43] concentrates on the true telegraph parameters and agrees with the exact maximum-likelihood estimate (posterior medians (0.31, 0.60, 20.9) for truth 0.3, 0.6, 20; Fig. 6a), the broadest posterior falling on the promoter inactivation rate, which a single snapshot constrains least. On a snapshot of a growing and dividing population with volume-scaled transcription and mid-cycle replication, the naive telegraph fit is biased ($k_{\text{on}} = 0.47$, $k_{\text{off}} = 1.95$, $k_{\text{tx}} = 71.5$) because volume and copy number inflate counts and their dispersion, whereas ABC with the population simulator as forward model recovers the truth ((0.30, 0.68, 20.8); Fig. 6b). Drug parameters follow from population data in the same way: from a kill curve and the concordance of sister fates, the posterior over the maximal death rate and the protection threshold is a ridge (weighted correlation of the two log parameters -0.85) that contains the truth (Fig. 6c), and posterior predictive kill curves simulated at the posterior median reproduce the observed one (Fig. 6d). The two summaries constrain the product of death rate and unprotected fraction, so separating the parameters requires measuring the protective protein: an experimental-design question the forward model answers before an experiment is run (the tolerance and acceptance-rate schedules of all three ABC runs are in Supplementary Fig. 2).

Calibration to time-lapse data is tested on held-out data

A simulator is useful for treatment questions only if it can be fitted to the data a laboratory produces and then predict what it was not fitted to. We therefore calibrated the resistance-gene model to time-lapse measurements in which U2OS cells were tracked for three days under 7, 10 and 13 μM cisplatin and every cell present at drug addition was scored as dying, dividing or surviving without dividing [6] (Supplementary Table 4). Six parameters, the two promoter switching rates, the maximal death rate, the EC_{50} and Hill coefficient of death and the IC_{50} of growth arrest, were fitted to the fate fractions at 7 and 13 μM by simulated moments with common random

328 numbers, under a weak prior keeping the memory of the resistant state near the two to three generations that the
 329 same study inferred from lineage correlations (Methods). The calibrated model reproduces the training fractions
 330 with a root-mean-square error of 0.077 (memory 2.8 generations, resistant fraction 15.2%, EC_{50} 16.7 μ M; Fig. 7a,f).
 331 Three related numbers appear in this Article and are not interchangeable. Every root-mean-square error quoted in
 332 the text is the residual of the seed-averaged fate fractions pooled over every fraction of the concentrations named:
 333 four seeds per condition for the calibration, here and at the held-out concentration, and three for the comparison
 334 of the two killing mechanisms below, whose training errors alone are single-seed, because the optimiser evaluates
 335 one seed per candidate. The colour scale of Fig. 7e is the quantity the optimiser minimised, which is evaluated on
 336 a single seed and adds the memory prior: at the optimum it is 0.067, made up of 0.061 of single-seed training error
 337 and 0.006 of prior penalty. The comparison of the two killing mechanisms in Supplementary Note 12 uses that
 338 single-seed training error for both models, so that they are scored by the same rule as the fit itself. That error
 339 should be read against the precision of the measurement: the observed fractions carry a mean binomial standard
 340 error of 0.028 and repeated simulations differ by 0.015, so the fit is about 2.4 times the combined noise and the
 341 model does not reproduce the training data within error. At the held-out concentration the root-mean-square
 342 error is 0.093, and the discrepancy is specific rather than diffuse: the model matches the fraction surviving
 343 without dividing but over-predicts death and under-predicts division (observed died, divided and survived 0.28,
 344 0.31 and 0.41; predicted 0.39, 0.20 and 0.41; Fig. 7a), so it captures the ordering of the three concentrations while
 345 misplacing the split between death and division at the intermediate one.

346 Six parameters fitted to six numbers are not all determined by them. Moving one parameter at a time from
 347 the fitted values, with the others held there, the fate fractions constrain the death parameters to within about
 348 a factor of 2 but leave the two switching rates free over roughly a factor of 6 (Supplementary Fig. 3a), which
 349 is why the memory of the resistant state was given a prior rather than left to the fate data. Holding the death
 350 parameters at their fitted values, memories from 1.0 to 9.0 generations paired with resistant fractions from 1%
 351 to 25% all describe the fate fractions within the noise of the measurement (Supplementary Fig. 3b), because a
 352 shorter memory with more resistant cells and a longer memory with fewer are hard to tell apart from fates alone.
 353 The parameters quoted above should therefore be read as one member of a family of fits that describe these fate
 354 fractions equally well, and the biological reading of the memory rests on the lineage correlations that set the prior,
 355 not on the fate fractions. Halving the integration step changes the simulated fractions by 0.005, less than the
 356 seed-to-seed spread, so the discretisation is not a factor (Supplementary Fig. 3c).

357 A second mechanism can produce the same fates. Cisplatin kills preferentially during replication, so where a cell
 358 sits in its division cycle when the drug arrives is a competing explanation both for the dose-invariance of death
 359 timing and for the cells that survive three days without dividing: a hazard confined to part of the cycle spares
 360 whatever never enters it, with no heritable state at all. We therefore refitted the model with a cell-cycle-dependent
 361 hazard, adding one parameter, the fraction of the hazard that is independent of cycle position, with killing
 362 otherwise peaking at mid-cycle, under the same objective, bounds, budget and seeds (Methods). The two fits
 363 describe the training fractions equally well (0.064 against 0.061, against a mean sampling error of 0.028 in the
 364 measurement), and the cycle-dependent fit does not improve the held-out prediction either (0.106 against 0.095).
 365 The fitted cycle-independent fraction is 0.34, but values from 0.20 to 0.80 all sit within one sampling error of the
 366 best fit, so the fate fractions do not determine it (Supplementary Note 12 and Supplementary Fig. 4). Three fate
 367 counts per concentration cannot separate inherited protection from cell-cycle position; the measurement that can
 368 is the cycle position of each cell at the moment the drug arrives, which a cycle reporter records directly.

369 Properties that were not fitted can then be compared with the data. The predicted kill curves (Fig. 7b) are
 370 1.5-fold steeper at 13 μ M in their first two days than in their second half, but they do not plateau, as HCT116

cells do after about 100 h. Death fates are correlated between sisters (0.09), first cousins (0.06) and second cousins (0.05), and the correlation decays to 0.02 for third cousins (95% confidence interval -0.026 to +0.069, Student's *t* on two degrees of freedom over three seeds, a range that contains zero) and 0.00 for unrelated cells (Fig. 7c). That is the pattern reported for HCT116 lineages, in which fate correlations are present for first and second cousins and absent for third cousins, with memory inherited over at least two to three generations [6]; the comparison is qualitative on both sides, because the source reports these correlations graphically and gives no numeric coefficients or confidence intervals to match. One caveat belongs beside this agreement. Cells here draw their cycle times independently, so no cycle timing is inherited and every correlation between related cells is carried by the inherited promoter state. Cell-cycle duration is in fact heritable over several generations, and its inheritance alone produces strong cousin and second-cousin correlations even when mother-daughter correlations are weak [44–46]. Coupled to a cycle-gated hazard, that is a complete alternative explanation for the correlations fitted here, and the model as built cannot produce it. The mean time to death changes by 8% between 7 and 13 μ M while the death fraction rises from 0.13 to 0.59 (Fig. 7d), the dose-invariance of single-cell timing that motivated the original study, and it arises here because dose changes the fraction of unprotected cells rather than their kinetics. The plateau the kill curves miss is the clearest failure of the minimal model: reproducing it requires cells to stay in the resistant state for longer than the memory of 2.8 generations allows, and that memory was set by the prior taken from the lineage correlations rather than by the kill curves. The two observations therefore pull in opposite directions within this minimal model: correlations that vanish by the third cousin call for a short memory, and a kill curve that flattens calls for a long one. A single switching gene cannot satisfy both, which is a concrete and testable limitation, and points to a drug-induced stabilisation of the resistant state, for which there is independent evidence [2], as the mechanism a lineage-tracing experiment should look for. The calibration landscape (Fig. 7e) shows how the fitted death rate and EC_{50} are constrained by the fate fractions alone; the protocol applies to any observable a laboratory records.

Schedule optimisation is consistent with clinical outcomes

With a calibrated model the framework becomes a schedule testbed. We first asked which schedule best delays regrowth of a melanoma-like population with a rare, heritable pre-resistant state [2] treated at a fractional-killing dose, comparing continuous dosing, the intermittent schedule of the SWOG S1320 trial, in which patients received a continuous 8-week lead-in and were then randomised to continuous dosing or to cycles of 3 weeks off and 5 weeks on [47], and adaptive therapy, in which the drug is withdrawn when the population has halved and resumed when it regrows to its initial size [48]. Because adaptive therapy deliberately lets the tumour regrow, schedules are compared by the time after randomisation at which the population exceeds 120% of its pre-treatment size (loss of control), with the RECIST-like time to progression from the nadir reported alongside (Methods).

The answer depends on two properties of the resistant cells (Fig. 8a–d). When the resistant state is fully protective and carries no fitness cost, as lineage data indicate for pre-existing tolerance [6], resistant cells grow at the same rate with or without drug, so a holiday cannot cost the resistant subpopulation anything and can only let sensitive cells regrow to be killed again. Even so, both interrupted schedules keep control longer than continuous dosing: the three lose control after 23, 26 and 27 weeks, with the intermittent schedule delivering 65% of the cumulative dose of continuous dosing over the full 60-week simulation. The reason is a second drug effect rather than the fitness of resistant cells: the drug stabilises the resistant state tenfold by suppressing the promoter inactivation rate (Methods), so continuous exposure keeps cells resistant that a holiday allows to revert. That effect is set in the model, not fitted, and it is what makes holidays worth anything at all in this arm. When resistant cells pay a 50% fitness cost in the absence of drug, a coarse stand-in for the drug dependence that made intermittent dosing superior in vemurafenib-treated patient-derived xenografts, where withdrawal drove established resistant tumours

414 to regress rather than merely grow more slowly [49], holidays let sensitive cells outcompete them: none of the three
 415 schedules loses control within the 52-week follow-up (52+, 52+ and 52+ weeks, censored), and adaptive therapy
 416 achieves this with 36% of the cumulative dose of continuous dosing. When resistant cells are only partly protected
 417 and grow at half speed under drug, holidays release them: continuous dosing keeps control for 52+ weeks against
 418 49 for the intermittent schedule and 43 for adaptive therapy, the ranking observed in the trial, where continuous
 419 dosing gave a median progression-free survival of 9.0 months against 5.5 months for intermittent dosing. Gating
 420 three quarters of the death hazard to a window of the cell cycle, as in the persister simulations above, leaves
 421 these comparisons standing, whether or not the maximal hazard is raised to hold the cycle-averaged hazard fixed
 422 (Supplementary Note 13 and Supplementary Fig. 5): the best schedule is unchanged in 3 of 3 mechanisms and the
 423 model still favours continuous dosing under partial protection (52 against 43 weeks, from 52 and 49). The memory
 424 of the resistant state, five net population doublings, is an assumption rather than a measurement, and it is what
 425 decides whether resistant cells revert during a three-week holiday; scanning it from two to twelve doublings leaves
 426 the ranking unchanged and narrows the advantage of holidays from 9 weeks at 2 net doublings of memory to
 427 1.7 weeks at 12 (Supplementary Note 15). That the benefit of holidays should depend on a cost of resistance
 428 is established by compartment models of adaptive therapy, which also show that tumour turnover modulates
 429 how large that cost must be [50]. Here the dependence emerges from single cells whose resistance is a heritable
 430 expression state rather than a fixed compartment, the cost and the degree of protection are properties of cells
 431 that lineage experiments measure, and the third case shows that partial protection reverses the ranking even
 432 when a cost is absent. The optimiser makes the dependence explicit (Fig. 8e): over intermittent schedules of
 433 any period and duty cycle after the lead-in, the best schedule is a period of 2.2 weeks with 39% of the time on
 434 drug (27 weeks to loss of control) without a fitness cost, a period of 1.0 weeks with 47% of the time on drug (52
 435 weeks to loss of control) with a fitness cost and a period of 1.0 weeks with 100% of the time on drug (52 weeks to
 436 loss of control) with partial protection. A bacterial-persistence model adapted to colorectal-cancer lines reaches
 437 the same conclusion by another route: an intermediate dose with release periods is optimal, because the rate of
 438 transition into the persistent state rises with concentration and then saturates [42]. The agreement is partial and
 439 the difference is informative. Where tolerance is drug-induced, the optimiser here also selects an intermediate
 440 dose over the highest one, for the same reason. Release periods, however, help there only when the drug changes
 441 the kinetics of the resistant state, either by stabilising it, as in the melanoma arm above, or by inducing it at
 442 a rate that saturates, as in that model; with a state whose decay rate is the same with and without drug, a
 443 holiday merely postpones the death of reverting cells while every cell grows, and the optimiser rejects it. The
 444 model therefore reproduces both the clinical and the preclinical outcome and shows their difference to depend on
 445 measurable properties of the resistant cells, their growth under drug and their fitness without it, both properties
 446 of cells rather than of the tumour, which growth and lineage experiments can bound, within the identifiability
 447 limits of Fig. 6c, before a schedule is chosen. The mechanism is set in each case, not inferred from the outcome, so
 448 this is a demonstration that the ranking follows from cell properties rather than a claim that these particular cells
 449 carry them; and because the trial and the xenograft also differ in drug and species (dabrafenib plus trametinib in
 450 patients, vemurafenib monotherapy in mice), their contrast is not attributable to resistant-cell fitness alone.

451 The second case study concerns glioblastoma, where the RTOG 0525 trial compared the standard temozolomide
 452 regimen (150–200 mg/m² on days 1–5 of 28-day cycles) with a dose-dense regimen (75–100 mg/m² on days 1–21)
 453 that delivers about twice the cumulative dose and was designed to deplete MGMT, and found no difference in
 454 survival for either MGMT-methylated or unmethylated tumours [51]. We simulated both regimens with one-
 455 compartment pharmacokinetics (elimination half-life 2.1 h, a plasma value from a three-compartment population
 456 model of plasma and cerebrospinal fluid [52]; an independent single-dose study gives 1.8 h [53]) on populations in
 457 which 1% (methylated) or 30% (unmethylated) of cells express MGMT. Those two fractions are illustrative values

458 that bracket low and high expression, not measurements, and the idealisation matters: promoter methylation
 459 shifts MGMT expression as a graded quantity across cells rather than switching a discrete subpopulation on,
 460 so the whole dose-dense prediction below rests on treating it as two states. With that caveat, with MGMT
 461 protein protecting cells from the drug [41] and either stable or consumed by it stoichiometrically, one molecule
 462 per lesion repaired, as a suicide enzyme is, so that the pool follows the cumulative exposure rather than the peak
 463 concentration (Fig. 8f,g). Lesions are not tracked as a state variable: they are taken to form in proportion to
 464 dose and to be repaired at once, so “one molecule per lesion” enters only through the rate of consumption, and no
 465 unrepaired damage accumulates. The rate constant of that consumption is not fitted to data; it is set so that
 466 consumption at the standard bolus peak matches the first-order parameterisation it replaces, which makes the two
 467 mechanisms comparable at the same initial rate and isolates how each integrates over time. With stable MGMT,
 468 the depth of response over six cycles is $0.20 \log_{10}$ units for the standard regimen and 1.00 for the dose-dense
 469 regimen in the methylated population (0.08 and 0.10 in the unmethylated one), and a 21-day regimen at the
 470 standard cumulative dose achieves no measurable kill ($0.00 \log_{10}$ units): the death hazard rises steeply with
 471 concentration, so spreading the same dose over more days loses efficacy, and the advantage of the dose-dense arm
 472 in the model is entirely its larger cumulative dose. With stable MGMT the two arms therefore differ little in the
 473 unmethylated population (0.08 against $0.10 \log_{10}$ units), as in the trial. Consumption changes that. Because one
 474 molecule is spent per lesion repaired, the pool follows the cumulative exposure, which the dose-dense regimen
 475 more than doubles: in the methylated population its response deepens to $1.52 \log_{10}$ units against 0.22 for the
 476 standard regimen, and in the unmethylated population, where 30% of cells start protected, the same depletion
 477 turns a regimen that barely worked into one that controls the population (1.08 against $0.18 \log_{10}$ units; final size
 478 0.1 against 1.1 times the initial size, and 4.9 and 3.6 with stable MGMT). Modelled as the suicide enzyme it
 479 is, MGMT therefore predicts a clear dose-dense advantage in both populations, which RTOG 0525 did not find.
 480 Read as a bound rather than as a fit, that disagreement is informative: either temozolomide consumes MGMT in
 481 tumours far more slowly than this stoichiometry allows, or the benefit of depleting it is offset by something the
 482 model leaves out. At the standard cumulative dose the number of dosing days per cycle that minimises the final
 483 population is 9 (methylated) and 9 (unmethylated) with stable MGMT, and 9 and 5 when MGMT is consumed
 484 (Fig. 8h). These simulations are not fits to the trial, whose end points depend on much that the model omits,
 485 dose-limiting toxicity and the dose reductions it forces among it. What they add is a sharper version of the
 486 rationale for dose-dense dosing: the benefit it is supposed to deliver exists in the model only at consumption rates
 487 that the trial’s null result appears to exclude. The quantity the argument rests on is the per-lesion molecular rate
 488 constant of that consumption, which has not been measured directly, although bulk depletion time-courses have:
 489 protracted temozolomide lowers alkyltransferase activity in peripheral blood mononuclear cells by about 63%
 490 at 14 days and 73% at 21 days [54], a compressed $1,000 \text{ mg/m}^2$ schedule takes it to 18% of initial levels while
 491 O^6 -methylguanine accumulates in both blood cells and tumour biopsies [55], and in orthotopic GBM43 xenografts
 492 tumour MGMT activity is depleted by day 6 on both the standard and the 21-day schedule and stays suppressed
 493 only on the 21-day one [56]. The last of these is the sharpest comparison available, and it makes the disagreement
 494 more specific than the trial alone does: the deeper and longer tumour depletion that the dose-dense schedule
 495 is supposed to deliver was confirmed in that model, and survival still did not improve. Either the depletion
 496 measured in bulk is not the per-cell depletion that matters for killing, or the benefit of depleting MGMT is offset
 497 by something the model leaves out.

Discussion

Biomodelling.jl 2.0 fills a gap in the single-cell toolkit: a simulator in which expression memory, cell-size effects, cell-cycle dependence and drug-tolerant persistence follow from explicit biophysics instead of being assumptions of a generative model. It is worth separating what is new here from what it confirms. No individual ingredient is new, and the claim is not that any of them is: stochastic kinetics in growing and dividing cells have been simulated before [27,28,33], lineage-tracing data have their own simulators [24,25], compartment models of adaptive therapy established that the benefit of treatment holidays depends on a cost of resistance and on turnover [48,50], and phenotypic-switching theory established that perturbing the retention of a resistant state can decide a treatment outcome [40]. What is new is that this is the first package to combine these mechanisms with joint inference and schedule optimisation in a single forward model. Version 1 of this framework already simulated stochastic gene networks in growing and dividing cells and generated synthetic scRNA-seq for imputation and network-inference benchmarking [33]; what version 2 adds to it is gene replication and cell-cycle copy number, division by size, adder or timer rules rather than a timer alone, inherited promoter states, the lineage table and its statistics, the drug and pharmacokinetic layer with state-dependent killing, likelihood-free inference, and the schedule optimiser. Relative to the closest external comparator, an agent-based framework that already couples stochastic kinetics to growth, adder division, partitioning and inference from lineage data [28,29], what is added is the drug and schedule layer, the observation models, gene replication, and the benchmark-generation use. Two consequences follow from the combination rather than from any part of it. It makes a resistance phenotype a heritable expression state with measurable kinetics instead of a compartment with an assumed rate, and it supplies ground truth of a kind observational data cannot: the same simulation gives known memory genes, known networks, known perturbation effects and the true lineage tree, which is what the quantification of how growth and sequencing destroy network inference, and of how far cell-cycle regression recovers it, depends on. It also makes the calibrate-then-predict protocol available for schedule questions, which turns a mechanism into a falsifiable prediction: here it identifies the growth of resistant cells without drug and their residual growth under drug as the two properties that decide whether holidays help, and points to the rate at which temozolomide consumes MGMT as the measurement that would separate the two glioblastoma regimens. That difference has practical consequences. As a hypothesis engine for non-genetic drug tolerance, the framework answers which schedule, which combination or which memory-modulating pretreatment minimises the long-term fitness of a treated population, with cell-level rules that can be matched to time-lapse and lineage-tracing data. As a ground-truth generator, it produces data with known memory genes, known networks, known perturbation effects and known lineage trees, which is what MemorySeq-style scores, lineage-aware models and virtual-cell benchmarks need in order to be evaluated fairly, and what existing simulators cannot supply. As a likelihood-free inference engine, it estimates mechanistic parameters from data that are now routinely collected, while quantifying the bias of simpler models.

The calibration and schedule studies set out a workflow for in silico single-cell pharmacology: declare the training and validation observables before fitting, fit by simulated moments or approximate Bayesian computation with common random numbers, validate on conditions and statistics that were not fitted, and only then compare regimens, reporting outcomes that map onto clinical end points. The case studies show why mechanism matters. The same question, continuous or intermittent dosing, has opposite answers depending on two properties of the resistant cells, their growth under drug and their fitness without it, and a trial result and a xenograft result that appeared to contradict each other are both reproduced once those properties are set. They are set in each case rather than inferred from the outcomes, so the case studies show that the ranking follows from cell properties, not that any particular tumour has them. Neither anchor is as one-sided as a single end point makes it look. In SWOG S1320 the progression-free survival advantage of continuous dosing (median 9.0 against 5.5 months,

hazard ratio 1.36 for intermittent against continuous, $P = 0.063$ against the trial’s pre-specified two-sided α of 0.2) did not carry through to overall survival, a secondary end point for which the trial was not powered and on which the two arms were equal at a median of 29.2 months [47]; and in the xenografts the resistant tumours were drug-addicted, regressing on withdrawal rather than merely growing more slowly [49], which is the extreme of the fitness cost that the second mechanism here represents. Because the framework simulates the lineage-resolved data from which such properties are measured, it can also say which experiment would settle the question before a trial is designed.

Several limitations delimit this release. Cells are well-mixed compartments without spatial structure or cell-cell signalling; the pharmacology is a one-compartment model; growth is exponential with phenomenological size control; and the inference layer, although general, relies on user-chosen summary statistics, so posterior calibration should be checked with predictive simulations. The persister case study is deliberately minimal, a single resistance gene with slow switching, and its calibration used fate fractions transcribed from a published study rather than raw tracks. Its death hazard was first written without cell-cycle dependence, which matters for a drug such as cisplatin whose lethality is enhanced during replication. Refitting with a cycle-dependent hazard shows that the two mechanisms are not separable by these data: a model in which killing is confined to the replication window accounts for the fate fractions as well as heritable protective expression does, and the fraction of the hazard that is cycle-independent is not determined by them at all. Cell-cycle position at the moment the drug arrives is therefore a live alternative to inherited protection, and the experiment that distinguishes them is a cycle reporter read at drug addition rather than more fate counts. Cells draw their interdivision times independently, so cell-cycle duration is not inherited, and every fate correlation between related cells in these simulations is carried by the inherited promoter state. That is a real restriction and not only a simplification. Cell-cycle duration is heritable over several generations in mammalian cells, and its inheritance alone produces strong correlations between cousins and second cousins even where mother-daughter correlations are weak [44–46]; combined with a cycle-gated hazard it is a complete competing explanation for the second-cousin fate correlations that the calibration in Fig. 7 reproduces, and the model as built cannot generate it, so the fit has no choice but to load all of that correlation onto expression memory. Chromatin states or cycle timing with their own inheritance rules would let the two be separated, and are the extension this points to. Six parameters fitted to six fractions leave the switching rates weakly determined, so the memory of the resistant state rests on the lineage correlations used as a prior rather than on the fates themselves; and that memory is carried by one slow promoter, a coarse proxy for chromatin-based inheritance whose single-exponential, symmetric form is a simplification of a process that can be multi-timescale and bistable. The clinical case studies use literature values for cell-cycle times, resistant fractions and pharmacokinetics rather than patient-specific measurements, so their conclusions concern mechanisms and rankings, not absolute survival times. In the glioblastoma study MGMT is consumed stoichiometrically, one molecule per lesion repaired, so its depletion follows the cumulative exposure as a suicide enzyme’s should. Two things about that model should be read with care. Lesions are not a state variable: they are assumed to form in proportion to dose and to be repaired immediately, so “one molecule per lesion” enters only through the consumption rate and no unrepaired damage accumulates. And the rate constant of that consumption was set so that the initial consumption rate at the standard bolus peak matches the first-order parameterisation it replaces, which makes the two mechanisms comparable but leaves the constant itself unmeasured; bulk depletion time-courses in blood cells and in tumour tissue exist [54–56] and bound it loosely, and that is why we name the per-lesion constant as the quantity that separates the regimens. MGMT is in any case only one determinant of a temozolomide response that is also mismatch-repair-dependent, and the 1% and 30% expressing fractions are illustrative values that bracket low and high expression rather than measured quantities: promoter methylation shifts expression as a graded quantity across cells rather than switching a discrete subpopulation on, and the

dose-dense prediction rests entirely on this two-state idealisation. The memory of the resistant state is expressed in these tumour-scale studies as a number of net population doublings, not of cell divisions, and so is not the division-counted memory of memory-sequencing experiments. A bridge to the Catalyst.jl ecosystem is not part of this release. The benchmark results should be read as demonstrations of what the framework makes possible rather than as verdicts on the methods compared, which were run with default settings on small networks.

Spatial coupling through neighbourhood-dependent rates, cell-cell communication, chromatin states with their own inheritance rules and neural simulation-based inference are natural extensions of the existing forward model; the spatial half of that list is the province of agent-based tumour simulators such as PhysiCell and PhysiBoSS [31,32], which model tissue mechanics and microenvironment but not stochastic transcription in dividing cells, so the two are complementary rather than competing. Because the framework is deterministic given a seed, fast enough to simulate tens of thousands of cells in minutes, and writes AnnData files, it can be dropped into single-cell analysis pipelines as a source of mechanistically grounded synthetic data, and into virtual-cell training and evaluation as a source of causal, out-of-distribution contexts.

Methods

Reaction model and propensities

A model is a set of species and reaction channels. Species are molecules (partitioned at division) or promoter states (inherited). For a reaction with reactant counts x_i , stoichiometric coefficients ν_i , rate constant k and volume exponent e , the mass-action propensity is

$$a = k V^e \prod_i \binom{x_i}{\nu_i},$$

where by default $e = 1 - \sum_i \nu_i$, so that zero-order production scales with volume, first-order reactions are volume independent and bimolecular reactions scale as $1/V$; this keeps concentrations invariant under growth. Regulated reactions multiply this by $b + (1 - b)F(\mathbf{c})$, where b is a basal fraction and F combines Hill terms of the regulator concentrations $c_j = x_j/V$, $h_j = c_j^{n_j}/(K_j^{n_j} + c_j^{n_j})$ for activators and $1 - h_j$ for inhibitors, either as a product (AND logic) or as $1 - \prod_j (1 - h_j)$ (OR logic). Promoter switching is represented by explicit promoter species (for instance $G_{\text{off}} \rightleftharpoons G_{\text{on}}$) grouped into promoter groups. Custom propensities can be supplied as functions of counts, volume, parameters and time.

Stochastic simulation kernels

Within one update step the cell volume and gene copy number are held fixed and the reaction network is advanced by one of four kernels: the direct method [57] with propensities updated through a species-reaction dependency graph; fixed-step Poisson tau-leaping [58]; a hybrid scheme in which a leap that would drive any species negative is discarded and that interval is simulated exactly instead (the “switching” scheme of version 1, with the discarded interval resolved exactly rather than clipped); and adaptive tau-leaping with the step-size selection of [59] and the critical-reaction partition of [60], falling back to a fixed number of exact events when the proposed leap is shorter than $10/a_0$. The hybrid scheme is approximate rather than exact: a leap is kept only when it leaves every species non-negative, so an accepted step follows the Poisson leap truncated to the non-negative region, and only the discarded intervals follow the exact process. The discrepancy falls with τ and is bounded by the Kolmogorov-Smirnov test of Fig. 2d, which the scheme passes at $\tau = 0.05$ and fails at $\tau = 0.2$ for the models tested; the direct method is used wherever exactness is required (Supplementary Table 1 lists the kernel correctness

622 tests). Every kernel takes an explicit random number generator.

623 Cells, growth, replication and division

624 Each cell carries its volume V , age, generation, lineage identifiers, gene copy state, growth rate, molecule counts,
 625 a cell-specific parameter vector and its own random number generator (seeded from the parent's generator, which
 626 makes results independent of the number of threads). Volume grows exponentially, $\dot{V} = \lambda V$, with an optional
 627 log-normal spread of λ between cells. Division is triggered by a sizer ($V \geq V_{\text{div}}$), an adder ($V - V_{\text{birth}} \geq \Delta$) or
 628 a timer ($\text{age} \geq T$), each with a per-cell noise factor drawn at birth. At division the first daughter inherits a
 629 volume fraction $f \sim \mathcal{N}(0.5, \sigma^2)$ (clipped to $[0.05, 0.95]$) and each molecule is assigned to it independently with
 630 probability f (binomial partitioning); a beta-binomial option models over-dispersed partitioning. Promoter groups
 631 are inherited by both daughters when unreplicated and split one copy per daughter, in random assignment, when
 632 replicated. Gene replication, if enabled, doubles the promoter counts (and the propensity of reactions flagged as
 633 copy-number dependent) once the cell has completed a set fraction of its cycle, and is reset at division. Populations
 634 can be started from per-founder states and volumes (for instance a snapshot of a burnt-in constant-size population).
 635 Population control is either constant size (every second daughter replaces a uniformly chosen cell, the behaviour
 636 of version 1), free growth (all daughters kept, with uniform subsampling above a cap and book-keeping of the true
 637 size) or logistic (free growth with an additional hazard $\lambda N/K$).

638 Drug and perturbation layer

639 A perturbation combines a dose schedule $d(t)$ (constant, pulsed, piecewise, one-compartment bolus pharmacoki-
 640 netics or an arbitrary function) with effects evaluated for every cell at every step: a death hazard

$$h = h_{\max} \frac{d^m}{\text{EC}_{50}^m + d^m} \cdot \frac{K^q}{K^q + c_P^q},$$

641 where c_P is the concentration of a protective species (survival probability over a step is $e^{-h \Delta t}$); growth inhibition
 642 by $1/(1 + (d/\text{IC}_{50})^m)$; a dose-independent fitness cost of a cell state, $1 - c_{\max} c_P^q/(K^q + c_P^q)$, on the growth rate;
 643 multiplicative modulation of any parameter by a function of the dose; and genetic perturbations (knockdown,
 644 knockout, over-expression) of any parameter in a random fraction of cells, the flag being inherited by daughters.

645 Two further effects couple the drug to the cell cycle and to the protein that resists it. A cell-cycle-dependent
 646 hazard multiplies h by

$$\beta + (1 - \beta) \exp \left[-\frac{1}{2} \left(\frac{\varphi - \varphi_0}{w} \right)^2 \right],$$

647 where φ is the cell's progress through its division cycle (0 at birth, 1 at division), so that killing peaks at φ_0 and
 648 falls to a fraction β of the peak elsewhere; $\beta = 1$ is the cycle-independent hazard and $\beta = 0$ confines killing to a
 649 window of width w around φ_0 . This represents an agent whose lesions are converted into death during replication.
 650 Stoichiometric consumption of a protective protein removes molecules of it at rate

$$k dV \frac{c_P}{c_P + K_m},$$

651 by a Poisson draw over the step. Because lesions form in proportion to the dose, the rate is set by lesion
 652 formation alone while the protein is abundant, so the pool follows the cumulative exposure rather than the peak
 653 concentration, and falls with what remains once the pool runs low; resynthesis is whatever the reaction model
 654 provides. This is the behaviour of a suicide enzyme such as MGMT, which is inactivated by the very lesion it

655 repairs.

656 Lineage statistics

657 The lineage table records, for every cell, its parent, founder (clone), generation, birth and end times, fate (divided,
658 died, removed, alive) and its volume and state at birth and at the end of its life. From it we compute Pearson
659 correlations of expression at division between mothers and daughters, sisters and cousins; the lineage autocorrelation
660 r_g between a cell and its ancestor g generations earlier, and the memory timescale τ from $r_g = e^{-g/\tau}$; single-lineage
661 traces by following the ancestral line of a surviving cell; the coefficient of variation squared of concentration across
662 the population versus along single lineages; Luria-Delbrück fluctuation tests (variance of a phenotype across clones
663 grown from single cells compared with the binomial expectation); and MemorySeq-style clonal scores [61], the
664 variance of clone-mean expression divided by its permutation-null expectation.

665 Observation models

666 True counts are converted to observations by a scRNA-seq model in which each cell captures every molecule
667 independently with a Beta-distributed efficiency (mean and coefficient of variation set by the protocol), optional
668 multinomial down-sampling to a Poisson-distributed depth, optional additional dropout and batch scaling; an
669 smFISH model with a fixed detection efficiency; and a time-lapse model that samples a species along lineages at a
670 fixed interval with additive or multiplicative measurement noise. Because counts scale with volume, library size
671 correlates with cell size, reproducing the transcriptome-size confound of real data. Outputs are written as CSV,
672 as AnnData-compatible HDF5 files, or as Newick lineage trees.

673 Inference

674 Snapshot counts of a gene in non-dividing cells follow the Beta-Poisson law of the telegraph model [34], which we
675 evaluate through Kummer’s confluent hypergeometric function with rescaled series summation, and maximise by
676 a grid-initialised Nelder-Mead search in log-parameter space. For any model, including populations, we implement
677 sequential Monte Carlo approximate Bayesian computation [62] with an adaptive tolerance schedule (the α
678 quantile of the previous generation’s distances), Gaussian perturbation kernels of variance twice the weighted
679 population variance, importance weights against the prior, and parallel evaluation of the user-supplied distance
680 function; for count distributions we use the Wasserstein-1 distance between simulated and observed samples [43].

681 Calibration to laboratory data and validation strategy

682 Reference values were transcribed from the text and tables of the cited articles into `paper/data/` with their
683 provenance (Supplementary Table 4). For the cisplatin case study the model is a resistance gene with slow promoter
684 switching whose protein product protects against a dose-dependent death hazard and against a dose-dependent
685 growth arrest (Methods above), with a 24-h cell cycle (at the fast end of reported U2OS doubling times and
686 consistent with the intermitotic times tracked by [6]). Its six free parameters (activation and inactivation rates,
687 maximal death rate, EC_{50} and Hill coefficient of death, IC_{50} of growth arrest) were fitted by the method of
688 simulated moments: the root-mean-square difference between simulated and observed fractions of cells that died,
689 divided or survived without dividing during 72 h of drug, for the training concentrations (7 and 13 μM), plus a
690 weak penalty keeping the memory of the resistant state near the two to three generations inferred from lineage
691 correlations, was minimised with the bounded optimiser described below using common random numbers (the
692 same seeds for every candidate). Validation used only quantities not fitted: the fate fractions at the held-out
693 concentration (10 μM), the biphasic kill curve and its plateau, the extent of lineage correlations in fate (sisters,

first, second and third cousins) and the invariance of single-cell division and death times across concentrations. The protocol is deliberately generic: any observable computable from a `PopulationResult` (kill curves, fate fractions, lineage correlations, snapshot distributions) can enter the distance, and the split between training and validation conditions is declared before fitting.

Clinical regimens, treatment outcomes and schedule optimisation

Regimens are built from the schedule primitives of the perturbation layer: `PulsedDose` with a finite number of cycles, `PiecewiseDose` for arbitrary calendars (the S1320 schedule), `daily_boluses` and `cycle_days` for daily oral doses with one-compartment pharmacokinetics, and `AdaptiveDose`, a feedback schedule that withdraws the drug when the population falls below a fraction of its initial size and resumes it when the population regrows above another fraction, evaluated with the current population size at every step [48]. Outcomes computed from a simulation are the net growth rate, the depth of response ($\log_{10}$ of the ratio of the initial to the smallest population), the time to progression (first time the population exceeds a threshold, 1.2, times its running minimum, the analogue of RECIST progression from the nadir; censored at the end of the simulation), the extinction probability over replicates and the cumulative dose actually applied. `optimize_schedule` minimises any such outcome over a box of schedule parameters: a log-spaced grid (or a Latin hypercube when the grid is large) is evaluated with the same random seeds for every candidate, and the best point is refined by Nelder-Mead in a bounded transform of the box; constraints such as a fixed cumulative dose enter as penalties or as the parameterisation (for instance a per-day dose equal to the budget divided by the number of dosing days). The case studies are run at the scale of a tumour rather than a culture: the cell cycle of the model is the net doubling time of the tumour (4 weeks for melanoma, 40 days for glioblastoma). The glioblastoma value lies within the volume doubling times measured for untreated tumours, a median of 21 days in one magnetic-resonance series and 50 days in another [63,64]; the melanoma value is an assumption, since net doubling times of treated cutaneous melanoma are not established, and it sets the time axis of the melanoma case study rather than the ranking of schedules, which follows from the resistance mechanism. Growth is exponential throughout, which is a simplification: the same measurements favour Gompertzian slowing in large glioblastomas [64], so absolute regrowth times in Fig. 8 are shorter than a saturating model would give. Gene-expression rates are scaled to the hour (mRNA half-life 7 h and protein half-life 35 h, generic mammalian values; at the tumour time scale both species turn over many times per generation, so the results are insensitive to them and all heritable memory resides in the promoter), and founders are drawn from the stationary distribution of promoter states. The melanoma case uses a pre-resistant fraction of 1:200 (a per-marker frequency [2] used as a proxy for the functional pre-resistant fraction, which is not the same quantity, and 5- to 50-fold higher than the 1 in 1,000 to 1 in 10,000 at which drug-naïve precursors of resistant fates have been traced [38]; the arms of Fig. 8 that lose control therefore do so earlier than the traced frequency would give, while those censored at 52 weeks are unaffected), with a memory of five net doublings (the model's generation is the net doubling of the tumour, a modelling choice rather than the cell-division-counted memory of memory sequencing [61]) and drug effects at the fractional-killing dose: sensitive cells arrest and die with a half-life of two weeks, cells in the resistant state are protected, and the drug stabilises the resistant state ten-fold (the drug-induced reprogramming of [2]). Three variants of the resistant state are compared: full protection without cost; a 50% fitness cost in the absence of drug [49]; and partial protection, in which resistant cells grow at half speed under drug. Progression is scored in two ways: from the nadir, as in RECIST (a 20% increase in diameter, taken as a 1.73-fold increase in cell number, after the 8-week lead-in; a well-mixed cell count cannot represent RECIST's absolute 5 mm threshold, its new-lesion and non-target criteria or its sum-of-diameters basis, and it assumes cell number proportional to volume, so this is an analogue rather than the clinical end point), and as loss of control, the time after randomisation at which the population exceeds 120% of its pre-treatment

size, the end point of adaptive-therapy studies, which is the one used to compare schedules that deliberately let the tumour regrow. The temozolomide case uses a one-compartment elimination half-life of 2.1 h [52] (a plasma value from a three-compartment population model of plasma and cerebrospinal fluid; that study measured cerebrospinal fluid rather than tumour tissue and found its exposure to be 20% of plasma exposure, so a plasma half-life stands here as a surrogate for what the tumour sees), the regimens of RTOG 0525 [51] (200 mg/m² on days 1-5 or 100 mg/m² on days 1-21 of 28-day cycles, six cycles), MGMT-expressing fractions of 1% and 30% for methylated and unmethylated promoters (illustrative values bracketing low and high expression rather than measured fractions; methylation shifts graded expression rather than a bimodal positive fraction), and MGMT either stable or consumed stoichiometrically by the drug, one molecule per lesion repaired ($k = 300$, $K_m = 150$ in the expression above). The rate constant was chosen so that the consumption rate at the standard bolus peak equals that of a dose-scaled first-order degradation, which isolates the consequence of letting depletion follow the cumulative exposure instead of the peak concentration; this is the suicide-enzyme behaviour that motivated dose-dense schedules.

Random regulatory networks

Random networks are built with a chosen number of activating and inhibiting edges under Erdős-Rényi or preferential-attachment (scale-free) sampling; regulation acts on the transcription rate of constitutive genes or on the promoter activation rate of telegraph genes through AND-combined Hill functions; per-gene rate constants can be drawn log-uniformly from ranges.

Validation protocols

Exactness was tested by Kolmogorov-Smirnov distances between simulated stationary distributions and the analytic laws (Poisson for birth-death, Beta-Poisson for the telegraph model, negative binomial for bursty protein production [65,66]), by agreement between kernels, and by a two-sample comparison against JumpProcesses.jl [67] on the same model.

The population and lineage layers were then validated against exact solutions for growing, dividing populations, which exist for this class of model and take different forms for a single lineage and for a population snapshot [30,35–37]. Four analytically solvable members of that class were simulated in both modes: constitutive production, production in bursts of fixed size, and a two-state promoter, each with a memoryless interdivision time; and volume-scaled synthesis from a gene that replicates at mid-cycle, with a deterministic cycle and exponential volume growth. Single lineages were generated with a mother-machine population control that keeps one daughter at random at every division, so that the recorded cells are independent; population snapshots were generated by free branching growth with uniform subsampling, which preserves the snapshot distribution. For the memoryless models the exact stationary law of the scheme the simulator runs is the leading eigenvector of one update step, exact kinetics over dt followed by division with probability $1 - e^{-dt/T}$, with the division operator carrying a factor of two in the population mode; it was computed by matrix exponentiation and checked against the closed-form factorial moments, and the step size was scanned to show the discretisation of the interdivision time vanishing. For the replication model the count distribution at every cycle phase is exactly Poisson, and the cycle-phase weights of the population were obtained by iterating the age map of the scheme from the founder ages used, so that no assumption about convergence to a stable age distribution enters. Each simulated distribution was compared with the exact law of its own mode and with that of the other mode, the second as a check that the comparison can tell them apart (Supplementary Note 14). Population-level properties were compared with the analytic expectations for cell-size scaling and concentration homeostasis [27], partitioning noise [68], lineage versus

population statistics [11] and cell-cycle-dependent burst parameters [10]. Persister simulations were compared with the quantitative signatures reported by [6], [42] and [9], and with the phenotypic-selection model of [41]. All scripts, seeds and outputs are in the `paper/` directory of the repository.

Identifiability and numerical robustness of the calibration

Because six parameters were fitted to six measured fractions, the calibration was examined in three ways. A conditional profile moves one parameter at a time over its full range on a log grid of eleven points, holding the others at their fitted values, and records the root-mean-square error of the fate fractions without the memory prior; parameters whose error stays within the binomial standard error of the measurement over a wide range are not determined by these data. A two-dimensional slice does the same over the memory of the resistant state and the fraction of cells in it, the two quantities that trade off most directly, with the death parameters held at their fitted values; a slice was used in preference to a profile with re-optimisation because, at this cost per evaluation, the inner search did not reliably recover the fitted solution and reported spuriously poor fits at short memory. Numerical robustness was checked by halving the integration step and comparing the simulated fractions with the seed-to-seed spread at the fitted parameters.

Because the lethality of cisplatin is enhanced during replication, the position of a cell in its division cycle when the drug arrives is a competing explanation for the dose-invariance of death timing and for the fraction of cells that survive without dividing. To test it, the same six parameters were refitted together with a seventh, the cycle-independent fraction β of the death hazard, with the peak of killing at the middle of the cycle ($\varphi_0 = 0.5$, $w = 0.15$), under the same objective, bounds, optimiser budget and seeds as the main fit. The two fits are compared by the root-mean-square error of the fate fractions without the memory prior, by their predictions at the held-out concentration, and by a conditional profile over β with the remaining parameters held at the cycle-dependent fit. The same gate ($\beta = 0.25$, $\varphi_0 = 0.5$, $w = 0.15$) was then applied to the persister simulations of Fig. 4 and the melanoma simulations of Fig. 8a-e, everything else unchanged, to test whether their conclusions depend on the hazard being cycle-blind (Supplementary Note 13). The window position is a modelling choice rather than a claim about either agent: what is tested is robustness to a strong cell-cycle gate of any kind.

Statistics and reproducibility

No statistical method was used to predetermine sample size, and no experimental data were generated in this study; sample sizes are the numbers of simulated cells, founders or seeds stated in each figure legend and in Supplementary Table 2. No data were excluded. Randomisation and blinding do not apply: every result is a deterministic function of an explicit random seed, and all seeds are fixed in the scripts that generate the figures. Replicates are independent simulations differing only in seed; where a mean and a standard deviation are reported, the number of seeds is given in the legend. Kolmogorov-Smirnov distances are compared with the 99% critical value for the stated sample size. Distributional agreement between kernels and with an independent implementation was assessed by two-sample Kolmogorov-Smirnov statistics. Correlations are Pearson correlations unless stated otherwise; correlations of binary fates between related cells are phi coefficients. Posterior summaries are medians, because the posteriors are asymmetric on the log scale. Schedules within a scan are evaluated on common random numbers, the same seeds for every point of the grid, so they are compared by the paired differences over those seeds rather than by an unpaired standard error, which the shared founder population would make invalid; a margin is called resolved when a two-sided paired Student's t-test rejects equality at the 5% level, which on four degrees of freedom requires the margin to exceed 2.776 of its paired standard errors rather than two. Multithreaded runs give results identical to single-threaded runs, because each cell carries its own

819 random number generator seeded from its parent.

820 Use of artificial intelligence

821 The software and the first draft of this Article were produced with Claude (Anthropic), a large language model,
822 under the direction and supervision of the author. Specifically, Claude wrote the Julia source of the package, the
823 simulation and plotting scripts, the test suite and the first draft of the text, working from the author's specification
824 of the scope, the models, the validation strategy and the case studies. The author defined every scientific question
825 addressed here, chose the models, parameters and reference data, independently verified all numerical results
826 against the generated output tables, checked every claim against its cited source, and wrote the final version of
827 the manuscript; the author takes full responsibility for the content. The tools used were Claude Opus 5, Claude
828 Opus 4.8 and Claude Fable 5.1 (Anthropic), accessed through the Claude Code command-line interface between
829 19 and 22 September 2026; the model used for each commit is recorded in that commit's trailer in the repository,
830 so the record is auditable rather than summary. The model was not used to generate citations: every reference
831 was located by the author and its bibliographic details and quantitative claims were checked against the primary
832 record. The model is not an author, in keeping with journal policy on generative artificial intelligence, and the
833 author is accountable for the accuracy and integrity of everything reported here.

834 No display item in this Article is generated by a generative model. Every figure and table is produced deter-
835 ministically from simulation output by the scripts in the `paper/` directory: the panels by `matplotlib` from the
836 comma-separated tables in `paper/output/`, and every number quoted in the Results and the Supplementary
837 Information by `paper/manuscript/fill_numbers.py`, which substitutes values read from those tables into the
838 manuscript source. Figure 1 is a schematic drawn by `paper/fig1_schematic.py`. Re-running the scripts from
839 the fixed seeds reproduces all of them.

840 Data availability

841 The laboratory and clinical reference values used for calibration and validation were transcribed from the
842 published articles cited in Supplementary Table 4 and are provided, with their provenance, in the `paper/data/`
843 directory of the software repository (<https://github.com/ayoublasri/Biomodelling.jl>). All other data in this Article
844 are simulated: they are regenerated exactly, from fixed seeds, by the scripts in the `paper/` directory of that
845 repository, so those scripts rather than a deposited table are the primary record. The archived release carries
846 them (10.5281/zenodo.22894286).

847 Code availability


848 Biomodelling.jl 2.0 is open source under the MIT licence at <https://github.com/ayoublasri/Biomodelling.jl>, with
849 documentation, tutorials and a test suite that runs under continuous integration. Every figure is regenerated
850 from fixed seeds by the scripts in `paper/`. After instantiating the environment (`julia --project=paper -e`
851 `'using Pkg; Pkg.develop(path="."); Pkg.instantiate()'`), the simulation data are produced by `julia`
852 `--project=paper -t 4 paper/run_all.jl`, the Python network-inference and imputation benchmarks of Fig.
853 5b–d by `python paper/scripts/benchmarks_python.py`, and the identifiability analysis of Supplementary
854 Fig. 3 by `julia --project=paper paper/scripts/fig7_calibration.jl profile`. The remaining opt-in
855 analyses, the cell-cycle refits and scans (`fig7_calibration.jl cycle`, `fig4_persisters.jl cycle seeds`,
856 `fig8_schedules.jl melanoma_cycle melanoma_memory`) and the exact-solution validation of Supplementary
857 Fig. 6 (`fig9_validation.jl`, followed by `python paper/scripts/exact_solutions.py`), are listed in
858 `paper/README.md` with their invocations. The panels are then assembled by `python paper/fig1_schematic.py`,

python paper/plot_all.py and python paper/plot_supplement.py. Full instructions, including the four-thread invocation that produces the timings of Supplementary Table 3, are in paper/README.md. Version 2.0.0, the version used for every result in this Article, is archived at Zenodo under the concept DOI 10.5281/zenodo.22894286, which resolves to the current release.

Software

Biomodelling.jl 2.0 is written in Julia (≥ 1.10) and depends on Distributions.jl, PoissonRandom.jl, QuadGK.jl, Optim.jl, SpecialFunctions.jl and StatsBase.jl; HDF5.jl is an optional extension. Benchmarks used absolute Pearson and Spearman correlations and scikit-learn random forests re-implementing GENIE3 [69], MAGIC [70] and kNN-smoothing [71]. The test suite (Aqua quality checks and statistical tests) runs under GitHub Actions on Julia 1.10 and the latest release.


Figure legends



870

Fig. 1 | The Biomodelling.jl 2.0 framework. **a**, One cell. Reactions run at the cell's current volume, with propensities scaled so that concentrations are invariant under growth; genes replicate at a set point of the cycle; division splits every molecule binomially while both daughters inherit the promoter state. Inset: exponential volume growth, halved at each division. **b**, Populations under free growth, constant-size or logistic control, with dose schedules and a state-dependent death hazard evaluated for every cell at every step, and the lineage record from which heritability, autocorrelation, memory timescales, clone diversity and fluctuation tests are computed; crosses mark cells killed by the drug. **c**, Observation models turn true counts into scRNA-seq counts with


capture, depth, dropout and batch effects, into smFISH spots, or into time-lapse reporter traces; outputs are CSV, AnnData and Newick. **d**, Uses of the forward model: detection of memory genes, benchmarking of network and perturbation inference against known ground truth, likelihood-free parameter inference, and schedule optimisation over the period and the fraction of time on drug (star, optimum). **e**, A complete experiment in twelve lines of code.



883

Fig. 2 | The engine matches exact laws and an independent implementation. **a**, Stationary distribution of a birth-death process (bars, 50,000 simulated cells; line, Poisson law). **b**, Telegraph model in an intermediate and a bursty regime against the Beta-Poisson law. **c**, Bursty protein production against the negative binomial. **d**, Kolmogorov-Smirnov distance to the exact law for every kernel (n = 20,000; dotted line, 99% critical value) with wall-clock time per 10^4 cells. **e**, Telegraph model simulated with Biomodelling.jl and with JumpProcesses.jl on the same model; the two-sample Kolmogorov-Smirnov statistic lies below its 99% critical value. **f**, Runtime of population simulations against the number of cells (10 genes) and genes (1,000 cells) for the exact and hybrid kernels.

889



892

893 **Fig. 3 | Growth, division and partitioning generate memory, size scaling and cell-cycle effects.** **a**, A
 894 single lineage: volume, mRNA and protein of a slowly switching gene, with the promoter-on state shaded. **b**,
 895 mRNA count against cell volume for a fast-switching gene without replication (points; line, mean concentration
 896 $\times V$) and with gene replication at mid-cycle (binned means). **c**, Correlation at division between mothers
 897 and daughters, sisters and cousins as a function of the promoter switching rate (protein concentration, solid;
 898 mother-daughter mRNA, dashed); dotted line, inverse cell-cycle time. **d**, Memory timescale against the measured
 899 cell-cycle time for two switching rates; dotted lines, $1/(k_{\text{on}} + k_{\text{off}})$. **e**, Squared coefficient of variation of a stable
 900 protein's concentration across the population and along single lineages, for increasing partitioning noise and
 901 cell-cycle-time noise. **f**, Telegraph parameters fitted to cells before replication (one gene copy), after replication
 902 (two copies) and pooled. The transcription rate is plotted per unit cell volume and divided by twenty, so that all
 903 three series share an axis and the dotted lines mark the true per-allele values ($k_{\text{on}} = 0.5$, $k_{\text{off}} = 1.5$, and 2 for the
 904 scaled transcription rate); the raw fitted transcription rates are in `paper/output/fig3f_copynumber.csv`.





Fig. 4 | A single resistance gene with slow promoter switching reproduces persister biology. **a**, Population size under continuous, pulsed and reduced-dose schedules (shading, treatment window). **b**, Kill curves at increasing dose. **c**, Net population decay rate over the first 30 time units of treatment against dose, with the mean division time of cells born under drug (inset), which is unchanged. **d**, Probability that a cell's lineage survives the treatment (at least one descendant alive at the end of the run) given that its sister's or cousin's lineage survived, against the marginal survival probability, for cells born in the 25 time units before the drug; memory gene and fast-switching control, with the fold enrichment above each bar. **e**, Effective number of clones (barcodes) before and after drug for pre-existing and drug-induced tolerance (mean $\pm$ s.d. of three seeds). **f**, Surviving clones after drug without memory disruption, with promoter switching accelerated 20-fold from two cell-cycle times before the drug until the drug arrives, and with the acceleration maintained during exposure (mean $\pm$ s.d. of five seeds). **g**, Fraction of high-expressing cells before, during and after a drug pulse for slow and fast switching. **h**, Net population growth rate over a 200-time-unit course as a function of release period (time units without drug between 20-time-unit exposures) and dose, for pre-existing tolerance without and with a fitness cost of the resistant state and for drug-induced tolerance; each cell of the map is the mean of five independent

seeds and the box marks the schedule with the lowest mean growth rate; every schedule is run on the same five seeds and compared by paired differences, and the margin over the best alternative with a non-zero release period is resolved in only one of the three mechanisms (Supplementary Note 4).



923

Fig. 5 | Ground-truth benchmarks. **a**, Clonal variance score against the known memory time of 40 independent genes, from true concentrations and from sequenced counts; dotted line, cell-cycle time. **b**, Area under the precision-recall curve of correlation (Pearson, Spearman) and GENIE3 network inference on the same 30-gene network observed at fixed volume, in a dividing population (raw counts, concentrations, and concentrations regressed on the cell-cycle covariates copy number and age) and after sequencing (raw and library-size normalised counts); dotted line, random ranking. **c**, Effect of kNN-smoothing and MAGIC imputation before inference. **d**, Coefficient of determination of predicted $\log_2$ fold changes after knocking down the five most connected regulators, for a no-change baseline and for propagation of observational correlations.

931





Fig. 6 | Inference with the forward model. **a**, Posterior distributions of telegraph parameters from ABC-SMC on a snapshot of non-dividing cells, with the truth and the exact maximum-likelihood estimate. **b**, The same for a snapshot of a growing and dividing population: the naive telegraph fit is biased while division-aware ABC recovers the truth. **c**, Posterior particles for the maximal death rate and the protection threshold inferred from a kill curve and sister fate concordance: the two parameters trade off along a ridge that contains the truth. **d**, Posterior predictive kill curves simulated at the posterior median against the observed curve.



939

Fig. 7 | Calibration to laboratory time-lapse data and validation on held-out conditions. **a**, Fates over 72 h of cisplatin of U2OS cells present at drug addition, observed by time-lapse microscopy (grey; ref. [6]) and simulated by the calibrated model (colours; mean $\pm$ s.d. of four seeds); the model was calibrated on 7 and 13 μM and the 10 μM condition is predicted. **b**, Predicted population size at the three concentrations and at 11.5 μM , from two days before the drug to seven days of exposure. **c**, Correlation of death fates between sisters, first, second and third cousins present at drug addition, and between unrelated cells (mean $\pm$ s.d. of three seeds), against the observation that correlations extend to second but not third cousins. That observation is reported graphically in ref. [6] with no numeric coefficients, so the comparison is of pattern rather than of value. **d**, Mean time to death of cells present at drug addition against concentration (mean $\pm$ s.d. of four seeds); division times are not shown because too few cells born under drug divided again to estimate them (at most four per condition). **e**, Calibration landscape: every parameter set evaluated by the optimiser, coloured by the objective it minimised, which is the single-seed training error plus the memory prior and is not the seed-averaged error quoted in the text. **f**, Calibrated parameters.

days) temozolomide regimens of the RTOG 0525 trial and a 21-day regimen at the standard cumulative dose, with one-compartment pharmacokinetics (elimination half-life 2.1 h, a plasma value; an independent study gives 1.8 h) and stoichiometric consumption of MGMT by the drug, one molecule per lesion repaired, for populations in which 1% (methylated promoter, solid) or 30% (unmethylated, dashed) of cells express MGMT. **g**, Depth of response over six cycles for the three regimens, the two populations and MGMT either stable or consumed by the drug (mean $\pm$ s.d. of two seeds); a bar at zero marks a regimen that achieves no net kill (the 21-day schedule at the standard cumulative dose). **h**, Final population size against the number of dosing days per 28-day cycle at the standard cumulative dose.

Acknowledgements

Version 1 of the framework was developed with Marc Sturrock and Vahid Shahrezaei, whose contributions to that work are described in ref. [33]. This Article and version 2.0 of the software were authored with Claude (Anthropic), a large language model, under the direction and supervision of the author; what the model did and what the author verified independently are set out in the Methods, under Use of artificial intelligence.

Author contributions

A.L. conceived the study, defined its scope, the validation strategy and the case studies, supervised the work, verified all numerical results and biological interpretations against their sources, and wrote the final version of the manuscript. The implementation of the software, the simulations, the analyses, the figures and the first draft of the manuscript were produced with Claude (Anthropic) under the author's direction, as described in the Methods. Claude is not an author and no display item was produced by a generative model.

Competing interests

The author declares no competing interests.

Additional information

Supplementary information accompanies this Article. Correspondence and requests for materials should be addressed to A.L.

References

- Sharma SV, Lee DY, Li B, Quinlan MP, Takahashi F, Maheswaran S, et al. A chromatin-mediated reversible drug-tolerant state in cancer cell subpopulations. *Cell*. 2010;141: 69–80. doi:10.1016/j.cell.2010.02.027
- Shaffer SM, Dunagin MC, Torborg SR, Torre EA, Emert B, Krepler C, et al. Rare cell variability and drug-induced reprogramming as a mode of cancer drug resistance. *Nature*. 2017;546: 431–435. doi:10.1038/nature22794
- Hata AN, Niederst MJ, Archibald HL, Gomez-Caraballo M, Siddiqui FM, Mulvey HE, et al. Tumor cells can follow distinct evolutionary paths to become resistant to epidermal growth factor receptor inhibition. *Nature Medicine*. 2016;22: 262–269. doi:10.1038/nm.4040

- 995 4. Oren Y, Tsabar M, Cuoco MS, Amir-Zilberstein L, Cabanos HF, Hütter J-C, et al. Cycling cancer persister
996 cells arise from lineages with distinct programs. *Nature*. 2021;596: 576–582. doi:10.1038/s41586-021-03796-6
- 997 5. França GS, Baron M, King BR, Bossowski JP, Bjornberg A, Pour M, et al. Cellular adaptation to cancer
998 therapy along a resistance continuum. *Nature*. 2024;631: 876–883. doi:10.1038/s41586-024-07690-9
- 999 6. Iyer A, Alva A, Granada AE, Chakrabarti S. Inheritable cell-states shape drug-persister correla-
1000 tions and population dynamics in cancer cells. *PLoS Computational Biology*. 2025;21: e1013446.
doi:10.1371/journal.pcbi.1013446
- 1001 7. Lin Y, Chen X, Wu L, Zhou Y, Lin Y. Widespread transcriptional memory shapes heritable states and func-
1002 tional heterogeneity in cancer and stem cells. *Cell Systems*. 2026;17: 101621. doi:10.1016/j.cels.2026.101621
- 1003 8. Ghosh S, Chakrabarti S, Raju A. Identifying memory gene expression from single sample scRNA-seq data
1004 using power law signatures. *bioRxiv*. 2025. doi:10.1101/2025.01.15.633183
- 1005 9. Harmange G, Hueros RAR, Schaff DL, Emert B, Saint-Antoine M, Kim LC, et al. Disrupting cellular
1006 memory to overcome drug resistance. *Nature Communications*. 2023;14: 7130. doi:10.1038/s41467-023-
41811-8
- 1007 10. Sukys A, Grima R. Cell-cycle dependence of bursty gene expression: Insights from fitting mechanistic
1008 models to single-cell RNA-seq data. *Nucleic Acids Research*. 2025;53: gkaf295. doi:10.1093/nar/gkaf295
- 1009 11. Zhang Z, Zabaikina I, Nieto C, Vahdat Z, Bokes P, Singh A. Stochastic gene expression in proliferating
1010 cells: Differing noise intensity in single-cell and population perspectives. *PLoS Computational Biology*.
2025;21: e1013014. doi:10.1371/journal.pcbi.1013014
- 1011 12. Okochi Y, Sawazaki Y, Kondo Y, Naoki H. Inferring division-associated stochasticity from time-series
1012 single-cell transcriptomes. *bioRxiv*. 2026. doi:10.64898/2026.04.14.718485
- 1013 13. Ahlmann-Eltze C, Huber W, Anders S. Deep-learning-based gene perturbation effect prediction does not yet
1014 outperform simple linear baselines. *Nature Methods*. 2025;22: 1657–1661. doi:10.1038/s41592-025-02772-6
- 1015 14. Vollenweider MS, Bühlmann P. Signal, bounds, and baselines: Principles for rigorous evaluation of
1016 high-dimensional biological perturbation prediction. *bioRxiv*. 2026. doi:10.64898/2026.04.20.719650
- 1017 15. Dibaeinia P, Babu S, Knudson M, ElSheikh A, Wen Y, Liu H, et al. Virtual cells need context, not just
1018 scale. *bioRxiv*. 2026. doi:10.64898/2026.02.04.703804
- 1019 16. Bunne C, Roohani Y, Rosen Y, Gupta A, Zhang X, Roed M, et al. How to build the virtual cell with artificial
1020 intelligence: Priorities and opportunities. *Cell*. 2024;187: 7045–7063. doi:10.1016/j.cell.2024.11.015
- 1021 17. Zappia L, Phipson B, Oshlack A. Splatter: Simulation of single-cell RNA sequencing data. *Genome Biology*.
1022 2017;18: 174. doi:10.1186/s13059-017-1305-0
- 1023 18. Song D, Wang Q, Yan G, Liu T, Sun T, Li JJ. scDesign3 generates realistic in silico data for multimodal
1024 single-cell and spatial omics. *Nature Biotechnology*. 2024;42: 247–252. doi:10.1038/s41587-023-01772-1
- 1025 19. Zhang X, Xu C, Yosef N. Simulating multiple faceted variability in single cell RNA sequencing. *Nature*
1026 *Communications*. 2019;10: 2611. doi:10.1038/s41467-019-10500-w
- 1027 20. Dibaeinia P, Sinha S. SERGIO: A single-cell expression simulator guided by gene regulatory networks. *Cell*
1028 *Systems*. 2020;11: 252–271. doi:10.1016/j.cels.2020.08.003
- 1029 21. Cannoodt R, Saelens W, Deconinck L, Saeys Y. Spearheading future omics analyses using dyngen, a multi-
1030 modal simulator of single cells. *Nature Communications*. 2021;12: 3942. doi:10.1038/s41467-021-24152-2

- 1031 22. Li H, Zhang Z, Squires M, Chen X, Zhang X. scMultiSim: Simulation of single-cell multi-omics and spatial
data guided by gene regulatory networks and cell–cell interactions. *Nature Methods*. 2025;22: 982–993.
1032 doi:10.1038/s41592-025-02651-0
- 1033 23. Zinati Y, Takiddeen A, Emad A. GRouNdGAN: GRN-guided simulation of single-cell RNA-seq data using
causal generative adversarial networks. *Nature Communications*. 2024;15: 4055. doi:10.1038/s41467-024-
1034 48516-6
- 1035 24. Pan X, Li H, Zhang X. TedSim: Temporal dynamics simulation of single-cell RNA sequencing data and
cell division history. *Nucleic Acids Research*. 2022;50: 4272–4288. doi:10.1093/nar/gkac235
1036
- 1037 25. Jones MG, Khodaverdian A, Quinn JJ, Chan MM, Hussmann JA, Wang R, et al. Inference of single-cell
phylogenies from lineage tracing data using Cassiopeia. *Genome Biology*. 2020;21: 92. doi:10.1186/s13059-
1038 020-02000-8
- 1039 26. Pratapa A, Jalihal AP, Law JN, Bharadwaj A, Murali TM. Benchmarking algorithms for gene reg-
ulatory network inference from single-cell transcriptomic data. *Nature Methods*. 2020;17: 147–154.
1040 doi:10.1038/s41592-019-0690-6
- 1041 27. Bertaux F, Marguerat S, Shahrezaei V. Division rate, cell size and proteome allocation: Impact on gene
expression noise and implications for the dynamics of genetic circuits. *Royal Society Open Science*. 2018;5:
1042 172234. doi:10.1098/rsos.172234
- 1043 28. Piho P, Thomas P. AgentBasedModeling.jl: A tool for stochastic simulation of structured population
dynamics. *PLoS Computational Biology*. 2026;22: e1014366. doi:10.1371/journal.pcbi.1014366
1044
- 1045 29. Piho P, Thomas P. Feedback between stochastic gene networks and population dynamics enables cellular
decision-making. *Science Advances*. 2024;10: eadl4895. doi:10.1126/sciadv.adl4895
1046
- 1047 30. Thomas P, Shahrezaei V. Coordination of gene expression noise with cell size: Analytical results for
agent-based models of growing cell populations. *Journal of the Royal Society Interface*. 2021;18: 20210274.
1048 doi:10.1098/rsif.2021.0274
- 1049 31. Ghaffarizadeh A, Heiland R, Friedman SH, Mumenthaler SM, Macklin P. PhysiCell: An open source
physics-based cell simulator for 3-D multicellular systems. *PLoS Computational Biology*. 2018;14: e1005991.
1050 doi:10.1371/journal.pcbi.1005991
- 1051 32. Ponce-de-Leon M, Montagud A, Noël V, Meert A, Pradas G, Barillot E, et al. PhysiBoSS 2.0: A
sustainable integration of stochastic Boolean and agent-based modelling frameworks. *npj Systems Biology
and Applications*. 2023;9: 54. doi:10.1038/s41540-023-00314-4
1052
- 1053 33. Lasri A, Shahrezaei V, Sturrock M. Benchmarking imputation methods for network inference using a novel
method of synthetic scRNA-seq data generation. *BMC Bioinformatics*. 2022;23: 236. doi:10.1186/s12859-
1054 022-04778-9
- 1055 34. Peccoud J, Ycart B. Markovian modeling of gene-product synthesis. *Theoretical Population Biology*.
1995;48: 222–234. doi:10.1006/tpbi.1995.1027
1056
- 1057 35. Beentjes CHL, Perez-Carrasco R, Grima R. Exact solution of stochastic gene expression mod-
els with bursting, cell cycle and replication dynamics. *Physical Review E*. 2020;101: 032403.
1058 doi:10.1103/PhysRevE.101.032403
- 1059 36. Jia C, Grima R. Coupling gene expression dynamics to cell size dynamics and cell cycle events:
Exact and approximate solutions of the extended telegraph model. *iScience*. 2023;26: 105746.
1060 doi:10.1016/j.isci.2022.105746

- 1061 37. Thomas P. Making sense of snapshot data: Ergodic principle for clonal cell populations. *Journal of the*
1062 *Royal Society Interface*. 2017;14: 20170467. doi:10.1098/rsif.2017.0467
- 1063 38. Emert BL, Cote CJ, Torre EA, Dardani IP, Jiang CL, Jain N, et al. Variability within rare cell states enables
1064 multiple paths toward drug resistance. *Nature Biotechnology*. 2021;39: 865–876. doi:10.1038/s41587-021-
00837-3
- 1065 39. Torre EA, Arai E, Bayatpour S, Jiang CL, Beck LE, Emert BL, et al. Genetic screening for single-cell
1066 variability modulators driving therapy resistance. *Nature Genetics*. 2021;53: 76–85. doi:10.1038/s41588-
020-00749-z
- 1067 40. Gunnarsson EB, De S, Leder K, Foo J. Understanding the role of phenotypic switching in cancer drug
1068 resistance. *Journal of Theoretical Biology*. 2020;490: 110162. doi:10.1016/j.jtbi.2020.110162
- 1069 41. Lasri A, Juric V, Verreault M, Bielle F, Idhahbi A, Kel A, et al. Phenotypic selection through cell death:
1070 Stochastic modelling of O-6-methylguanine-DNA methyltransferase dynamics. *Royal Society Open Science*.
2020;7: 191243. doi:10.1098/rsos.191243
- 1071 42. Corigliano M, Di Bernardo A, Cosentino Lagomarsino M, Pompei S. Optimal treatment for drug-induced
1072 cancer persists involves release periods and intermediate drug doses. *PRX Life*. 2025. doi:10.1103/mh5v-
lcyw
- 1073 43. Bernton E, Jacob PE, Gerber M, Robert CP. Approximate Bayesian computation with the Wasserstein
1074 distance. *Journal of the Royal Statistical Society: Series B*. 2019;81: 235–269. doi:10.1111/rssb.12312
- 1075 44. Sandler O, Mizrahi SP, Weiss N, Agam O, Simon I, Balaban NQ. Lineage correlations of single cell division
1076 time as a probe of cell-cycle dynamics. *Nature*. 2015;519: 468–471. doi:10.1038/nature14318
- 1077 45. Kuchen EE, Becker NB, Claudino N, Höfer T. Hidden long-range memories of growth and cycle speed
1078 correlate cell cycles in lineage trees. *eLife*. 2020;9: e51002. doi:10.7554/eLife.51002
- 1079 46. Chakrabarti S, Paek AL, Reyes J, Lasick KA, Lahav G, Michor F. Hidden heterogeneity and
1080 circadian-controlled cell fate inferred from single cell lineages. *Nature Communications*. 2018;9: 5372.
doi:10.1038/s41467-018-07788-5
- 1081 47. Algazi AP, Othus M, Daud AI, Lo RS, Mehnert JM, Truong T-G, et al. Continuous versus intermittent
1082 BRAF and MEK inhibition in patients with BRAF-mutated melanoma: A randomized phase 2 trial.
Nature Medicine. 2020;26: 1564–1568. doi:10.1038/s41591-020-1060-8
- 1083 48. Zhang J, Cunningham JJ, Brown JS, Gatenby RA. Integrating evolutionary dynamics into treatment of
1084 metastatic castrate-resistant prostate cancer. *Nature Communications*. 2017;8: 1816. doi:10.1038/s41467-
017-01968-5
- 1085 49. Das Thakur M, Salangsang F, Landman AS, Sellers WR, Pryer NK, Levesque MP, et al. Modelling
1086 vemurafenib resistance in melanoma reveals a strategy to forestall drug resistance. *Nature*. 2013;494:
251–255. doi:10.1038/nature11814
- 1087 50. Strobl MAR, West J, Viosat Y, Damaghi M, Robertson-Tessi M, Brown JS, et al. Turnover modulates the
1088 need for a cost of resistance in adaptive therapy. *Cancer Research*. 2021;81: 1135–1147. doi:10.1158/0008-
5472.CAN-20-0806
- 1089 51. Gilbert MR, Wang M, Aldape KD, Stupp R, Hegi ME, Jaeckle KA, et al. Dose-dense temozolomide for
1090 newly diagnosed glioblastoma: A randomized phase III clinical trial. *Journal of Clinical Oncology*. 2013;31:
4085–4091. doi:10.1200/JCO.2013.49.6968

- 1091 52. Ostermann S, Csajka C, Buclin T, Leyvraz S, Lejeune F, Decosterd LA, et al. Plasma and cerebrospinal
fluid population pharmacokinetics of temozolomide in malignant glioma patients. *Clinical Cancer Research*.
1092 2004;10: 3728–3736. doi:10.1158/1078-0432.CCR-03-0807
- 1093 53. Rudek MA, Donehower RC, Statkevich P, Batra VK, Cutler DL, Baker SD. Temozolomide in pa-
tients with advanced cancer: Phase I and pharmacokinetic study. *Pharmacotherapy*. 2004;24: 16–25.
1094 doi:10.1592/phco.24.1.16.34800
- 1095 54. Tolcher AW, Gerson SL, Denis L, Geyer C, Hammond LA, Patnaik A, et al. Marked inactivation of
O6-alkylguanine-DNA alkyltransferase activity with protracted temozolomide schedules. *British Journal of*
1096 *Cancer*. 2003;88: 1004–1011. doi:10.1038/sj.bjc.6600827
- 1097 55. Middleton MR, Lee SM, Arance A, Wood M, Thatcher N, Margison GP. O6-methylguanine formation,
repair protein depletion and clinical outcome with a 4 hr schedule of temozolomide in the treatment of
advanced melanoma: Results of a phase II study. *International Journal of Cancer*. 2000;88: 469–473.
1098 doi:10.1002/1097-0215(20001101)88:3<469::AID-IJC21>3.0.CO;2-7
- 1099 56. Robinson CG, Palomo JM, Rahmathulla G, McGraw M, Donze J, Liu L, et al. Effect of alternative
temozolomide schedules on glioblastoma O6-methylguanine-DNA methyltransferase activity and survival.
1100 *British Journal of Cancer*. 2010;103: 498–504. doi:10.1038/sj.bjc.6605792
- 1101 57. Gillespie DT. Exact stochastic simulation of coupled chemical reactions. *The Journal of Physical Chemistry*.
1102 1977;81: 2340–2361. doi:10.1021/j100540a008
- 1103 58. Gillespie DT. Approximate accelerated stochastic simulation of chemically reacting systems. *The Journal*
1104 *of Chemical Physics*. 2001;115: 1716–1733. doi:10.1063/1.1378322
- 1105 59. Cao Y, Gillespie DT, Petzold LR. Efficient step size selection for the tau-leaping simulation method. *The*
1106 *Journal of Chemical Physics*. 2006;124: 044109. doi:10.1063/1.2159468
- 1107 60. Cao Y, Gillespie DT, Petzold LR. Avoiding negative populations in explicit Poisson tau-leaping. *The*
1108 *Journal of Chemical Physics*. 2005;123: 054104. doi:10.1063/1.1992473
- 1109 61. Shaffer SM, Emert BL, Reyes Hueros RA, Cote C, Harmange G, Schaff DL, et al. Memory sequencing
reveals heritable single-cell gene expression programs associated with distinct cellular behaviors. *Cell*.
1110 2020;182: 947–959. doi:10.1016/j.cell.2020.07.003
- 1111 62. Toni T, Welch D, Strelkowa N, Ipsen A, Stumpf MPH. Approximate Bayesian computation scheme for
parameter inference and model selection in dynamical systems. *Journal of the Royal Society Interface*.
1112 2009;6: 187–202. doi:10.1098/rsif.2008.0172
- 1113 63. Ellingson BM, Nguyen HN, Lai A, Nechifor RE, Zaw O, Pope WB, et al. Contrast-enhancing tumor
growth dynamics of preoperative, treatment-naive human glioblastoma. *Cancer*. 2016;122: 1718–1727.
1114 doi:10.1002/cncr.29957
- 1115 64. Stensj en AL, Solheim O, Kvistad KA, H aberg AK, Salvesen  , Berntsen EM. Growth dynamics of
untreated glioblastomas in vivo. *Neuro-Oncology*. 2015;17: 1402–1411. doi:10.1093/neuonc/nov029
- 1116 65. Friedman N, Cai L, Xie XS. Linking stochastic dynamics to population distribution: An analytical framework
of gene expression. *Physical Review Letters*. 2006;97: 168302. doi:10.1103/PhysRevLett.97.168302
- 1117 66. Shahrezaei V, Swain PS. Analytical distributions for stochastic gene expression. *Proceedings of the National*
1118 *Academy of Sciences*. 2008;105: 17256–17261. doi:10.1073/pnas.0803850105
- 1119 67. Rackauckas C, Nie Q. DifferentialEquations.jl – a performant and feature-rich ecosystem for solving
1120 differential equations in Julia. *Journal of Open Research Software*. 2017;5: 15. doi:10.5334/jors.151
- 1121
- 1122

- 1123 68. Huh D, Paulsson J. Non-genetic heterogeneity from stochastic partitioning at cell division. *Nature Genetics*.
1124 2011;43: 95–100. doi:10.1038/ng.729
- 1125 69. Huynh-Thu VA, Irrthum A, Wehenkel L, Geurts P. Inferring regulatory networks from expression data
1126 using tree-based methods. *PLoS ONE*. 2010;5: e12776. doi:10.1371/journal.pone.0012776
- 1127 70. Dijk D van, Sharma R, Nainys J, Yim K, Kathail P, Carr AJ, et al. Recovering gene interactions from
1128 single-cell data using data diffusion. *Cell*. 2018;174: 716–729. doi:10.1016/j.cell.2018.05.061
- 1129 71. Wagner F, Yan Y, Yanai I. K-nearest neighbor smoothing for high-throughput single-cell RNA-Seq data.
1130 *bioRxiv*. 2018. doi:10.1101/217737