

Supplementary Information

Ayoub Lasri

Mechanistic simulation of heritable expression states, cell division and drug response in single-cell populations. Supplementary Notes 1–13, Supplementary Figures 1–5 and Supplementary Tables 1–4.

Supplementary Note 1: Stationary laws used for validation

Birth-death. For $\emptyset \xrightarrow{k} X \xrightarrow{\gamma} \emptyset$ the stationary distribution is Poisson with mean k/γ .

Telegraph model. With activation k_{on} , inactivation k_{off} , transcription k_{tx} from the active state and degradation γ , and $a = k_{\text{on}}/\gamma$, $b = k_{\text{off}}/\gamma$, $\lambda = k_{\text{tx}}/\gamma$, the stationary distribution is the Beta-Poisson mixture

$$P(n) = \frac{\lambda^n \Gamma(a+n) \Gamma(a+b)}{n! \Gamma(a+b+n) \Gamma(a)} {}_1F_1(a+n; a+b+n; -\lambda),$$

whose mean is $\lambda a/(a+b)$ and whose variance is $\mu + \lambda^2 ab/[(a+b)^2(a+b+1)]$. We evaluate ${}_1F_1(\alpha; \beta; -\lambda)$ through Kummer’s transformation $e^{-\lambda} {}_1F_1(\beta - \alpha; \beta; \lambda)$, whose series has positive terms; partial sums are rescaled by 10^{-200} whenever they exceed 10^{200} so that the log of the sum is accumulated without overflow.

Bursty protein. Protein produced in geometric bursts of mean size b arriving at rate a and degraded at rate γ is negative binomial with shape a/γ and success probability $1/(1+b)$. The package emulates bursts through an mRNA of lifetime $1/(200\gamma)$ translated at rate $200\gamma b$.

Supplementary Note 2: Kernel details and correctness tests

The direct method draws the waiting time from $\text{Exp}(a_0)$ and the channel by linear search over the cumulative propensities; after each event only the propensities of reactions that depend on a changed species (the dependency graph) are recomputed, and the running total is refreshed every 512 events. Fixed-step tau-leaping draws $\text{Poisson}(a_j \tau)$ firings for every channel; the strict variant errors when a species would become negative, the hybrid variant discards the leap and simulates that interval with the direct method. Discarding is conditioned on the realised leap, so the accepted steps of the hybrid variant follow a truncated Poisson law and the kernel is approximate; Fig. 2d measures the resulting discrepancy against the exact stationary laws at two step sizes. Adaptive tau-leaping classifies as critical every channel within $n_c = 10$ firings of exhausting a reactant, selects τ from the non-critical channels through the bounds $\max(\varepsilon x_i/g_i, 1)/|\mu_i|$ and $\max(\varepsilon x_i/g_i, 1)^2/\sigma_i^2$, draws the time to the next critical event from $\text{Exp}(a_0^{\text{crit}})$, fires at most one critical channel, halves τ on rejection, and performs 100 exact events whenever $\tau < 10/a_0$.

Supplementary Table 1 lists the tests of the package test suite that validate the kernels: Kolmogorov-Smirnov distances below 0.015 (direct) or 0.02–0.03 (approximate kernels) to the Poisson, Beta-Poisson and negative-binomial laws with 8,000 to 12,000 cells; agreement of the adaptive kernel with the direct method on a stiff

30 dimerisation network within 5% in the mean; identical results with one and with several threads.

31 **Supplementary Note 3: Parameters of all simulations**


32 Supplementary Table 2. Parameters of every simulation in this Article.

Figure	Model	Parameters
2a	birth-death	$k = 12, \gamma = 1, T = 15$
2b	telegraph	$(k_{\text{on}}, k_{\text{off}}, k_{\text{tx}}) = (0.4, 0.6, 12)$ and $(0.05, 0.5, 40), \gamma = 1, T = 40$
2c	bursty protein	$a = 1.5, b = 6, \gamma = 1, T = 25$
2d-e	telegraph	$(0.4, 0.6, 12), n = 20\,000$
2f	random telegraph GRNs	$n_{\text{act}} = n_{\text{genes}}, n_{\text{inh}} = n_{\text{genes}}/2, \lambda = \ln 2/20, dt = 0.1, T = 20$
3	telegraph gene with protein	$k_{\text{tx}} = 30, k_{\text{dm}} = 1, k_{\text{tl}} = 4, k_{\text{dp}} = 0.2, k_{\text{on}} = k_{\text{off}}$ as indicated; $\lambda = \ln 2/20$, sizer $V_{\text{div}} = 2$ (cv 0.05), $\sigma_f = 0.02$, replication at 50% of the cycle where indicated; panel 3d varies the cell-cycle time over 10, 20 and 40; panel 3f uses $(k_{\text{on}}, k_{\text{off}}, k_{\text{tx}}) = (0.5, 1.5, 40)$ with 3,000 cells
3e	stable protein	$k_{\text{on}} = k_{\text{off}} = 5, k_{\text{tx}} = 20, k_{\text{dm}} = 1, k_{\text{tl}} = 2, k_{\text{dp}} = 0.05$
4	resistance gene	as Fig. 3 with $(k_{\text{on}}, k_{\text{off}}) = (0.002, 0.02)$ (memory, $p_{\text{on}} \approx 9\%$) or $(0.05, 0.5)$ (fast, same p_{on}); founders burnt in for 300 time units at constant size; death $h_{\text{max}} = 0.15, \text{EC}_{50} = 0.5, m = 2, K = 150, q = 4$; drug from $t = 40$ of the treatment stage; drug-induced variant $(0.0002, 0.02)$ with $k_{\text{on}} \rightarrow k_{\text{on}}(1 + 100d)$; fitness cost variant: growth $\times (1 - 0.5 c_P^4 / (150^4 + c_P^4))$; memory disruption: switching $\times 20$ on $[0, 40)$ or $[0, 120)$; schedule scan: 300 founders, 240 time units, 20-time-unit exposures; cycle-gated variant (Supplementary Fig. 5) adds $\beta = 0.25, \varphi_0 = 0.5, w = 0.15$ to the death hazard

Figure	Model	Parameters
5a	40 independent telegraph genes	$k_{\text{on}} = k_{\text{off}} \in [10^{-3}, 1]$ (log-spaced), $k_{\text{tx}} = 20$, $k_{\text{dm}} = 0.5$; 40 founders, 6 doublings
5b-c	random GRN	30 genes, 30 activations, 15 inhibitions, $k_{\text{tx}} \in [5, 30]$, $K \in [2, 10]$, $n = 2$, basal 0.05; 1,500 cells; sequencing capture 0.15 (cv 0.3)
5d	random GRN	20 genes, 22 activations, 10 inhibitions; knockdown factor 0.05; 800 cells
6a-b	telegraph	truth (0.3, 0.6, 20); priors log-uniform on $[0.01, 10]$, $[0.01, 10]$, $[1, 200]$; 200 (a) / 150 (b) particles
6c	resistance gene	truth $h_{\text{max}} = 0.5$, $K = 150$; priors log-uniform on $[0.05, 5]$ and $[20, 1000]$; 100 particles
7	resistance gene, hours	cell cycle 24 h (sizer, cv 0.1), $k_{\text{tx}} = 30$, $k_{\text{dm}} = 1$, $k_{\text{tl}} = 4$, $k_{\text{dp}} = 0.2$ per h, $K = 150$, $q = 4$, growth-arrest Hill coefficient 2; fitted: $k_{\text{on}}, k_{\text{off}} \in [10^{-3}, 0.3]$, $h_{\text{max}} \in [0.005, 0.5]$ per h, $\text{EC}_{50} \in [3, 40]$ μM , $m \in [1, 6]$, $\text{IC}_{50} \in [1, 40]$ μM ; 300 founders burnt in for 10 cycles, 48 h drug-free, 72 h drug, integration step $dt = 0.5$ h; 160 Latin-hypercube points and 80 Nelder-Mead iterations with seed 1; identifiability (Supplementary Fig. 3) by 11-point log-grid conditional profiles of the six parameters, a 7×7 memory $\times$ resistant-fraction slice, and the step halved to 0.25 h with three seeds, via the <code>profile</code> part of the script; the cycle-dependent variant (Supplementary Fig. 4) adds a seventh parameter $\beta \in [0, 1]$ with $\varphi_0 = 0.5$, $w = 0.15$, fitted under the same budget by the <code>cycle</code> part

Figure	Model	Parameters
8a-e	melanoma-like	$k_{tx} = 3$, $k_{dm} = 0.1$, $k_{tl} = 0.4$, $k_{dp} = 0.02$ per h; net doubling 4 weeks, memory 5 net doublings, pre-resistant fraction 0.005, $h_{max} = 0.0023$ per h, $EC_{50} = 0.3$, $m = 2$, protection $K = 150$, $q = 4$, growth arrest $IC_{50} = 0.3$ with the same protection, $k_{off} \rightarrow k_{off}/(1 + 9d)$, fitness cost 0.5, partial protection: additional unprotected growth inhibition with $IC_{50} = 1$; 2,000 founders with promoter states drawn from the stationary distribution, 60 weeks, dt 4 h; progression at $1.73 \times$ nadir after the 8-week lead-in; cycle-gated variant (Supplementary Fig. 5) adds $\beta = 0.25$, $\varphi_0 = 0.5$, $w = 0.15$
8f-h	MGMT model	same expression kinetics; net doubling 40 days, memory 4 net doublings, MGMT-expressing fraction 0.01 or 0.30, $h_{max} = 0.03$ per h at peak, $EC_{50} = 0.4$ of the standard bolus peak, $m = 2$, $K = 150$, $q = 4$; elimination half-life 2.1 h; stoichiometric MGMT consumption at rate $k d V c / (c + K_m)$ with $k = 300$, $K_m = 150$; 1,500 founders at the stationary promoter distribution, six 28-day cycles, dt 1 h

Supplementary Note 4: Additional persister simulations



Supplementary Fig. 1 | Single-cell timing distributions underlying Fig. 4c. a, Division times of cells born under drug and **b**, times from drug start (or from birth, for cells born under drug) to death, per dose.

The division-time distribution is set by the sizer and is the same at every dose, because in this model the drug kills but does not slow growth. Times to death are broad (coefficient of variation 0.72 at dose 0.5 and 0.82 at dose 2) and their mean shifts by much less than the population decay rate, because most deaths occur among low-expressing cells whose hazard is already close to $h_{\max}$ at dose 0.5 ($EC_{50} = 0.5$, $m = 2$); raising the dose mainly shortens the survival of cells with intermediate protection. The dose dependence of population decay is therefore carried by the fraction of cells that are protected, not by the kinetics of death of unprotected cells, which is the interpretation¹ give of their single-cell tracking data.

Under continuous dosing at the highest dose (release period 0, dose 2 in Fig. 4h), the long-term growth rate is +0.004 per time unit for pre-existing tolerance, -0.013 when the resistant state carries a 50% growth cost and +0.012 for drug-induced tolerance; the schedules with the lowest long-term growth rate are release period 0, dose 2 (growth rate +0.004 per time unit), release period 10, dose 2 (growth rate -0.016 per time unit) and release period 10, dose 1 (growth rate +0.009 per time unit) respectively. For pre-existing tolerance without a cost, every release period raises the net growth rate; with the fitness cost, release periods of 5 and 10 time units are within 0.003 per time unit of continuous dosing at doses of 1 and 2; for drug-induced tolerance, dose 1 gives a lower net growth rate than dose 2 at every release period, and release periods of 20 time units or more raise the growth rate in every model. All values are single realisations with 300 founder cells.


Supplementary Note 5: Benchmark details

Network inference methods: absolute Pearson and Spearman correlations of $\log(1 + x)$ (Pearson) or ranks (Spearman); GENIE3 as random-forest importances (200 trees, `sqr` features) fitted on $\log(1 + x)$ with scikit-learn. Imputation: kNN-smoothing (one step, $k = 15$, 10 principal components of the Freeman-Tukey transformed, library-size normalised counts) and MAGIC with default parameters on square-root normalised counts. Metrics: area under the precision-recall and receiver operating characteristic (ROC) curves against the undirected (correlation)

or directed (GENIE3) ground-truth adjacency. The cell-cycle-regressed dataset is the residual of $\log(1 + c)$ (concentration c) after ordinary least-squares regression on gene copy number, cell age and age squared, per gene.

Supplementary Note 6: Inference diagnostics

Summary statistics and distances. For Fig. 6a,b the distance between a simulated and the observed sample of counts is their Wasserstein-1 distance², computed as the mean absolute difference between the two empirical quantile functions on a grid of 201 probabilities, divided by the observed mean; each simulated sample contains 600 cells (a) or the final snapshot of 500 founders grown for 80 time units (b), against 2,000 observed cells. The distance between two independent samples at the true parameters (the noise floor) is reported in the summary tables of the repository. For Fig. 6c the summaries of a treated population are the log surviving fraction on a grid of 25 time points from the start of the drug to 60 time units later (populations that went extinct contribute a floor of 10^{-3}) and the concordance of death fates between sisters born in the 20 time units before the drug; the distance is the root-mean-square difference of these summaries. Each generation's tolerance is the median of the previous generation's accepted distances, particles are perturbed with a Gaussian kernel of variance twice the weighted variance of the previous generation, and the first generation samples the prior.



Supplementary Fig. 2 | ABC-SMC diagnostics. **a**, Tolerance and **b**, acceptance rate per generation for the three runs of Fig. 6.

The run on non-dividing cells (10 generations, 200 particles) reached a final tolerance of 0.148 with an acceptance rate of 1% in the last generation; the division-aware run (10 generations, 150 particles, each particle simulating a population of 500 founders for 80 time units) reached 0.0841 at 1%; the drug-parameter run (5 generations, 100 particles, 200 founders each) reached 0.0727 at 8%. Posterior predictive checks were made by simulating the forward model at the posterior median (Fig. 6d for the drug parameters); the posterior median rather than the mean is reported throughout because the posteriors are asymmetric on the log scale.

Supplementary Note 7: Runtime

Supplementary Table 3 lists the wall-clock times of the population simulations of Fig. 2f (200 steps of $dt = 0.1$ for random telegraph networks of 10, 50 and 100 genes, four threads of an Intel Xeon processor at 2.1 GHz, Julia 1.10). Throughput in cell-steps per second falls with the number of reactions because every step of every cell has to sample the events of all channels; the hybrid kernel is 2 to 5 times faster than the direct method at these tolerances, the gap widening with the number of reactions. The single-cell interface used for the comparison with JumpProcesses.jl (Fig. 2e) allocates a workspace per simulation and is slower in serial execution than JumpProcesses.jl on that one-gene model; the population interface reuses one workspace per thread.

Supplementary Table 3. Runtime of population simulations.

genes	cells	kernel	threads	seconds (200 steps)	cell-steps / s
10	100	DirectSSA	4	0.38	52,398
10	100	HybridSSATau	4	0.11	180,019
10	1000	DirectSSA	4	1.04	193,080
10	1000	HybridSSATau	4	0.58	346,004
10	5000	DirectSSA	4	4.78	209,388
10	5000	HybridSSATau	4	2.57	388,683
50	1000	DirectSSA	4	9.72	20,581
50	1000	HybridSSATau	4	2.57	77,671
100	1000	DirectSSA	4	31.41	6,367
100	1000	HybridSSATau	4	6.37	31,373

Supplementary Note 8: Laboratory and clinical reference data

Supplementary Table 4. Reference values used for calibration and validation (transcribed from the cited articles; files in `paper/data/`).

Quantity	Value	Source
U2OS cells tracked at cisplatin addition (7 / 10 / 13 μ M)	232 / 240 / 296	¹ , Table 2
of which died within 72 h	30 / 66 / 176	same
of which divided within 72 h	100 / 75 / 67	same
of which survived without dividing	102 / 99 / 53	same
HCT116 cells at drug addition; death fraction	275; 0.64	¹ , Table 1
Plateau of the HCT116 kill curve	after ~100 h	¹ , Fig 1b
Lineage correlations of fate	present for sisters, first and second cousins; absent for third cousins	¹ , Fig 5a
Pre-resistant melanoma cells	1:50 to 1:500 per marker; EGFR-high cells give 7.9 ± 0.9 fold more resistant colonies	³

Quantity	Value	Source
N15-0385 glioblastoma doubling time	50 h	4
Temozolomide elimination half-life	2.1 h	5
RTOG 0525 regimens	150-200 mg/m ² days 1-5 vs 75-100 mg/m ² days 1-21 of 28-day cycles; median OS 16.6 vs 14.9 months	6
SWOG S1320 regimens	continuous vs 3 weeks off / 5 weeks on after an 8-week lead-in; median PFS 9.0 vs 5.5 months	7

Supplementary Note 9: Migration from version 1

The v1 API (`Donne`, `ssa`, `tauleap`, `tauleapswitch`, `adaptive_tauleap`, `exponential_growth`, `NamedTuple` reactions) remains available through a deprecated compatibility layer that converts `NamedTuple` reactions to `ReactionModels` (reactions whose names contain `act`, `inhib`, `comb_a` or `comb_i` become Hill kinetics with `rate = [k, n, K]`; species whose names contain `on/off` become promoter groups) and maps the algorithms onto the new kernels. Output shapes are preserved, including the leading `:NULL` column. See `docs/src/migration.md`.

Supplementary Note 10: Relation to previous work

The ingredients of this framework have precedents, and several of its results confirm conclusions reached earlier by other means. The table separates, claim by claim, what was already established from what this work adds. Prior results the framework reproduces are evidence that it is behaving correctly, not claims of novelty.


Result in this Article	Already established	What this work adds
Stochastic reaction kinetics inside growing, dividing cells (Fig. 3)	Simulators of stochastic expression with growth and division exist ^{8,9} , including version 1 of this framework ¹⁰	Gene replication, promoter inheritance, a drug layer, lineage statistics, observation models and inference in one forward model, so that a resistance phenotype is a heritable expression state with measurable kinetics rather than an assumed compartment
Simulation of division trees for lineage-tracing methods	TedSim couples expression to division history ¹¹ ; Cassiopeia simulates topologies, heritable fitness and CRISPR barcodes ¹²	Molecular content on the same tree: partitioning at division, promoter states, drug-induced death, and per-cell fates with the molecules that caused them

Result in this Article	Already established	What this work adds
Heritable expression states decide which cells survive a drug (Fig. 4)	Measured directly in barcoded and time-lapse experiments ^{1,3,13,14}	A generative model in which memory is one promoter timescale, reproducing the reported signatures quantitatively and predicting which of them discriminate pre-existing from induced tolerance
Cell-cycle-aware inference of transcriptional kinetics is necessary (Fig. 3f, Fig. 6b)	Established from data and theory ^{15–17}	A simulator that generates the data such methods assume, and a quantification of the bias incurred by a division-blind fit
Benchmarks of network inference from single-cell data	Extensive benchmarks exist ^{10,18,19}	The growth and cell-cycle confound with known ground truth: how much accuracy is lost in a dividing population, and how much cell-cycle regression recovers (Fig. 5b)
The benefit of treatment holidays depends on a cost of resistance and on turnover (Fig. 8a-e)	Compartment models of adaptive therapy ^{20,21}	The same dependence from single cells whose resistance is a heritable expression state; cost and degree of protection become cell properties that lineage experiments measure; partial protection reverses the ranking even without a cost
Perturbing the retention of a resistant state changes the outcome (Fig. 4f)	Phenotypic-switching theory ²² and memory-disrupting compounds ¹³	The timing requirement (disruption must continue during exposure) and a lineage-resolved readout, the number of surviving clones
An intermediate dose can be optimal against drug-induced tolerance (Fig. 4h)	Reported for induced persisters ²³	Reproduced; the accompanying benefit of release periods does not appear here, which the text attributes to the decay of the induced state in this model rather than to a disagreement about data
MGMT expression selects glioblastoma cells under temozolomide (Fig. 8f-h)	Phenotypic selection with stable inheritance ⁴	Pharmacokinetics, the RTOG 0525 regimens, consumption of MGMT by the drug, and the fractionation that minimises the final population at fixed cumulative dose

Result in this Article	Already established	What this work adds
Dose-dense temozolomide does not improve survival	Clinical result ⁶	A mechanistic account for unmethylated tumours; with MGMT consumption the model instead predicts a dose-dense advantage in the methylated population, a testable divergence from the trial, and points to the rate at which the drug consumes MGMT as the measurement that separates the regimens

Supplementary Note 11: Identifiability and numerical robustness of the calibration

Six parameters were fitted to six measured fractions, so the fit cannot determine all six. Three checks bound what the calibration does and does not establish.



Supplementary Fig. 3 | What the fate fractions determine. **a**, Root-mean-square error of the simulated fate fractions at the training concentrations when one parameter is moved away from its fitted value over its full range and the others are held at theirs, plotted against the ratio to the fitted value; the dotted line marks one binomial standard error of the measurement above the best fit. **b**, The same error over the memory of the resistant state and the fraction of cells in it, with the death parameters held at their fitted values; red outlines mark the combinations that lie within one standard error of the best. **c**, Simulated fates at 13 μM with the

115 integration step halved (mean $\pm$ s.d. of three seeds).


116 The death parameters are the better determined: the maximal death rate stays within one standard error of the
117 best fit over about a twofold range, and the EC₅₀, Hill coefficient and growth-arrest IC₅₀ over narrower ranges
118 still (Supplementary Fig. 3a). The two promoter switching rates are not determined: each can move by roughly
119 sixfold with no penalty the data can detect. The two-dimensional slice shows why (Supplementary Fig. 3b). A
120 short memory with many resistant cells and a long memory with few produce nearly the same fate fractions, so
121 the error surface is flat along that direction, and a majority of the combinations tested (memory one to nine
122 generations, resistant fraction 1% to 40%) fit within the noise of the measurement; those that do span one to nine
123 generations of memory and 1% to 25% of cells resistant. The memory quoted in the main text is therefore set by
124 the prior taken from the lineage correlations, not by the fate fractions, and the resistant fraction that accompanies
125 it should not be read as a measurement.

126 Halving the integration step from 0.5 h to 0.25 h changes each simulated fate fraction by less than the seed-to-seed
127 spread (Supplementary Fig. 3c), so the discretisation of the death hazard is not a source of error at this step size.

128 A profile in which the remaining parameters are re-optimised at each fixed memory would be preferable to a slice,
129 but at seven seconds per objective evaluation the inner search could not be given enough budget to be reliable: it
130 returned fits three times worse than the known solution at short memory, which reflects the optimiser rather than
131 the data. The slice avoids that failure at the cost of conditioning on the fitted death parameters.

132 **Supplementary Note 12: Cell-cycle-dependent killing as an alternative** 133 **to heritable expression**

134 The lethality of cisplatin is enhanced during replication, so where a cell sits in its division cycle when the drug
135 arrives is a competing explanation for two of the observations the calibrated model reproduces: the near-invariance
136 of the mean time to death across concentrations, and the large fraction of cells that survive three days of drug
137 without dividing. A model in which killing is confined to a window of the cycle would spare exactly those cells
138 that never enter that window, without any heritable expression state.



139

140 **Supplementary Fig. 4 | Cell-cycle-dependent killing tested against the fate data.** **a**, Root-mean-square
 141 error of the fate fractions at the training concentrations as the cycle-independent fraction β of the death hazard is
 142 moved from 0 (killing confined to the replication window) to 1 (cycle-independent), with the remaining parameters
 143 held at the cycle-dependent fit; dotted line, one binomial standard error of the measurement above the best fit;
 144 dashed line, the fitted value. **b**, Simulated against observed fate fractions at all three concentrations for the
 145 cycle-independent and the cycle-dependent fit (squares, the held-out 10 μ M condition; line, equality).

146 Refitting under the same objective, bounds and optimiser budget gives a cycle-independent fraction of 0.34 and a
 147 root-mean-square error of 0.064 on the training fractions, against 0.061 for the cycle-independent model; at the
 148 held-out concentration the errors are 0.106 and 0.095. The profile over that fraction is flat: every value from 0.20
 149 to 0.80 lies within one binomial standard error of the measurement (0.028) of the best fit (Supplementary Fig.
 150 4a). Three fate counts per concentration therefore cannot say whether cells are spared because they inherited a
 151 protective state or because they were outside the replication window, and the parameters of either fit should be
 152 read with that in mind.


153 The two mechanisms are not exclusive, and the experiment that separates them is not a fate count: it is a
 154 measurement of cell-cycle position at the moment the drug arrives, which time-lapse imaging with a cycle reporter
 155 provides directly. Until that is done, the heritable-expression reading of the calibrated parameters should be
 156 understood as one of two accounts that these fate fractions cannot tell apart.

157 **Supplementary Note 13: Cell-cycle-dependent killing in the persister** 158 **and melanoma case studies**

159 Supplementary Note 12 shows that the cisplatin fate data cannot decide whether cells are spared because they
 160 inherited a protective state or because they were outside the replication window. The persister simulations of
 161 Fig. 4 and the melanoma simulations of Fig. 8a–e use a hazard with no cell-cycle dependence at all, so the same

question applies to them: do their conclusions rest on that choice? We repeated both with three quarters of the death hazard confined to a window around mid-cycle ($\beta = 0.25$, $\varphi_0 = 0.5$, $w = 0.15$), leaving everything else unchanged.

The window position is a modelling choice rather than a claim about either agent. The persister study uses a generic drug and commits to no mechanism; the melanoma study is built on a BRAF/MEK inhibitor, which does not create replication-coupled lesions as a platinum drug does, although it does act on cells that are traversing the cycle. What these runs test is therefore robustness to a strong cell-cycle gate of any kind, not the phase specificity of a particular drug.



170

Supplementary Fig. 5 | The case studies under cell-cycle-gated killing. **a**, Net population decay rate over the first 30 time units of treatment against dose (solid, left axis) and the mean time to death of killed cells (dashed, right axis), with the hazard cycle-blind and with three quarters of it gated to a mid-cycle window. **b**, Concordance of lineage fate between sisters and between cousins, above the value expected if fates were independent, for the memory gene and the fast-switching control. **c**, Weeks after randomisation to loss of control for the three melanoma mechanisms and the three clinical schedules (mean $\pm$ s.d. of four seeds, censored at the end of follow-up), cycle-blind (solid) and cycle-gated (pale).

Population decay still rises with dose (-0.001 to +0.023 per time unit over doses 0.5 to 2, against +0.023 to +0.050 without the gate), at roughly half the rate, and the mean time to death of killed cells stays nearly dose-invariant (8.4 to 9.9 against 6.2 to 8.5 time units), so the dose-dependence of decay alongside dose-invariant single-cell

181 timing is not an artefact of a cycle-blind hazard. The kin correlations deserve more care, because sisters are born
 182 together and therefore occupy the same phase of the cycle: a gated hazard can in principle correlate their fates
 183 with no heritable expression state at all. It does not: in the fast-switching control, which has no usable memory,
 184 the sister concordance above independence rises only from +0.000 to +0.008, while the memory gene sits at
 185 +0.054, 7-fold larger. What the gate does change is the schedule scan: the best dose is unchanged in 3 of 3 models,
 186 but the best release period is not. With a fitness cost of resistance and with drug-induced tolerance, the holidays
 187 that won without the gate collapse to continuous dosing, because a gated hazard kills the reverting cells more
 188 slowly on re-exposure. In the melanoma study the best schedule is unchanged in 3 of 3 mechanisms and the model
 189 still favours continuous dosing under partial protection (52 against 43 weeks, from 52 and 49); the full ordering of
 190 the three schedules survives in 2 of 3 mechanisms (under partial protection the two losing schedules trade places).

- 191 1. Iyer, A., Alva, A., Granada, A. E. & Chakrabarti, S. Inheritable cell-states shape drug-persister correlations
and population dynamics in cancer cells. *PLoS Computational Biology* **21**, e1013446 (2025).
- 192 2. Bernton, E., Jacob, P. E., Gerber, M. & Robert, C. P. Approximate Bayesian computation with the
Wasserstein distance. *Journal of the Royal Statistical Society: Series B* **81**, 235–269 (2019).
- 193 3. Shaffer, S. M. *et al.* Rare cell variability and drug-induced reprogramming as a mode of cancer drug
resistance. *Nature* **546**, 431–435 (2017).
- 194 4. Lasri, A. *et al.* Phenotypic selection through cell death: Stochastic modelling of O-6-methylguanine-DNA
methyltransferase dynamics. *Royal Society Open Science* **7**, 191243 (2020).
- 195 5. Ostermann, S. *et al.* Plasma and cerebrospinal fluid population pharmacokinetics of temozolomide in
malignant glioma patients. *Clinical Cancer Research* **10**, 3728–3736 (2004).
- 196 6. Gilbert, M. R. *et al.* Dose-dense temozolomide for newly diagnosed glioblastoma: A randomized phase III
clinical trial. *Journal of Clinical Oncology* **31**, 4085–4091 (2013).
- 197 7. Algazi, A. P. *et al.* Continuous versus intermittent BRAF and MEK inhibition in patients with BRAF-
mutated melanoma: A randomized phase 2 trial. *Nature Medicine* **26**, 1564–1568 (2020).
- 198 8. Bertaux, F., Marguerat, S. & Shahrezaei, V. Division rate, cell size and proteome allocation: Impact on
gene expression noise and implications for the dynamics of genetic circuits. *Royal Society Open Science* **5**,
172234 (2018).
- 199 9. Piho, P. & Thomas, P. AgentBasedModeling.jl: A tool for stochastic simulation of structured population
dynamics. *PLoS Computational Biology* **22**, e1014366 (2026).
- 200 10. Lasri, A., Shahrezaei, V. & Sturrock, M. Benchmarking imputation methods for network inference using a
novel method of synthetic scRNA-seq data generation. *BMC Bioinformatics* **23**, 236 (2022).
- 201 11. Pan, X., Li, H. & Zhang, X. TedSim: Temporal dynamics simulation of single-cell RNA sequencing data
and cell division history. *Nucleic Acids Research* **50**, 4272–4288 (2022).
- 202 12. Jones, M. G. *et al.* Inference of single-cell phylogenies from lineage tracing data using Cassiopeia. *Genome
Biology* **21**, 92 (2020).
- 203 13. Harmange, G. *et al.* Disrupting cellular memory to overcome drug resistance. *Nature Communications* **14**,
7130 (2023).
- 204 14. Oren, Y. *et al.* Cycling cancer persister cells arise from lineages with distinct programs. *Nature* **596**,
576–582 (2021).
- 205 15. Sukys, A. & Grima, R. Cell-cycle dependence of bursty gene expression: Insights from fitting mechanistic
models to single-cell RNA-seq data. *Nucleic Acids Research* **53**, gkaf295 (2025).
- 206 16. Zhang, Z. *et al.* Stochastic gene expression in proliferating cells: Differing noise intensity in single-cell and
population perspectives. *PLoS Computational Biology* **21**, e1013014 (2025).

- 207 17. Okochi, Y., Sawazaki, Y., Kondo, Y. & Naoki, H. Inferring division-associated stochasticity from time-series single-cell transcriptomes. *bioRxiv* <https://doi.org/10.64898/2026.04.14.718485> (2026) doi:10.64898/2026.04.14.718485.
- 208 18. Pratapa, A., Jalihal, A. P., Law, J. N., Bharadwaj, A. & Murali, T. M. Benchmarking algorithms for gene regulatory network inference from single-cell transcriptomic data. *Nature Methods* **17**, 147–154 (2020).
- 209 19. Dibaenia, P. & Sinha, S. SERGIO: A single-cell expression simulator guided by gene regulatory networks. *Cell Systems* **11**, 252–271 (2020).
- 210 20. Zhang, J., Cunningham, J. J., Brown, J. S. & Gatenby, R. A. Integrating evolutionary dynamics into treatment of metastatic castrate-resistant prostate cancer. *Nature Communications* **8**, 1816 (2017).
- 211 21. Strobl, M. A. R. *et al.* Turnover modulates the need for a cost of resistance in adaptive therapy. *Cancer Research* **81**, 1135–1147 (2021).
- 212 22. Gunnarsson, E. B., De, S., Leder, K. & Foo, J. Understanding the role of phenotypic switching in cancer drug resistance. *Journal of Theoretical Biology* **490**, 110162 (2020).
- 213 23. Corigliano, M., Di Bernardo, A., Cosentino Lagomarsino, M. & Pompei, S. Optimal treatment for drug-induced cancer persists involves release periods and intermediate drug doses. *PRX Life* <https://doi.org/10.1103/mh5v-lcyw> (2025) doi:10.1103/mh5v-lcyw.